

# MEMORY HIERARCHIES IN CPU/GPU ARCHITECTURES

Siegfried Höfinger

ASC Research Center, TU Wien

October 20, 2025

→ <https://tinyurl.com/cudafordummies/i/12/notes-12.pdf>

CUDA 4 DUMMIES — OCT 22-23, 2025


# OUTLINE

MEMORY HIERARCHY

TAKE HOME MESSAGES

# MEMORY HIERARCHY

## CUDA C-PROGRAMMING GUIDE, PROGRAMMING MODEL




- CUDA threads have access to several memory spaces

→ <https://docs.nvidia.com/cuda/cuda-c-programming-guide>

# MEMORY HIERARCHY

## CUDA C-PROGRAMMING GUIDE, PROGRAMMING MODEL




- CUDA threads have access to several memory spaces
- Private local memory for each individual thread (small but very fast)

→ <https://docs.nvidia.com/cuda/cuda-c-programming-guide>

# MEMORY HIERARCHY

## CUDA C-PROGRAMMING GUIDE, PROGRAMMING MODEL




- CUDA threads have access to several memory spaces
- Private local memory for each individual thread (small but very fast)
- Shared memory for all threads within a thread block (pretty fast)

→ <https://docs.nvidia.com/cuda/cuda-c-programming-guide>

# MEMORY HIERARCHY

## CUDA C-PROGRAMMING GUIDE, PROGRAMMING MODEL




- CUDA threads have access to several memory spaces
- Private local memory for each individual thread (small but very fast)
- Shared memory for all threads within a thread block (pretty fast)
- Global memory accessible by all threads (slow)

→ <https://docs.nvidia.com/cuda/cuda-c-programming-guide>

# MEMORY HIERARCHY

## CUDA C-PROGRAMMING GUIDE, PROGRAMMING MODEL




- CUDA threads have access to several memory spaces
- Private local memory for each individual thread (small but very fast)
- Shared memory for all threads within a thread block (pretty fast)
- Global memory accessible by all threads (slow)
- Special ROMs, constant and texture memory

→ <https://docs.nvidia.com/cuda/cuda-c-programming-guide>

# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, UNIFIED MEMORY PROGRAMMING



Unified Memory:

GPUs of compute capability  $\geq 6.x$  and  
 $\geq$  CUDA 8.0


Separate Device/Host Memory:

GPUs of compute capability  $< 6.x$  and  
 $<$  CUDA 8.0

→ <https://docs.nvidia.com/cuda/cuda-c-programming-guide>

# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, UNIFIED MEMORY PROGRAMMING



Unified Memory:

GPUs of compute capability  $\geq 6.x$  and  
 $\geq$  CUDA 8.0

Separate Device/Host Memory:

GPUs of compute capability  $< 6.x$  and  
 $<$  CUDA 8.0

→ <https://docs.nvidia.com/cuda/cuda-c-programming-guide>

# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, UNIFIED MEMORY PROGRAMMING CONT.

- Unified memory (managed) for recent GPUs (Pascal class or newer) greatly simplifies programming<sup>4</sup>/porting<sup>2</sup> the GPU

→ <https://docs.nvidia.com/cuda/cuda-c-programming-guide>

# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, UNIFIED MEMORY PROGRAMMING CONT.

- Unified memory (managed) for recent GPUs (Pascal class or newer) greatly simplifies programming<sup>4</sup>/porting<sup>2</sup> the GPU
- Single common address space accessible from any processor in a system

→ <https://docs.nvidia.com/cuda/cuda-c-programming-guide>

# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, UNIFIED MEMORY PROGRAMMING CONT.

- Unified memory (managed) for recent GPUs (Pascal class or newer) greatly simplifies programming<sup>4</sup>/porting<sup>2</sup> the GPU
- Single common address space accessible from any processor in a system
- Strict distinction into device- and host-memory is getting irrelevant to the programmer

→ <https://docs.nvidia.com/cuda/cuda-c-programming-guide>

# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, UNIFIED MEMORY PROGRAMMING CONT.

- Unified memory (managed) for recent GPUs (Pascal class or newer) greatly simplifies programming<sup>4</sup>/porting<sup>2</sup> the GPU
- Single common address space accessible from any processor in a system
- Strict distinction into device- and host-memory is getting irrelevant to the programmer
- Single-pointer-to-data model, no longer any need to copy forth/back data into/from GPU memory

# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, UNIFIED MEMORY PROGRAMMING CONT.

- Unified memory (managed) for recent GPUs (Pascal class or newer) greatly simplifies programming<sup>4</sup>/porting<sup>2</sup> the GPU
- Single common address space accessible from any processor in a system
- Strict distinction into device- and host-memory is getting irrelevant to the programmer
- Single-pointer-to-data model, no longer any need to copy forth/back data into/from GPU memory
- On-demand page migration (hardware supported) — Page Migration Engine

# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, UNIFIED MEMORY PROGRAMMING CONT.

- Unified memory (managed) for recent GPUs (Pascal class or newer) greatly simplifies programming<sup>4</sup>/porting<sup>2</sup> the GPU
- Single common address space accessible from any processor in a system
- Strict distinction into device- and host-memory is getting irrelevant to the programmer
- Single-pointer-to-data model, no longer any need to copy forth/back data into/from GPU memory
- On-demand page migration (hardware supported) — Page Migration Engine
- Simply invoked via `cudaMallocManaged()`

# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, UNIFIED MEMORY PROGRAMMING CONT.

### Multiple Thread Blocks Matrix Addition

```
__global__ void MatAdd( float **A, float **B, float **C )
{
    int i, j;
    i = (blockIdx.x * blockDim.x) + threadIdx.x;
    j = (blockIdx.y * blockDim.y) + threadIdx.y;
    C[i][j] = A[i][j] + B[i][j];
}

int main()
{
    int i;
    dim3 threadsPerBlock, numBlocks;
    float **A, **B, **C;
    // unified memory allocation, B and C analogous
    cudaMallocManaged(&A, N * sizeof(float *));
    for (i = 0; i < N; i++) {
        cudaMallocManaged(&A[i], N * sizeof(float));
    }
    // kernel invocation
    MatAdd <<< numBlocks, threadsPerBlock >>> (A, B, C);
    cudaDeviceSynchronize();
    ...
    cudaFree(A);
}
```

# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, UNIFIED MEMORY PROGRAMMING CONT.

### Multiple Thread Blocks Matrix Addition

```
__global__ void MatAdd( float **A, float **B, float **C )
{
    int i, j;
    i = (blockIdx.x * blockDim.x) + threadIdx.x;
    j = (blockIdx.y * blockDim.y) + threadIdx.y;
    C[i][j] = A[i][j] + B[i][j];
}

int main()
{
    int i;
    dim3 threadsPerBlock, numBlocks;
    float **A, **B, **C;
    // unified memory allocation, B and C analogous
    cudaMallocManaged(&A, N * sizeof(float *));
    for (i = 0; i < N; i++) {
        cudaMallocManaged(&A[i], N * sizeof(float));
    }
    // kernel invocation
    MatAdd <<< numBlocks, threadsPerBlock >>> (A, B, C);
    cudaDeviceSynchronize();
    ...
    cudaFree(A);
}
```

Standard  
malloc-like  
allocation  
on the host



# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, UNIFIED MEMORY PROGRAMMING CONT.

### Multiple Thread Blocks Matrix Addition

Standard  
malloc-like  
allocation  
on the host

```
__global__ void MatAdd( float **A, float **B, float **C ) {  
    int i, j;  
    i = (blockIdx.x * blockDim.x) + threadIdx.x;  
    j = (blockIdx.y * blockDim.y) + threadIdx.y;  
    C[i][j] = A[i][j] + B[i][j];  
}  
  
int main()  
{  
    int i;  
    dim3 threadsPerBlock, numBlocks;  
    float **A, **B, **C;  
    // unified memory allocation, B and C analogous  
    cudaMallocManaged(&A, N * sizeof(float *));  
    for (i = 0; i < N; i++) {  
        cudaMallocManaged(&A[i], N * sizeof(float));  
    }  
    // kernel invocation  
    MatAdd <<< numBlocks, threadsPerBlock >>> (A, B, C);  
    cudaDeviceSynchronize();  
    ...  
    cudaFree(A);  
}
```

host point-  
ers directly  
usable in  
kernel code

# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, UNIFIED MEMORY PROGRAMMING CONT.

### Multiple Thread Blocks Matrix Addition

```
__global__ void MatAdd( float **A, float **B, float **C ) {  
    int i, j;  
    i = (blockIdx.x * blockDim.x) + threadIdx.x;  
    j = (blockIdx.y * blockDim.y) + threadIdx.y;  
    C[i][j] = A[i][j] + B[i][j];  
}  
  
int main()  
{  
    int i;  
    dim3 threadsPerBlock, numBlocks;  
    float **A, **B, **C;  
    // unified memory allocation, B and C analogous  
    cudaMallocManaged(&A, N * sizeof(float *));  
    for (i = 0; i < N; i++) {  
        cudaMallocManaged(&A[i], N * sizeof(float));  
    }  
    // kernel invocation  
    MatAdd <<< numBlocks, threadsPerBlock >>> (A, B, C);  
    cudaDeviceSynchronize();  
    ...  
    cudaFree(A);  
}
```

Standard malloc-like allocation on the host

host pointers directly usable in kernel code

ensure proper kernel completion

# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, UNIFIED MEMORY PROGRAMMING CONT.

### Multiple Thread Blocks Matrix Addition

```
__global__ void MatAdd( float **A, float **B, float **C ) {  
    int i, j;  
    i = (blockIdx.x * blockDim.x) + threadIdx.x;  
    j = (blockIdx.y * blockDim.y) + threadIdx.y;  
    C[i][j] = A[i][j] + B[i][j];  
}  
  
int main()  
{  
    int i;  
    dim3 threadsPerBlock, numBlocks;  
    float **A, **B, **C;  
    // unified memory allocation, B and C analogous  
    cudaMallocManaged(&A, N * sizeof(float *));  
    for (i = 0; i < N; i++) {  
        cudaMallocManaged(&A[i], N * sizeof(float));  
    }  
    // kernel invocation  
    MatAdd <<< numBlocks, threadsPerBlock >>> (A, B, C);  
    cudaDeviceSynchronize();  
    ...  
    cudaFree(A);  
}
```

Standard  
malloc-like  
allocation  
on the host

free memory  
when done

host point-  
ers directly  
usable in  
kernel code

ensure  
proper ker-  
nel comple-  
tion

# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, UNIFIED MEMORY PROGRAMMING CONT.

### 2nd Form — Managed Global Memory

```
#define ARRAYDIM 268435456

// global managed declaration on GPU
__device__ __managed__ float x[ARRAYDIM], y[ARRAYDIM], z[ARRAYDIM];

__global__ void KrnlDmmmy()
{
    int i;
    i = (blockIdx.x * blockDim.x) + threadIdx.x;
    x[i] = y[i] + z[i];
    return;
}

int main()
{
    ...
    // kernel invocation
    KrnlDmmmy <<< numBlocks, threadsPerBlock >>> ();
    cudaDeviceSynchronize();
    ...
}
```

# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, UNIFIED MEMORY PROGRAMMING CONT.

### 2nd Form — Managed Global Memory

Global variables directly usable on device & host

```
#define ARRAYDIM 268435456

// global managed declaration on GPU
__device__ __managed__ float x[ARRAYDIM], y[ARRAYDIM], z[ARRAYDIM];

__global__ void KrnlDmmy()
{
    int i;
    i = (blockIdx.x * blockDim.x) + threadIdx.x;
    x[i] = y[i] + z[i];
    return;
}

int main()
{
    ...
    // kernel invocation
    KrnlDmmy <<< numBlocks, threadsPerBlock >>> ();
    cudaDeviceSynchronize();
    ...
}
```

# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, UNIFIED MEMORY PROGRAMMING CONT.

### 2nd Form — Managed Global Memory

Global variables directly usable on device & host

```
#define ARRAYDIM 268435456

// global managed declaration on GPU
__device__ __managed__ float x[ARRAYDIM], y[ARRAYDIM], z[ARRAYDIM];

__global__ void KrnlDmmy()
{
    int i;
    i = (blockIdx.x * blockDim.x) + threadIdx.x;
    x[i] = y[i] + z[i];
    return;
}

int main()
{
    ...
    // kernel invocation
    KrnlDmmy <<< numBlocks, threadsPerBlock >>> ();
    cudaDeviceSynchronize();
    ...
}
```

void kernel call

# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, UNIFIED MEMORY PROGRAMMING CONT.

### 2nd Form — Managed Global Memory

Global variables directly usable on device & host

```
#define ARRAYDIM 268435456

// global managed declaration on GPU
__device__ __managed__ float x[ARRAYDIM], y[ARRAYDIM], z[ARRAYDIM];

__global__ void KrnlDmmmy()
{
    int i;
    i = (blockIdx.x * blockDim.x) + threadIdx.x;
    x[i] = y[i] + z[i];
    return;
}

int main()
{
    ...
    // kernel invocation
    KrnlDmmmy <<< numBlocks, threadsPerBlock >>> (0);
    cudaDeviceSynchronize();
    ...
}
```

void kernel call

ensure proper kernel completion

# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, UNIFIED MEMORY PROGRAMMING CONT.

### Unified Memory Example Version 1

```
#define ARRAYDIM 268435456

__global__ void KrnlDmmy(float *x, float *y, float *z)
{
    int i;
    i = (blockIdx.x * blockDim.x) + threadIdx.x;
    x[i] = (float) i;
    y[i] = (float) (i + 1);
    z[i] = x[i] + y[i];
    return;
}

int main()
{
    ...
    // unified memory allocation, b and c analogous
    cudaMallocManaged(&a, ARRAYDIM * sizeof(float));
    // kernel invocation
    KrnlDmmy <<< numBlocks, threadsPerBlock >>> (a, b, c);
    cudaDeviceSynchronize();
    ...
}
```

# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, UNIFIED MEMORY PROGRAMMING CONT.

### Unified Memory Example Version 1

```
#define ARRAYDIM 268435456

__global__ void KrnlDmmy(float *x, float *y, float *z)
{
    int i;
    i = (blockIdx.x * blockDim.x) + threadIdx.x;
    x[i] = (float) i;
    y[i] = (float) (i + 1); ←
    z[i] = x[i] + y[i];
    return;
}

int main()
{
    ...
    // unified memory allocation, b and c analogous
    cudaMallocManaged(&a, ARRAYDIM * sizeof(float));
    // kernel invocation
    KrnlDmmy <<< numBlocks, threadsPerBlock >>> (a, b, c);
    cudaDeviceSynchronize();
    ...
}
```

Kernel initialization and calculation

# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, UNIFIED MEMORY PROGRAMMING CONT.

### Unified Memory Example Version 1

```
#define ARRAYDIM 268435456

__global__ void KrlIDmmy(float *x, float *y, float *z)
{
    int i;
    i = (blockIdx.x * blockDim.x) + threadIdx.x;
    x[i] = (float) i;
    y[i] = (float) (i + 1); ←
    z[i] = x[i] + y[i];
    return;
}

int main()
{
    ...
    // unified memory allocation, b and c analogous
    cudaMallocManaged(&a, ARRAYDIM * sizeof(float)); ←
    // kernel invocation
    KrlIDmmy <<< numBlocks, threadsPerBlock >>> (a, b, c);
    cudaDeviceSynchronize();
    ...
}
```

Kernel initialization and calculation

1 GB arrays

# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, UNIFIED MEMORY PROGRAMMING CONT.

```
cuda-zen sh@n3073-004:~$ nvcc ./unified_memory_example_1.cu
cuda-zen sh@n3073-004:~$ ./a.out
cuda-zen sh@n3073-004:~$ nsys nvprof ./a.out
```

[4/7] Executing 'cudaapisum' stats report

Time (%)	Total Time (ns)	Num Calls	Avg (ns)	Med (ns)	Min (ns)	Max (ns)	StdDev (ns)	Name
39.5	238225943	1	238225943.0	238225943.0	238225943	238225943	0.0	cudaDeviceSynchronize
38.9	234132874	3	78044291.3	43351.0	22130	234067393	135119970.0	cudaMallocManaged
21.5	129631998	3	43210666.0	37602092.0	37551931	54477975	9757808.1	cudaFree
0.1	417977	1	417977.0	417977.0	417977	417977	0.0	cudaLaunchKernel

[5/7] Executing 'gpukernsum' stats report

Time (%)	Total Time (ns)	Instances	Avg (ns)	Med/Min/Max	StdDev (ns)	GridXYZ	BlockXYZ	Name
100.0	238223963	1	238223963.0	238223963.0	0.0	1048576	1 1	KrnldMmmy(float *, float *, float *)

→ [https://tinyurl.com/cudafordummies/i/l2/unified\\_memory\\_example\\_1.cu](https://tinyurl.com/cudafordummies/i/l2/unified_memory_example_1.cu)

# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, UNIFIED MEMORY PROGRAMMING CONT.

```
cuda-zen sh@n3073-004:~$ nvcc ./unified_memory_example_1.cu
cuda-zen sh@n3073-004:~$ ./a.out
cuda-zen sh@n3073-004:~$ nsys nvprof ./a.out
```

[4/7] Executing 'cudaapisum' stats report

Time (%)	Total Time (ns)	Num Calls	Avg (ns)	Med (ns)	Min (ns)	Max (ns)	StdDev (ns)	Name
39.5	238225943	1	238225943.0	238225943.0	238225943	238225943	0.0	cudaDeviceSynchronize
38.9	234132874	3	78044291.3	43351.0	22130	234067393	135119970.0	cudaMallocManaged
21.5	129631998	3	43210666.0	37602092.0	37551931	54477975	9757808.1	cudaFree
0.1	417977	1	417977.0	417977.0	417977	417977	0.0	cudaLaunchKernel

[5/7] Executing 'gpukernsum' stats report

Time (%)	Total Time (ns)	Instances	Avg (ns)	Med/Min/Max	StdDev (ns)	GridXYZ	BlockXYZ	Name
100.0	238223963	1	238223963.0	238223963.0	0.0	1048576	1 1	KrnldMmmy(float *, float *, float *)

→ [https://tinyurl.com/cudafordummies/i/l2/unified\\_memory\\_example\\_1.cu](https://tinyurl.com/cudafordummies/i/l2/unified_memory_example_1.cu)

## MEMORY HIERARCHY CONT.

### CUDA C-PROGRAMMING GUIDE, UNIFIED MEMORY PROGRAMMING CONT.

- Straightforward profiling with `nsys nvprof` on the command line, however, suppressed information about unified memory, especially numbers of page faults etc.

→ <https://devblogs.nvidia.com/unified-memory-cuda-beginners>

## MEMORY HIERARCHY CONT.

### CUDA C-PROGRAMMING GUIDE, UNIFIED MEMORY PROGRAMMING CONT.

- Straightforward profiling with `nsys nvprof` on the command line, however, suppressed information about unified memory, especially numbers of page faults etc.
- In a real application, the GPU is likely to use the data more often and for more complex computations, so the overhead from page faulting may become negligible; to quantify, we...

# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, UNIFIED MEMORY PROGRAMMING CONT.

- Straightforward profiling with `nsys nvprof` on the command line, however, suppressed information about unified memory, especially numbers of page faults etc.
  - In a real application, the GPU is likely to use the data more often and for more complex computations, so the overhead from page faulting may become negligible; to quantify, we...
1. Could separate the actual calculation from the initialization with the help of a second kernel

→ <https://devblogs.nvidia.com/unified-memory-cuda-beginners>

# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, UNIFIED MEMORY PROGRAMMING CONT.

- Straightforward profiling with `nsys nvprof` on the command line, however, suppressed information about unified memory, especially numbers of page faults etc.
- In a real application, the GPU is likely to use the data more often and for more complex computations, so the overhead from page faulting may become negligible; to quantify, we...
  1. Could separate the actual calculation from the initialization with the help of a second kernel
  2. Could repeat the kernel doing the calculation many times

→ <https://devblogs.nvidia.com/unified-memory-cuda-beginners>

# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, UNIFIED MEMORY PROGRAMMING CONT.

- Straightforward profiling with `nsys nvprof` on the command line, however, suppressed information about unified memory, especially numbers of page faults etc.
- In a real application, the GPU is likely to use the data more often and for more complex computations, so the overhead from page faulting may become negligible; to quantify, we...
  1. Could separate the actual calculation from the initialization with the help of a second kernel
  2. Could repeat the kernel doing the calculation many times
  3. Could use unified memory prefetching to explicitly move the data to the GPU

→ <https://devblogs.nvidia.com/unified-memory-cuda-beginners>

# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, UNIFIED MEMORY PROGRAMMING CONT.

```
cuda-zen sh@n3073-004:~$ nvcc ./unified_memory_example_2.cu
cuda-zen sh@n3073-004:~$ nsys nvprof ./a.out
```

[4/7] Executing 'cudaapisum' stats report

Time (%)	Total Time (ns)	Num Calls	Avg (ns)	Med (ns)	Min (ns)	Max (ns)	StdDev (ns)	Name
43.0	262687803	2	131343901.5	131343901.5	80061098	182626705	72524836.2	cudaDeviceSynchronize
36.3	221685218	3	73895072.7	37931.0	23340	221623947	127936958.3	cudaMallocManaged
19.5	118880186	3	39626728.7	38862579.0	37921894	42095713	2189322.9	cudaFree
1.2	7069296	2	3534648.0	3534648.0	26820	7042476	4960817.9	cudaLaunchKernel

[5/7] Executing 'gpukernsum' stats report

Time (%)	Total Time (ns)	Instances	Avg (ns)	Med/Min/Max	StdDev (ns)	GridXYZ	BlockXYZ	Name
69.5	182625001	1	182625001.0	182625001.0	0.0	1048576	1 1 256	KrnldmmyInit(float *, float *, float *)
30.5	80056695	1	<u>80056695.0</u>	80056695.0	0.0	1048576	1 1 256	KrnldmmyCalc(float *, float *, float *)

→ [https://tinyurl.com/cudafordummies/i/l2/unified\\_memory\\_example\\_2.cu](https://tinyurl.com/cudafordummies/i/l2/unified_memory_example_2.cu)

# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, UNIFIED MEMORY PROGRAMMING CONT.

```
cuda-zen sh@n3073-004:~$ nvcc ./unified_memory_example_2.cu
cuda-zen sh@n3073-004:~$ nsys nvprof ./a.out
```

[4/7] Executing 'cudaapisum' stats report

Time (%)	Total Time (ns)	Num Calls	Avg (ns)	Med (ns)	Min (ns)	Max (ns)	StdDev (ns)	Name
43.0	262687803	2	131343901.5	131343901.5	80061098	182626705	72524836.2	cudaDeviceSyncro
36.3	221685218	3	73895072.7	37931.0	23340	221623947	127936958.3	cudaMallocManaged
19.5	118880186	3	39626728.7	38862579.0	37921894	42095713	2189322.9	cudaFree
1.2	7069296	2	3534648.0	3534648.0	26820	7042476	4960817.9	cudaLaunchKernel

Compute kernel much faster now !

[5/7] Executing 'gpukernsum' stats report

Time (%)	Total Time (ns)	Instances	Avg (ns)	Med/Min/Max	StdDev (ns)	GridXYZ	BlockXYZ	Name
69.5	182625001	1	182625001.0	182625001.0	0.0	1048576	1 1 256	KrnldmmyInit(float *, float *, float *)
30.5	80056695	1	80056695.0	80056695.0	0.0	1048576	1 1 256	KrnldmmyCalc(float *, float *, float *)

→ [https://tinyurl.com/cudafordummies/i/l2/unified\\_memory\\_example\\_2.cu](https://tinyurl.com/cudafordummies/i/l2/unified_memory_example_2.cu)

# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, UNIFIED MEMORY PROGRAMMING CONT.

```
cuda-zen sh@n3073-004:~$ nvcc ./unified_memory_example_2.cu
cuda-zen sh@n3073-004:~$ nsys nvprof ./a.out
```

[4/7] Executing 'cudaapisum' stats report

Time (%)	Total Time (ns)	Min (ns)	Max (ns)	StdDev (ns)	Name
43.0	262	80061098	182626705	72524836.2	cudaDeviceSynchron
36.3	221	23340	221623947	127936958.3	cudaMallocManaged
19.5	118	37921894	42095713	2189322.9	cudaFree
1.2	7	26820	7042476	4960817.9	cudaLaunchKernel

Still very low effective memory bandwidth,  
 $3 \times 268435456 \times 4 \times 10^9 = 80056695 \times 10^{-9}$  = 40.24GB/s  
from theoretical 1600GB/s

Compute kernel much faster now !

[5/7] Executing 'gpu

Time (%)	Total Time (ns)	Instances	Avg (ns)	Med/Min/Max	StdDev (ns)	GridXYZ	BlockXYZ	Name
69.5	182625001	1	182625001.0	182625001.0	0.0	1048576	1 1 256	KrnldmmyInit(float *, float *, float *)
30.5	80056695	1	80056695.0	80056695.0	0.0	1048576	1 1 256	KrnldmmyCalc(float *, float *, float *)

→ [https://tinyurl.com/cudafordummies/i/l2/unified\\_memory\\_example\\_2.cu](https://tinyurl.com/cudafordummies/i/l2/unified_memory_example_2.cu)

# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, UNIFIED MEMORY PROGRAMMING CONT.

```
cuda-zen sh@n3073-004:~$ nvcc ./unified_memory_example_3.cu
cuda-zen sh@n3073-004:~$ nsys nvprof ./a.out
```

[4/7] Executing 'cudaapisum' stats report

Time (%)	Total Time (ns)	Num Calls	Avg (ns)	Med (ns)	Min (ns)	Max (ns)	StdDev (ns)	Name
59.4	502130428	101	4971588.4	2366751.0	2362901	180230574	19454564.9	cudaDeviceSynchronize
26.9	227483529	3	75827843.0	40981.0	23201	227419347	131282093.8	cudaMallocManaged
13.5	114443761	3	38147920.3	38122817.0	38117096	38203848	48519.2	cudaFree
0.1	893207	101	8843.6	3850.0	3650	457938	45211.6	cudaLaunchKernel

[5/7] Executing 'gpukernsum' stats report

Time (%)	Total Time (ns)	Instances	Avg (ns)	Med (ns)	Min (ns)	Max (ns)	StdDev (ns)	GridXYZ	BlockXYZ	Name
64.1	321527279	100	3215272.8	2363150.5	2359535	85412082	8302785.2	1048576	1 1	KrnldMmmyCalc(float *, float
35.9	180201758	1	180201758.0	180201758.0	180201758	180201758	0.0	1048576	1 1	KrnldMmmyInit(float *, float

→ [https://tinyurl.com/cudafordummies/i/l2/unified\\_memory\\_example\\_3.cu](https://tinyurl.com/cudafordummies/i/l2/unified_memory_example_3.cu)

# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, UNIFIED MEMORY PROGRAMMING CONT.

```
cuda-zen sh@n3073-004:~$ nvcc ./unified_memory_example_3.cu
cuda-zen sh@n3073-004:~$ nsys nvprof ./a.out
```

[4/7] Executing 'cudaapisum' stats report

Time (%)	Total Time (ns)	Num Calls	Avg (ns)	Med (ns)	Min (ns)	Max (ns)	StdDev (ns)
59.4	502130428	101	4971588.4	2366751.0	2362901	180230574	19454564.9
26.9	227483529	3	75827843.0	40981.0	23201	227419347	131282093.8
13.5	114443761	3	38147920.3	38122817.0	38117096	38203848	48519.2
0.1	893207	101	8843.6	3850.0	3650	457938	45211.6

100× greatly improves effective memory bandwidth, 1002GB/s and compute performance

[5/7] Executing 'gpukernsum' stats report

Time (%)	Total Time (ns)	Instances	Avg (ns)	Med (ns)	Min (ns)	Max (ns)	StdDev (ns)	GridXYZ	BlockXYZ	Name
64.1	321527279	100	3215272.8	2363150.5	2359535	85412082	8302785.2	1048576	1 1	KrnldMmmyCalc(float *, float)
35.9	180201758	1	180201758.0	180201758.0	180201758	180201758	0.0	1048576	1 1	KrnldMmmyInit(float *, float)

→ [https://tinyurl.com/cudafordummies/i/l2/unified\\_memory\\_example\\_3.cu](https://tinyurl.com/cudafordummies/i/l2/unified_memory_example_3.cu)

# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, UNIFIED MEMORY PROGRAMMING CONT.

```
cuda-zen sh@n3073-004:~$ nvcc ./unified_memory_example_4.cu
cuda-zen sh@n3073-004:~$ nsys nvprof ./a.out
```

[4/7] Executing 'cudaapisum' stats report

Time (%)	Total Time (ns)	Num Calls	Avg (ns)	Med (ns)	Min (ns)	Max (ns)	StdDev (ns)	Name
55.1	422794342	101	4186082.6	2430620.0	2362338	181284947	17798470.6	cudaDeviceSynchronize
29.2	224184902	3	74728300.7	38891.0	23080	224122931	129379545.3	cudaMallocManaged
14.8	113834964	3	37944988.0	37947335.0	37937445	37950184	6685.9	cudaFree
0.8	6191521	3	2063840.3	1278401.0	960186	3952934	1643721.8	cudaMemPrefetchAsync
0.1	891612	101	8827.8	3800.0	3600	474048	46837.1	cudaLaunchKernel

[5/7] Executing 'gpukernsum' stats report

Time (%)	Total Time (ns)	Instances	Avg (ns)	Med (ns)	Min (ns)	Max (ns)	StdDev (ns)	GridXYZ	BlockXYZ	Name
57.0	240257784	100	2402577.8	2426684.5	2358605	2455405	32694.9	1048576	1 1	KrnldMmmyCalc(float *, float
43.0	181291027	1	181291027.0	181291027.0	181291027	181291027	0.0	1048576	1 1	KrnldMmmyInit(float *, float

→ [https://tinyurl.com/cudafordummies/i/l2/unified\\_memory\\_example\\_4.cu](https://tinyurl.com/cudafordummies/i/l2/unified_memory_example_4.cu)

# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, UNIFIED MEMORY PROGRAMMING CONT.

```
cuda-zen sh@n3073-004:~$ nvcc ./unified_memory_example_4.cu
cuda-zen sh@n3073-004:~$ nsys nvprof ./a.out
```

[4/7] Executing 'cudaapisum' stats report

Time (%)	Total Time (ns)	Num Calls	Avg (ns)	Med (ns)	Min (ns)	Max (ns)	StdDev (ns)	Name
55.1	422794342	101	4186082.6	2430620.0	2362338	181284947	17798470.6	cu
29.2	224184902	3	74728300.7	38891.0	23080	224122931	129379545.3	cu
14.8	113834964	3	37944988.0	37947335.0	37937445	37950184	6685.9	cu
0.8	6191521	3	2063840.3	1278401.0	960186	3952934	1643721.8	cu
0.1	891612	101	8827.8	3800.0	3600	474048	46837.1	cu

Prefetching: fastest,  
1341 GB/s

[5/7] Executing 'gpukernsum' stats report

Time (%)	Total Time (ns)	Instances	Avg (ns)	Med (ns)	Min (ns)	Max (ns)	StdDev (ns)	GridXYZ	BlockXYZ	Name
57.0	240257784	100	2402577.8	2426684.5	2358605	2455405	32694.9	1048576	1 1	KrnldMmmyCalc(float *, float
43.0	181291027	1	181291027.0	181291027.0	181291027	181291027	0.0	1048576	1 1	KrnldMmmyInit(float *, float

→ [https://tinyurl.com/cudafordummies/i/l2/unified\\_memory\\_example\\_4.cu](https://tinyurl.com/cudafordummies/i/l2/unified_memory_example_4.cu)

# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, UNIFIED MEMORY PROGRAMMING CONT.

```
cuda-zen sh@n3073-004:~$ nvcc ./unified_memory_example_5.cu
cuda-zen sh@n3073-004:~$ nsys nvprof ./a.out
```

[4/7] Executing 'cudaapisum' stats report

Time (%)	Total Time (ns)	Num Calls	Avg (ns)	Med (ns)	Min (ns)	Max (ns)	StdDev (ns)	Name
65.0	439291852	101	4349424.3	3800.0	3590	438878684	43669650.2	cudaLaunchKernel
35.0	236640315	101	2342973.4	2352459.0	1377152	2358178	97072.8	cudaDeviceSynchronize

[5/7] Executing 'gpukernsum' stats report

Time (%)	Total Time (ns)	Instances	Avg (ns)	Med (ns)	Min (ns)	Max (ns)	StdDev (ns)	GridXYZ	BlockXYZ	Name
99.4	234905893	100	2349058.9	2348844.5	2345838	2355021	1351.8	1048576	1 1 256	KrnldMmmyCalc()
0.6	1385416	1	1385416.0	1385416.0	1385416	1385416	0.0	1048576	1 1 256	KrnldMmmyInit()

→ [https://tinyurl.com/cudafordummies/i/l2/unified\\_memory\\_example\\_5.cu](https://tinyurl.com/cudafordummies/i/l2/unified_memory_example_5.cu)

# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, UNIFIED MEMORY PROGRAMMING CONT.

```
cuda-zen sh@n3073-004:~$ nvcc ./unified_memory_example_5.cu
cuda-zen sh@n3073-004:~$ nsys nvprof ./a.out
```

[4/7] Executing 'cudaapisum' stats report

Time (%)	Total Time (ns)	Num Calls	Avg (ns)	Med (ns)	Min (ns)	Max (ns)	StdDev (ns)	Name
65.0	439291852	101	4349424.3	3800.0	3590	438878684	43669650.2	cuda
35.0	236640315	101	2342973.4	2352459.0	1377152	2358178	97072.8	cuda

Globally managed ex  
aequo, 1371GB/s

[5/7] Executing 'gpukernsum' stats report

Time (%)	Total Time (ns)	Instances	Avg (ns)	Med (ns)	Min (ns)	Max (ns)	StdDev (ns)	GridXYZ	BlockXYZ	Name
99.4	234905893	100	2349058.9	2348844.5	2345838	2355021	1351.8	1048576	1 1 256	KrnldMmmyCalc()
0.6	1385416	1	1385416.0	1385416.0	1385416	1385416	0.0	1048576	1 1 256	KrnldMmmyInit()

→ [https://tinyurl.com/cudafordummies/i/l2/unified\\_memory\\_example\\_5.cu](https://tinyurl.com/cudafordummies/i/l2/unified_memory_example_5.cu)

# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, UNIFIED MEMORY PROGRAMMING CONT.

### GPU Memory Oversubscription

```
#define DBLEARRAY24GB 3221225472

__global__ void KrnlDmmmy(double *a, double *b, double *c)
{
    unsigned int i;
    i = (blockIdx.x * blockDim.x) + threadIdx.x;
    a[i] = (double) i;
    b[i] = (double) 3221225472 - i;
    c[i] = a[i] + b[i];
    return;
}

int main()
{
    ...
    // unified memory allocation a[], b[], c[]
    cudaMallocManaged(&a, DBLEARRAY24GB * sizeof(double));
    // kernel invocation
    KrnlDmmmy <<< numBlocks, threadsPerBlock >>> (a, b, c);
    cudaDeviceSynchronize();
    check_array(c);
    cudaFree(c); cudaFree(b); cudaFree(a);
    return(0);
}
```

# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, UNIFIED MEMORY PROGRAMMING CONT.

### GPU Memory Oversubscription

```
#define DBLEARRAY24GB 3221225472

__global__ void KrnlDmmy(double *a, double *b, double *c)
{
    unsigned int i;
    i = (blockIdx.x * blockDim.x) + threadIdx.x;
    a[i] = (double) i;
    b[i] = (double) 3221225472 - i;
    c[i] = a[i] + b[i];
    return;
}

int main()
{
    ...
    // unified memory allocation a[], b[], c[]
    cudaMallocManaged(&a, DBLEARRAY24GB * sizeof(double)); ←
    // kernel invocation
    KrnlDmmy <<< numBlocks, threadsPerBlock >>> (a, b, c);
    cudaDeviceSynchronize();
    check_array(c);
    cudaFree(c); cudaFree(b); cudaFree(a);
    return(0);
}
```

24 GB array(s)

# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, UNIFIED MEMORY PROGRAMMING CONT.

### GPU Memory Oversubscription

```
#define DBLEARRAY24GB 3221225472

__global__ void KrnlDmmmy(double *a, double *b, double *c)
{
    unsigned int i;
    i = (blockIdx.x * blockDim.x) + threadIdx.x;
    a[i] = (double) i;
    b[i] = (double) 3221225472 - i;
    c[i] = a[i] + b[i];
    return;
}

int main()
{
    ...
    // unified memory allocation a[], b[], c[]
    cudaMallocManaged(&a, DBLEARRAY24GB * sizeof(double));
    // kernel invocation
    KrnlDmmmy <<< numBlocks, threadsPerBlock >>> (a, b, c);
    cudaDeviceSynchronize();
    check_array(c);
    cudaFree(c); cudaFree(b); cudaFree(a);
    return(0);
}
```

Kernel calculation (40 GB onboard)

24 GB array(s)

# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, UNIFIED MEMORY PROGRAMMING CONT.

### GPU Memory Oversubscription

```
#define DBLEARRAY24GB 3221225472

__global__ void KrnlDmmmy(double *a, double *b, double *c)
{
    unsigned int i;
    i = (blockIdx.x * blockDim.x) + threadIdx.x;
    a[i] = (double) i;
    b[i] = (double) 3221225472 - i;
    c[i] = a[i] + b[i];
    return;
}

int main()
{
    ...
    // unified memory allocation a[], b[], c[]
    cudaMallocManaged(&a, DBLEARRAY24GB * sizeof(double));
    // kernel invocation
    KrnlDmmmy <<< numBlocks, threadsPerBlock >>> (a, b, c);
    cudaDeviceSynchronize();
    check_array(c);
    cudaFree(c); cudaFree(b); cudaFree(a);
    return(0);
}
```

Simple correctness check

Kernel calculation (40 GB onboard)

24 GB array(s)

# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, UNIFIED MEMORY PROGRAMMING CONT.

```
cuda-zen sh@n3073-004:~$ nvcc ./unified_memory_oversubscription.cu
cuda-zen sh@n3073-004:~$ nsys nvprof ./a.out
```

[5/7] Executing 'gpukernsum' stats report

Time (%)	Total Time (ns)	Instances	Avg (ns)	Med (ns)	Min (ns)	Max (ns)	StdDev (ns)	GridXYZ	BlockXYZ
100.0	6274870199	1	6274870199.0	6274870199.0	6274870199	6274870199	0.0	6291456	1 1 1

[6/7] Executing 'gpumemtimesum' stats report

Time (%)	Total Time (ns)	Count	Avg (ns)	Med (ns)	Min (ns)	Max (ns)	StdDev (ns)	Operation
100.0	5909981859	96714	61107.8	6656.0	2495	556067	97099.5	[CUDA Unified Memory memcpy DtoH]

[7/7] Executing 'gpumemsizesum' stats report

Total (MB)	Count	Avg (MB)	Med (MB)	Min (MB)	Max (MB)	StdDev (MB)	Operation
49417.290	96714	0.511	0.061	0.004	2.097	0.780	[CUDA Unified Memory memcpy DtoH]

→ [https://tinyurl.com/cudafordummies/i/l2/unified\\_memory\\_oversubscription.cu](https://tinyurl.com/cudafordummies/i/l2/unified_memory_oversubscription.cu)

# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, UNIFIED MEMORY PROGRAMMING CONT.

```
cuda-zen sh@n3073-004:~$ nvcc ./unified_memory_oversubscription.cu  
cuda-zen sh@n3073-004:~$ nsys nvprof ./a.out
```

[5/7] Executing 'gpukernsum' stats report

Time (%)	Total Time (ns)	Instances	Avg (ns)	Med (ns)	Min (ns)	Max (ns)	StdDev (ns)	GridXYZ	BlockXYZ
100.0	6274870199	1	6274870199.0	6274870199.0	6274870199	6274870199	0.0	6291456	1 1 1

All time in  
memory transfer

[6/7] Executing 'gpumemtimesum' stats report

Time (%)	Total Time (ns)	Count	Avg (ns)	Med (ns)	Min (ns)	Max (ns)	StdDev (ns)	Operation
100.0	5909981859	96714	61107.8	6656.0	2495	556067	97099.5	[CUDA Unified Memory memcpy DtoH]

[7/7] Executing 'gpumemsizesum' stats report

Total (MB)	Count	Avg (MB)	Med (MB)	Min (MB)	Max (MB)	StdDev (MB)	Operation
49417.290	96714	0.511	0.061	0.004	2.097	0.780	[CUDA Unified Memory memcpy DtoH]

→ [https://tinyurl.com/cudafordummies/i/l2/unified\\_memory\\_oversubscription.cu](https://tinyurl.com/cudafordummies/i/l2/unified_memory_oversubscription.cu)

# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, UNIFIED MEMORY PROGRAMMING CONT.

```
cuda-zen sh@n3073-004:~$ nvcc ./unified_memory_oversubscription.cu  
cuda-zen sh@n3073-004:~$ nsys nvprof ./a.out
```

[5/7] Executing 'gpukernsum' stats report

Time (%)	Total Time (ns)	Instances	Avg (ns)	Med (ns)	Min (ns)	Max (ns)	StdDev (ns)	GridXYZ	BlockXYZ
100.0	6274870199	1	6274870199.0	6274870199.0	6274870199	6274870199	0.0	6291456	1 1 1

All time in  
memory transfer

[6/7] Executing 'gpumemtimesum' stats report

Time (%)	Total Time (ns)	Count	Avg (ns)	Med (ns)	Min (ns)	Max (ns)	StdDev (ns)	Operation
100.0	5909981859	96714	61107.8	6656.0	2495	556067	97099.5	[CUDA Unified Memory memcpy DtoH]

48 GB reported

[7/7] Executing 'gpumemsizesum' stats report

Total (MB)	Count	Avg (MB)	Med (MB)	Min (MB)	Max (MB)	StdDev (MB)	Operation
49417.290	96714	0.511	0.061	0.004	2.097	0.780	[CUDA Unified Memory memcpy DtoH]

→ [https://tinyurl.com/cudafordummies/i/l2/unified\\_memory\\_oversubscription.cu](https://tinyurl.com/cudafordummies/i/l2/unified_memory_oversubscription.cu)

# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, PROGRAMMING MODEL CONT.

How did this work  
in the days be-  
fore CUDA man-  
aged unified mem-  
ory...



→ <https://docs.nvidia.com/cuda/cuda-c-programming-guide>

# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, DEVICE MEMORY

On pre-Pascal-class GPUs a strict distinction is made between host- and device-memory. Kernels operate out of device-memory, so the CUDA runtime provides functions to allocate, deallocate, and copy data into device-memory as well as transfer data between host-memory and device-memory. A typical sequence of operations is:

1. Declare and allocate arrays in host- and device-memory
2. Initialize host data
3. Transfer data from the host to the device
4. Execute one or more kernels
5. Transfer back results from the device to the host

→ <https://docs.nvidia.com/cuda/cuda-c-programming-guide>

→ <https://devblogs.nvidia.com/easy-introduction-cuda-c-and-c>

# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, DEVICE MEMORY CONT.

```
#define ARRAYDIM 268435456
__global__ void KrnlDmmy(float *x, float *y, float *z)
{
    int i;
    i = (blockIdx.x * blockDim.x) + threadIdx.x;
    z[i] = x[i] + y[i];
    return;
}

int main()
{
    // host memory allocation, hb, hc analogous
    ha = (float *) malloc(ARRAYDIM * sizeof(float));
    // device memory allocation, db, dc analogous
    cudaMalloc(&da, ARRAYDIM * sizeof(float));
    // host to device memory transfer
    cudaMemcpy(da, ha, ARRAYDIM * sizeof(float), cudaMemcpyHostToDevice);
    // kernel invocation
    KrnlDmmy <<< numBlocks, threadsPerBlock >>> (da, db, dc);
    // device to host memory transfer
    cudaMemcpy(hc, dc, ARRAYDIM * sizeof(float), cudaMemcpyDeviceToHost);
    // free device/host memory
    cudaFree(da);
    free(ha);
}
```

# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, DEVICE MEMORY CONT.

```
#define ARRAYDIM 268435456
__global__ void KrnlDmmy(float *x, float *y, float *z)
{
    int i;
    i = (blockIdx.x * blockDim.x) + threadIdx.x;
    z[i] = x[i] + y[i];
    return;
}

int main()
{
    // host memory allocation, hb, hc analogous
    ha = (float *) malloc(ARRAYDIM * sizeof(float));
    // device memory allocation, db, dc analogous
    cudaMalloc(&da, ARRAYDIM * sizeof(float));
    // host to device memory transfer
    cudaMemcpy(da, ha, ARRAYDIM * sizeof(float), cudaMemcpyHostToDevice);
    // kernel invocation
    KrnlDmmy <<< numBlocks, threadsPerBlock >>> (da, db, dc);
    // device to host memory transfer
    cudaMemcpy(hc, dc, ARRAYDIM * sizeof(float), cudaMemcpyDeviceToHost);
    // free device/host memory
    cudaFree(da);
    free(ha);
}
```

Kernel  
memory  
alloc-  
ation &  
set up

# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, DEVICE MEMORY CONT.

```
#define ARRAYDIM 268435456
__global__ void KrnlDmmy(float *x, float *y, float *z)
{
    int i;
    i = (blockIdx.x * blockDim.x) + threadIdx.x;
    z[i] = x[i] + y[i];
    return;
}

int main()
{
    // host memory allocation, hb, hc analogous
    ha = (float *) malloc(ARRAYDIM * sizeof(float));
    // device memory allocation, db, dc analogous
    cudaMalloc(&da, ARRAYDIM * sizeof(float));
    // host to device memory transfer
    cudaMemcpy(da, ha, ARRAYDIM * sizeof(float), cudaMemcpyHostToDevice);
    // kernel invocation
    KrnlDmmy <<< numBlocks, threadsPerBlock >>> (da, db, dc);
    // device to host memory transfer
    cudaMemcpy(hc, dc, ARRAYDIM * sizeof(float), cudaMemcpyDeviceToHost);
    // free device/host memory
    cudaFree(da);
    free(ha);
}
```

Kernel  
memory  
alloc-  
ation &  
set up

Stnd  
kernel  
call

# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, DEVICE MEMORY CONT.

```
#define ARRAYDIM 268435456
__global__ void KrnlDmmmy(float *x, float *y, float *z)
{
    int i;
    i = (blockIdx.x * blockDim.x) + threadIdx.x;
    z[i] = x[i] + y[i];
    return;
}

int main()
{
    // host memory allocation, hb, hc analogous
    ha = (float *) malloc(ARRAYDIM * sizeof(float));
    // device memory allocation, db, dc analogous
    cudaMalloc(&da, ARRAYDIM * sizeof(float));
    // host to device memory transfer
    cudaMemcpy(da, ha, ARRAYDIM * sizeof(float), cudaMemcpyHostToDevice);
    // kernel invocation
    KrnlDmmmy <<< numBlocks, threadsPerBlock >>> (da, db, dc);
    // device to host memory transfer
    cudaMemcpy(hc, dc, ARRAYDIM * sizeof(float), cudaMemcpyDeviceToHost);
    // free device/host memory
    cudaFree(da);
    free(ha);
}
```

Kernel  
memory  
alloc-  
ation &  
set up

Stnd  
kernel  
call

Device  
memory  
back-  
transfer

# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, DEVICE MEMORY CONT.

- Special forms `cudaMallocPitch()` and `cudaMalloc3D()` for 2D and 3D arrays
- Optimized for best performance when accessing via pointers
- Also good for device copies `cudaMemcpy2D()` and `cudaMemcpy3D()`
- Returned pitch (or stride) must be used to access array elements
- Optimally padded to meet alignment requirements
- Additional types of global memory, e.g.  
`__device__ float *devPointer`

# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, SHARED MEMORY

Matrix-Matrix Multiplication (simplest relationships):

$$\left( \begin{array}{cccc} A_{1,1} & A_{1,2} & \cdots & A_{1,N} \\ A_{2,1} & A_{2,2} & \cdots & A_{2,N} \\ \vdots & \vdots & \ddots & \vdots \\ A_{k,1} & A_{k,2} & \cdots & A_{k,N} \\ \vdots & \vdots & \ddots & \vdots \\ A_{N,1} & A_{N,2} & \cdots & A_{N,N} \end{array} \right) \left( \begin{array}{cccc} B_{1,1} & B_{1,2} & \cdots & B_{1,N} \\ B_{2,1} & B_{2,2} & \cdots & B_{2,N} \\ \vdots & \vdots & \ddots & \vdots \\ B_{k,1} & B_{k,2} & \cdots & B_{k,N} \\ \vdots & \vdots & \ddots & \vdots \\ B_{N,1} & B_{N,2} & \cdots & B_{N,N} \end{array} \right) = \left( \begin{array}{cccc} C_{1,1} & C_{1,2} & \cdots & C_{1,N} \\ C_{2,1} & C_{2,2} & \cdots & C_{2,N} \\ \vdots & \vdots & \ddots & \vdots \\ C_{k,1} & C_{k,2} & \cdots & C_{k,N} \\ \vdots & \vdots & \ddots & \vdots \\ C_{N,1} & C_{N,2} & \cdots & C_{N,N} \end{array} \right)$$

**A**                    **B**                    **C**

→ <https://docs.nvidia.com/cuda/cuda-c-programming-guide>

# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, SHARED MEMORY CONT.

### Matrix Matrix Multiplication Version 1

```
#define N 4800
__global__ void KrnlMMM(float **A, float **B, float **C)
{
    int i, j, k;
    float tmpC;
    i = ((blockIdx.x * blockDim.x) + threadIdx.x);
    j = ((blockIdx.y * blockDim.y) + threadIdx.y);
    tmpC = (float) 0;
    for (k=0; k<N; k++) {
        tmpC += A[i][k] * B[k][j];
    }
    C[i][j] = tmpC;
    return;
}

int main()
{
    threadsPerBlock.x = 16;
    threadsPerBlock.y = 16;
    numBlocks.x = N / threadsPerBlock.x;
    numBlocks.y = N / threadsPerBlock.y;
    KrnlMMM <<< numBlocks, threadsPerBlock >>> (A, B, C);
    ...
}
```

# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, SHARED MEMORY CONT.

### Matrix Matrix Multiplication Version 1

```
#define N 4800
__global__ void KrnlMMM(float **A, float **B, float **C)
{
    int i, j, k;
    float tmpC;
    i = ((blockIdx.x * blockDim.x) + threadIdx.x);
    j = ((blockIdx.y * blockDim.y) + threadIdx.y);
    tmpC = (float) 0;
    for (k=0; k<N; k++) {
        tmpC += A[i][k] * B[k][j];
    }
    C[i][j] = tmpC;
    return;
}

int main()
{
    threadsPerBlock.x = 16;
    threadsPerBlock.y = 16;
    numBlocks.x = N / threadsPerBlock.x;
    numBlocks.y = N / threadsPerBlock.y;
    KrnlMMM <<< numBlocks, threadsPerBlock >>> (A, B, C);
    ...
}
```

Each thread computes its specific  $C[i][j]$

# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, SHARED MEMORY CONT.

### Matrix Matrix Multiplication Version 1

Far-from-optimal memory access !

```
#define N 4800
__global__ void KrnlMMM(float **A, float **B, float **C)
{
    int i, j, k;
    float tmpC;
    i = ((blockIdx.x * blockDim.x) + threadIdx.x);
    j = ((blockIdx.y * blockDim.y) + threadIdx.y);
    tmpC = (float) 0;
    for (k=0; k<N; k++) {
        tmpC += A[i][k] * B[k][j];
    }
    C[i][j] = tmpC;
    return;
}

int main()
{
    threadsPerBlock.x = 16;
    threadsPerBlock.y = 16;
    numBlocks.x = N / threadsPerBlock.x;
    numBlocks.y = N / threadsPerBlock.y;
    KrnlMMM <<< numBlocks, threadsPerBlock >>> (A, B, C);
    ...
}
```

Each thread computes its specific  $C[i][j]$

# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, SHARED MEMORY CONT.

```
cuda-zen sh@n3073-004:~$ nvcc ./mmm_example_2.cu
cuda-zen sh@n3073-004:~$ nsys nvprof ./a.out
```

[5/7] Executing 'gpukernsum' stats report

Time (%)	Total Time (ns)	Instances	Avg (ns)	Med (ns)	Min (ns)	Max (ns)	StdDev (ns)	GridXYZ	BlockXYZ	Name
100.0	583635239	1	583635239.0	583635239.0	583635239	583635239	0.0	300	300	1 16 16 1 KrnlMMM(float **, float **

[6/7] Executing 'gpumemtimesum' stats report

Time (%)	Total Time (ns)	Count	Avg (ns)	Med (ns)	Min (ns)	Max (ns)	StdDev (ns)	Operation
51.3	93431969	2432	38417.8	5263.0	1727	363138	85310.1	[CUDA Unified Memory memcpy DtoH]
48.7	88846629	5887	15092.0	7039.0	3166	317537	24377.6	[CUDA Unified Memory memcpy HtoD]


[7/7] Executing 'gpumemsizesum' stats report

Total (MB)	Count	Avg (MB)	Med (MB)	Min (MB)	Max (MB)	StdDev (MB)	Operation
283.116	2432	0.116	0.025	0.004	1.032	0.246	[CUDA Unified Memory memcpy DtoH]
283.116	5887	0.048	0.025	0.004	1.024	0.091	[CUDA Unified Memory memcpy HtoD]

→ [https://tinyurl.com/cudafordummies/i/l2/mmm\\_example\\_2.cu](https://tinyurl.com/cudafordummies/i/l2/mmm_example_2.cu)

# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, SHARED MEMORY CONT.



→ <https://docs.nvidia.com/cuda/cuda-c-programming-guide>

# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, SHARED MEMORY CONT.

### Matrix Matrix Multiplication Version 2

```
#define N 4800
#define BS 16
__global__ void KrnlIMMM(float **A, float **B, float **C)
{
    int i, j, k, Blckk;
    __shared__ float BlckA[BS][BS], BlckB[BS][BS], BlckC[BS][BS];
    float tmpC;
    i = threadIdx.y;
    j = threadIdx.x;
    BlckC[i][j] = (float) 0;
    for (Blckk = 0; Blckk < (N/BS); Blckk++) {
        BlckA[i][j] = A[(blockIdx.y*BS)+i][(Blckk*BS)+j];
        BlckB[i][j] = B[(Blckk*BS)+i][(blockIdx.x*BS)+j];
        __syncthreads();
        tmpC = (float) 0;
        for (k = 0; k<BS; k++) {
            tmpC += BlckA[i][k] * BlckB[k][j];
        }
        BlckC[i][j] += tmpC;
        __syncthreads();
    }
    C[(blockIdx.y*BS+i)][(blockIdx.x*BS)+j] = BlckC[i][j];
    __syncthreads();
    return;
}
```

# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, SHARED MEMORY CONT.

### Matrix Matrix Multiplication Version 2

Declaration  
of shared  
arrays

```
#define N 4800
#define BS 16
__global__ void KrnlIMMM(float **A, float **B, float **C)
{
    int i, j, k, Blckk;
    __shared__ float BlckA[BS][BS], BlckB[BS][BS], BlckC[BS][BS];
    float tmpC;
    i = threadIdx.y;
    j = threadIdx.x;
    BlckC[i][j] = (float) 0;
    for (Blckk = 0; Blckk < (N/BS); Blckk++) {
        BlckA[i][j] = A[(blockIdx.y*BS)+i][(Blckk*BS)+j];
        BlckB[i][j] = B[(Blckk*BS)+i][(blockIdx.x*BS)+j];
        __syncthreads();
        tmpC = (float) 0;
        for (k = 0; k<BS; k++) {
            tmpC += BlckA[i][k] * BlckB[k][j];
        }
        BlckC[i][j] += tmpC;
        __syncthreads();
    }
    C[(blockIdx.y*BS+i)][(blockIdx.x*BS)+j] = BlckC[i][j];
    __syncthreads();
    return;
}
```

# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, SHARED MEMORY CONT.

### Matrix Matrix Multiplication Version 2

Declaration  
of shared  
arrays

Initialization  
by entire  
threadblock

```
#define N 4800
#define BS 16
__global__ void KrnlIMMM(float **A, float **B, float **C)
{
    int i, j, k, Blckk;
    __shared__ float BlckA[BS][BS], BlckB[BS][BS], BlckC[BS][BS];
    float tmpC;
    i = threadIdx.y;
    j = threadIdx.x;
    BlckC[i][j] = (float) 0;
    for (Blckk = 0; Blckk < (N/BS); Blckk++) {
        BlckA[i][j] = A[(blockIdx.y*BS)+i][(Blckk*BS)+j];
        BlckB[i][j] = B[(Blckk*BS)+i][(blockIdx.x*BS)+j];
        __syncthreads();
        tmpC = (float) 0;
        for (k = 0; k<BS; k++) {
            tmpC += BlckA[i][k] * BlckB[k][j];
        }
        BlckC[i][j] += tmpC;
        __syncthreads();
    }
    C[(blockIdx.y*BS+i)][(blockIdx.x*BS)+j] = BlckC[i][j];
    __syncthreads();
    return;
}
```

# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, SHARED MEMORY CONT.

### Matrix Matrix Multiplication Version 2

Declaration  
of shared  
arrays

Initialization  
by entire  
threadblock

```
#define N 4800
#define BS 16
__global__ void KrnlIMMM(float **A, float **B, float **C)
{
    int i, j, k, Blckk;
    __shared__ float BlckA[BS][BS], BlckB[BS][BS], BlckC[BS][BS];
    float tmpC;
    i = threadIdx.y;
    j = threadIdx.x;
    BlckC[i][j] = (float) 0;
    for (Blckk = 0; Blckk < (N/BS); Blckk++) {
        BlckA[i][j] = A[(blockIdx.y*BS)+i][(Blckk*BS)+j];
        BlckB[i][j] = B[(Blckk*BS)+i][(blockIdx.x*BS)+j];
        __syncthreads();
        tmpC = (float) 0;
        for (k = 0; k<BS; k++) {
            tmpC += BlckA[i][k] * BlckB[k][j];
        }
        BlckC[i][j] += tmpC;
        __syncthreads();
    }
    C[(blockIdx.y*BS+i)][(blockIdx.x*BS)+j] = BlckC[i][j];
    __syncthreads();
    return;
}
```

Loop com-  
bining A/B  
blocks in a  
dot-product  
like fashion

# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, SHARED MEMORY CONT.

### Matrix Matrix Multiplication Version 2

Declaration  
of shared  
arrays

Initialization  
by entire  
threadblock

Copy into  
shared  
memory  
arrays

```
#define N 4800
#define BS 16
__global__ void KrnlIMMM(float **A, float **B, float **C)
{
    int i, j, k, Blckk;
    __shared__ float BlckA[BS][BS], BlckB[BS][BS], BlckC[BS][BS];
    float tmpC;
    i = threadIdx.y;
    j = threadIdx.x;
    BlckC[i][j] = (float) 0;
    for (Blckk = 0; Blckk < (N/BS); Blckk++) {
        BlckA[i][j] = A[(blockIdx.y*BS)+i][(Blckk*BS)+j];
        BlckB[i][j] = B[(Blckk*BS)+i][(blockIdx.x*BS)+j];
        __syncthreads();
        tmpC = (float) 0;
        for (k = 0; k<BS; k++) {
            tmpC += BlckA[i][k] * BlckB[k][j];
        }
        BlckC[i][j] += tmpC;
        __syncthreads();
    }
    C[(blockIdx.y*BS+i)][(blockIdx.x*BS)+j] = BlckC[i][j];
    __syncthreads();
    return;
}
```

Loop com-  
bining A/B  
blocks in a  
dot-product  
like fashion

# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, SHARED MEMORY CONT.

### Matrix Matrix Multiplication Version 2

Declaration of shared arrays

Initialization by entire threadblock

Copy into shared memory arrays

```
#define N 4800
#define BS 16
__global__ void KrnlIMMM(float **A, float **B, float **C)
{
    int i, j, k, Blckk;
    __shared__ float BlckA[BS][BS], BlckB[BS][BS], BlckC[BS][BS];
    float tmpC;
    i = threadIdx.y;
    j = threadIdx.x;
    BlckC[i][j] = (float) 0;
    for (Blckk = 0; Blckk < (N/BS); Blckk++) {
        BlckA[i][j] = A[(blockIdx.y*BS)+i][(Blckk*BS)+j];
        BlckB[i][j] = B[(Blckk*BS)+i][(blockIdx.x*BS)+j];
        __syncthreads();
        tmpC = (float) 0;
        for (k = 0; k<BS; k++) {
            tmpC += BlckA[i][k] * BlckB[k][j];
        }
        BlckC[i][j] += tmpC;
        __syncthreads();
    }
    C[(blockIdx.y*BS+i)][(blockIdx.x*BS)+j] = BlckC[i][j];
    __syncthreads();
    return;
}
```

Loop combining A/B blocks in a dot-product like fashion

Thread-wise update of BlckC

# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, SHARED MEMORY CONT.

### Matrix Matrix Multiplication Version 2

Declaration of shared arrays

Initialization by entire threadblock

Copy into shared memory arrays

```
#define N 4800
#define BS 16
__global__ void KrnlIMMM(float **A, float **B, float **C)
{
    int i, j, k, Blckk;
    __shared__ float BlckA[BS][BS], BlckB[BS][BS], BlckC[BS][BS];
    float tmpC;
    i = threadIdx.y;
    j = threadIdx.x;
    BlckC[i][j] = (float) 0;
    for (Blckk = 0; Blckk < (N/BS); Blckk++) {
        BlckA[i][j] = A[(blockIdx.y*BS)+i][(Blckk*BS)+j];
        BlckB[i][j] = B[(Blckk*BS)+i][(blockIdx.x*BS)+j];
        __syncthreads();
        tmpC = (float) 0;
        for (k = 0; k<BS; k++) {
            tmpC += BlckA[i][k] * BlckB[k][j];
        }
        BlckC[i][j] += tmpC;
        __syncthreads();
    }
    C[(blockIdx.y*BS+i)][(blockIdx.x*BS)+j] = BlckC[i][j];
    __syncthreads();
    return;
}
```

Loop combining A/B blocks in a dot-product like fashion

Thread-wise update of BlckC

Back-copy from shared to global memory

# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, SHARED MEMORY CONT.

```
cuda-zen sh@n3073-004:~$ nvcc ./mmm_example_3.cu
cuda-zen sh@n3073-004:~$ nsys nvprof ./a.out
```

[5/7] Executing 'gpukernsum' stats report

Time (%)	Total Time (ns)	Instances	Avg (ns)	Med (ns)	Min (ns)	Max (ns)	StdDev (ns)	GridXY	Name
100.0	203907614	1	203907614.0	203907614.0	203907614	203907614	0.0	300 300 1 16 16 1	Krn1MMMM(float **, float **

Factor  
≈2.9×  
faster

[6/7] Executing 'gpumemtimesum' stats report

Time (%)	Total Time (ns)	Count	Avg (ns)	Med (ns)	Min (ns)	Max (ns)	StdDev (ns)	Operation
54.3	105519535	6791	15538.1	8831.0	3712	337602	28616.3	[CUDA Unified Memory memcpy HtoD]
45.7	88804386	2385	37234.5	4736.0	1535	545476	83237.8	[CUDA Unified Memory memcpy DtoH]

[7/7] Executing 'gpumemsizesum' stats report

Total (MB)	Count	Avg (MB)	Med (MB)	Min (MB)	Max (MB)	StdDev (MB)	Operation
283.116	2385	0.119	0.025	0.004	1.028	0.247	[CUDA Unified Memory memcpy DtoH]
283.116	6791	0.042	0.016	0.004	1.036	0.096	[CUDA Unified Memory memcpy HtoD]

→ [https://tinyurl.com/cudafordummies/i/l2/mmm\\_example\\_3.cu](https://tinyurl.com/cudafordummies/i/l2/mmm_example_3.cu)

# MEMORY HIERARCHY CONT.

## CUDA C-PROGRAMMING GUIDE, LOCAL MEMORY, REGISTERS

- Local memory for certain automatic variables
- Access to local memory is similar to global memory, i.e. high latency and low bandwidth
- Organized such that consecutive 32-bit words are accessed by consecutive thread IDs
- Analyzable with Nsight
- Fastest memory is registers, L1, for small, statically indexed arrays, e.g. A[16]

→ <https://docs.nvidia.com/cuda/cuda-c-programming-guide>

→ <https://stackoverflow.com/questions/10297067/in-a-cuda-kernel-how-do-i-store-an-array-in-local-thread-memory>

## TAKE HOME MESSAGES

- Unified memory — a big improvement

## TAKE HOME MESSAGES

- Unified memory — a big improvement
- Single address space also facilitates easier handling of 2D/3D arrays on the GPU

## TAKE HOME MESSAGES

- Unified memory — a big improvement
- Single address space also facilitates easier handling of 2D/3D arrays on the GPU
- Straightforward profiling with nsys nvprof on the CLI

## TAKE HOME MESSAGES

- Unified memory — a big improvement
- Single address space also facilitates easier handling of 2D/3D arrays on the GPU
- Straightforward profiling with nsys nvprof on the CLI
- Page migration engine facilitates seamless data transfer for array sizes exceeding largely the amount of RAM available on the GPU

## TAKE HOME MESSAGES

- Unified memory — a big improvement
- Single address space also facilitates easier handling of 2D/3D arrays on the GPU
- Straightforward profiling with nsys nvprof on the CLI
- Page migration engine facilitates seamless data transfer for array sizes exceeding largely the amount of RAM available on the GPU
- Using shared memory greatly improves kernel performance