

CUDA SDK — BASIC CONCEPTS

Siegfried Höfinger

ASC Research Center, TU Wien

October 20, 2025

→ <https://tinyurl.com/cudafordummies/ii/l3/notes-l3.pdf>

OUTLINE

BASIC CONVENTIONS

0_INTRODUCTION/SIMPLEPRINTF

1 UTILITIES

0_INTRODUCTION/SIMPLESTREAMS

4_CUDA_LIBRARIES/RANDOMFOG

5_DOMAIN_SPECIFIC/NBODY

4_CUDA_LIBRARIES/OCEANFFT

TAKE HOME MESSAGES

BASIC CONVENTIONS

CUDA SDK

- Nowadays on github (curated, re-structured, toolkit-dependent)

→ <https://github.com/nvidia/cuda-samples>

CUDA 4 DUMMIES — OCT 22-23, 2025

BASIC CONVENTIONS

CUDA SDK

- Nowadays on github (curated, re-structured, toolkit-dependent)
- Constantly growing collection of basic CUDA programming exercises (educational resource, not HPC-focussed)

→ <https://github.com/nvidia/cuda-samples>

BASIC CONVENTIONS

CUDA SDK

- Nowadays on github (curated, re-structured, toolkit-dependent)
- Constantly growing collection of basic CUDA programming exercises (educational resource, not HPC-focussed)
- Showcases all GPU features and explains fundamental CUDA concepts

→ <https://github.com/nvidia/cuda-samples>

BASIC CONVENTIONS

CUDA SDK

- Nowadays on github (curated, re-structured, toolkit-dependent)
- Constantly growing collection of basic CUDA programming exercises (educational resource, not HPC-focussed)
- Showcases all GPU features and explains fundamental CUDA concepts
- Open source and easily expandable with own projects

→ <https://github.com/nvidia/cuda-samples>

BASIC CONVENTIONS

CUDA SDK

- Nowadays on github (curated, re-structured, toolkit-dependent)
- Constantly growing collection of basic CUDA programming exercises (educational resource, not HPC-focussed)
- Showcases all GPU features and explains fundamental CUDA concepts
- Open source and easily expandable with own projects
- Excellent place for beginners to start looking around

→ <https://github.com/nvidia/cuda-samples>

BASIC CONVENTIONS

CUDA SDK

- Nowadays on github (curated, re-structured, toolkit-dependent)
- Constantly growing collection of basic CUDA programming exercises (educational resource, not HPC-focussed)
- Showcases all GPU features and explains fundamental CUDA concepts
- Open source and easily expandable with own projects
- Excellent place for beginners to start looking around
- Thematically organized into 9 major subject areas

→ <https://github.com/nvidia/cuda-samples>

BASIC CONVENTIONS CONT.

CUDA SDK CONT.

```
cuda-zen sh@n3073-004:~$ git clone --branch v12.9 --single-branch https://github.com/NVIDIA/cuda-samples
cuda-zen sh@n3073-004:~$ cd cuda-samples/Samples
cuda-zen sh@n3073-004: Samples$ ls

0_Introduction  2_Concepts_and_Techniques  4_CUDA_Libraries  6_Performance  8_Platform_Specific
1_Utils         3_CUDA_Features           5_Domain_Specific  7_libNVVM

cuda-zen sh@n3073-004: Samples$ cd 0_Introduction
cuda-zen sh@n3073-004: 0_Introduction$ ls

CMakeLists.txt      mergeSort          simpleHyperQ        simpleTemplates
README.md          simpleAWBarrier    simpleIPC          simpleTexture
UnifiedMemoryStreams simpleAssert       simpleLayeredTexture simpleTexture3D
asyncAPI           simpleAssert_nvrtc simpleMPI          simpleTextureDrv
clock               simpleAtomicIntrinsics simpleMultiCopy    simpleVoteIntrinsics
clock_nvrtc         simpleAtomicIntrinsics_nvrtc simpleMultiGPU    simpleZeroCopy
cudaOpenMP          simpleAttributes   simpleOccupancy   systemWideAtomics
fp16ScalarProduct simpleCUDA2GL     simpleP2P          template
matrixMul          simpleCallback     simplePitchLinearTexture vectorAdd
matrixMulDrv        simpleCooperativeGroups simplePrintf      vectorAddDrv
matrixMulDynlinkJIT simpleCubemapTexture simpleStreams    vectorAddMMAP
matrixMul_nvrtc    simpleDrvRuntime   simpleSurfaceWrite vectorAdd_nvrtc
```

→ <https://github.com/nvidia/cuda-samples>

BASIC CONVENTIONS CONT.

CUDA SDK CONT.

```
cuda-zen sh@n3073-004:~$ git clone --branch v12.9 --single-branch https://github.com/NVIDIA/cuda-samples
cuda-zen sh@n3073-004:~$ cd cuda-samples/Samples
cuda-zen sh@n3073-004: Samples$ ls

0_Introduction  2_Concepts_and_Techniques  4_CUDA_Libraries  6_Performance  8_Platform_Specific
1_Utils         3_CUDA_Features           5_Domain_Specific  7_libNVVM

cuda-zen sh@n3073-004: Samples$ cd 0_Introduction
cuda-zen sh@n3073-004: 0_Introduction$ ls

CMakeLists.txt      mergeSort          simpleHyperQ        simpleTemplates
README.md          simpleAWBarrier    simpleIPC          simpleTexture
UnifiedMemoryStreams simpleAssert       simpleLayeredTexture simpleTexture3D
asyncAPI           simpleAssert_nvrtc simpleMPI          simpleTextureDrv
clock               simpleAtomicIntrinsics simpleMultiCopy    simpleVoteIntrinsics
clock_nvrtc         simpleAtomicIntrinsics_nvrtc simpleMultiGPU    simpleZeroCopy
cudaOpenMP          simpleAttributes   simpleOccupancy   systemWideAtomics
fp16ScalarProduct simpleCUDA2GL     simpleP2P          template
matrixMul          simpleCallback     simplePitchLinearTexture vectorAdd
matrixMulDrv        simpleCooperativeGroups simplePrintf      vectorAddDrv
matrixMulDynlinkJIT simpleCubemapTexture simpleStreams    vectorAddMMAP
matrixMul_nvrtc    simpleDrvRuntime   simpleSurfaceWrite vectorAdd_nvrtc
```

→ <https://github.com/nvidia/cuda-samples>

BASIC CONVENTIONS CONT.

CUDA SDK CONT.

<u>0_Introduction</u>	CUDA samples for beginners that illustrate key concepts with using CUDA and CUDA runtime APIs
<u>1_Utils</u>	Utility samples that demonstrate how to query device capabilities and measure GPU/CPU bandwidth
<u>2_Concepts_and_Techniques</u>	CUDA related concepts and common problem solving techniques
<u>3_CUDA_Features</u>	Samples that demonstrate CUDA Features (cooperative groups, dynamic parallelism, graphs etc)
<u>4_CUDA_Libraries</u>	Samples that illustrate how to use CUDA platform libraries, NPP, NVJPEG, NVGRAPH cuBLAS, cuFFT, cuSPARSE, cuSOLVER and cuRAND
<u>5_Domain_Specific</u>	Samples from specific domains (graphics, finance, image processing)
<u>6_Performance</u>	Samples that illustrate performance optimizations
<u>7_libNVVM</u>	Samples of interfacing with NVVM IR (intermediate representations)
<u>8_Platform_Specific</u>	Samples of using CUDA for TEGRA architectures

→ <https://github.com/nvidia/cuda-samples>

BASIC CONVENTIONS CONT.

CUDA SDK CONT.

Consider for example function assert()

```
#include <stdio.h>
#include <assert.h>

int main()
{
    int i;

    for (i=0; i<10; i++) {
        assert(i < 5);
        printf("current i is %d \n", i);
    }

    return(0);
}
```

→ https://tinyurl.com/cudafordummies/ii/l3/smpl_assert.c

BASIC CONVENTIONS CONT.

CUDA SDK CONT.

Consider for example function assert()

```
#include <stdio.h>
#include <assert.h>

int main()
{
    int i;

    for (i=0; i<10; i++) {
        assert(i < 5);
        printf("current i is %d \n", i);
    }

    return(0);
}
```

```
cuda-zen sh@n3073-004:~$ gcc ./smpl_assert.c
cuda-zen sh@n3073-004:~$ ./a.out
current i is 0
current i is 1
current i is 2
current i is 3
current i is 4
a.out: smpl_assert.c:20: main: Assertion `i < 5' failed.
Aborted (core dumped)
```

→ https://tinyurl.com/cudafordummies/ii/l3/smpl_assert.c

BASIC CONVENTIONS CONT.

CUDA SDK CONT.

Consider for example function assert()

```
#include <stdio.h>
#include <assert.h>

int main()
{
    int i;

    for (i=0; i<10; i++) {
        assert(i < 5);
        printf("current i is %d \n", i);
    }

    return(0);
}
```

```
cuda-zen sh@n3073-004:~$ gcc ./smpl_assert.c
cuda-zen sh@n3073-004:~$ ./a.out
current i is 0
current i is 1
current i is 2
current i is 3
current i is 4
a.out: smpl_assert.c:20: main: Assertion `i < 5' failed.
Aborted (core dumped)
```

Developer's checkpoints: If expression is TRUE assert() does nothing. If FALSE, abortion and error message.

→ https://tinyurl.com/cudafordummies/ii/l3/smpl_assert.c

BASIC CONVENTIONS CONT.

CUDA SDK CONT.

There is also
a CUDA
SDK sample
in 0_Introduc-
tion/simpleAssert/

Consider for example function assert()

```
#include <stdio.h>
#include <assert.h>

int main()
{
    int i;

    for (i=0; i<10; i++) {
        assert(i < 5);
        printf("current i is %d \n", i);
    }

    return(0);
}
```

Developer's checkpoints: If expression is TRUE assert() does nothing. If FALSE, abortion and error message.

```
cuda-zen sh@n3073-004:~$ gcc ./smpl_assert.c
cuda-zen sh@n3073-004:~$ ./a.out
current i is 0
current i is 1
current i is 2
current i is 3
current i is 4
a.out: smpl_assert.c:20: main: Assertion `i < 5' failed.
Aborted (core dumped)
```

→ https://tinyurl.com/cudafordummies/ii/l3/smpl_assert.c

BASIC CONVENTIONS CONT.

CUDA SDK CONT. — 0_INTRODUCTION/SIMPLEASSERT

```
cuda-zen sh@n3073-004:~$ cd cuda-samples/Samples/0_Introduction/simpleAssert
cuda-zen sh@n3073-004: simpleAssert$ ls
CMakeLists.txt README.md simpleAssert.cu

cuda-zen sh@n3073-004: simpleAssert$ cd ..
cuda-zen sh@n3073-004: 0_Introduction$ cp -r ./simpleAssert ./my_simpleAssert
cuda-zen sh@n3073-004: 0_Introduction$ cd ./my_simpleAssert
cuda-zen sh@n3073-004: my_simpleAssert$ module purge
cuda-zen sh@n3073-004: my_simpleAssert$ module load gcc/12.2.0-gcc-8.5.0-7ih4ahb
cuda-zen sh@n3073-004: my_simpleAssert$ module load cuda/12.9.0-gcc-12.2.0-xznhx43
cuda-zen sh@n3073-004: my_simpleAssert$ module load cmake/3.29.1-gcc-12.2.0-yooqtj6
cuda-zen sh@n3073-004: my_simpleAssert$ mkdir build
cuda-zen sh@n3073-004: my_simpleAssert$ cd build
cuda-zen sh@n3073-004: build$ cmake ..
cuda-zen sh@n3073-004: build$ make
```

→ <https://github.com/nvidia/cuda-samples>

BASIC CONVENTIONS CONT.

CUDA SDK CONT. — 0_INTRODUCTION/SIMPLEASSERT CONT.

```
cuda-zen sh@n3073-004: build$ ./simpleAssert
simpleAssert starting...
OS_System_Type.release = 4.18.0-477.10.1.el8_8.x86_64
OS Info: <#1 SMP Tue May 16 07:35:04 EDT 2023>

GPU Device 0: "Ampere" with compute capability 8.0

Launch kernel to generate assertion failures

-- Begin assert output

simpleAssert.cu:63: void testKernel(int): block: [1,0,0], thread: [28,0,0] Assertion 'gtid < N' failed.
simpleAssert.cu:63: void testKernel(int): block: [1,0,0], thread: [29,0,0] Assertion 'gtid < N' failed.
simpleAssert.cu:63: void testKernel(int): block: [1,0,0], thread: [30,0,0] Assertion 'gtid < N' failed.
simpleAssert.cu:63: void testKernel(int): block: [1,0,0], thread: [31,0,0] Assertion 'gtid < N' failed.

-- End assert output

Device assert failed as expected, CUDA error message is: device-side assert triggered

simpleAssert completed, returned OK
```

→ <https://github.com/nvidia/cuda-samples>

BASIC CONVENTIONS CONT.

CUDA SDK CONT. — 0_INTRODUCTION/SIMPLEASSERT CONT.

simpleAssert.cu

```
__global__ void testKernel(int N)
{
    int gtid = blockIdx.x*blockDim.x + threadIdx.x;
    assert(gtid < N);
}
int main(int argc, char **argv)
{
    ...
    runTest(argc, argv);
    ...
}
void runTest(int argc, char **argv)
{
    int Nblocks = 2;
    int Nthreads = 32;
    cudaError_t error;
    ...
    findCudaDevice(argc, (const char **)argv);
    dim3 dimGrid(Nblocks);
    dim3 dimBlock(Nthreads);
    testKernel <<< dimGrid, dimBlock >>> (60);
    error = cudaDeviceSynchronize();
}
```

→ https://docs.nvidia.com/cuda/cuda-runtime-api/group__CUDART__TYPES.html

BASIC CONVENTIONS CONT.

CUDA SDK CONT. — 0_INTRODUCTION/SIMPLEASSERT CONT.

simpleAssert.cu

```
__global__ void testKernel(int N)
{
    int gtid = blockIdx.x*blockDim.x + threadIdx.x;
    assert(gtid < N);
}
int main(int argc, char **argv)
{
    ...
    runTest(argc, argv);
    ...
}
void runTest(int argc, char **argv)
{
    int Nblocks = 2;
    int Nthreads = 32;
    cudaError_t error;
    ...
    findCudaDevice(argc, (const char **)argv);
    dim3 dimGrid(Nblocks);
    dim3 dimBlock(Nthreads);
    testKernel <<< dimGrid, dimBlock >>> (60);
    error = cudaDeviceSynchronize();
}
```

2 threadblocks
with 32 threads

→ https://docs.nvidia.com/cuda/cuda-runtime-api/group__CUDART__TYPES.html

BASIC CONVENTIONS CONT.

CUDA SDK CONT. — 0_INTRODUCTION/SIMPLEASSERT CONT.

simpleAssert.cu

```
__global__ void testKernel(int N)
{
    int gtid = blockIdx.x*blockDim.x + threadIdx.x;
    assert(gtid < N);
}
int main(int argc, char **argv)
{
    ...
    runTest(argc, argv);
    ...
}
void runTest(int argc, char **argv)
{
    int Nblocks = 2;
    int Nthreads = 32;
    cudaError_t error;
    ...
    findCudaDevice(argc, (const char **)argv);
    dim3 dimGrid(Nblocks);
    dim3 dimBlock(Nthreads);
    testKernel <<< dimGrid, dimBlock >>> (60);
    error = cudaDeviceSynchronize();
}
```

New type (enum)
for error variable

2 threadblocks
with 32 threads

→ https://docs.nvidia.com/cuda/cuda-runtime-api/group__CUDART__TYPES.html

BASIC CONVENTIONS CONT.

CUDA SDK CONT. — 0_INTRODUCTION/SIMPLEASSERT CONT.

simpleAssert.cu

```
__global__ void testKernel(int N)
{
    int gtid = blockIdx.x*blockDim.x + threadIdx.x;
    assert(gtid < N);
}
int main(int argc, char **argv)
{
    ...
    runTest(argc, argv);
    ...
}
void runTest(int argc, char **argv)
{
    int Nblocks = 2;
    int Nthreads = 32;
    cudaError_t error;
    ...
    findCudaDevice(argc, (const char **)argv);
    dim3 dimGrid(Nblocks);
    dim3 dimBlock(Nthreads);
    testKernel <<< dimGrid, dimBlock >>> (60);
    error = cudaDeviceSynchronize();
}
```

New type (enum)
for error variable

Generic function
to identify GPU

2 threadblocks
with 32 threads

→ https://docs.nvidia.com/cuda/cuda-runtime-api/group__CUDART__TYPES.html

BASIC CONVENTIONS CONT.

CUDA SDK CONT. — 0_INTRODUCTION/SIMPLEASSERT CONT.

simpleAssert.cu

```
__global__ void testKernel(int N)
{
    int gtid = blockIdx.x*blockDim.x + threadIdx.x;
    assert(gtid < N);
}
int main(int argc, char **argv)
{
    ...
    runTest(argc, argv);
    ...
}
void runTest(int argc, char **argv)
{
    int Nblocks = 2;
    int Nthreads = 32;
    cudaError_t error;
    ...
    findCudaDevice(argc, (const char **)argv);
    dim3 dimGrid(Nblocks);
    dim3 dimBlock(Nthreads);
    testKernel <<< dimGrid, dimBlock >>> (60);
    error = cudaDeviceSynchronize();
}
```

New type (enum)
for error variable

Generic function
to identify GPU

2 threadblocks
with 32 threads

C++ style of initializing variables
dimGrid and dim-
Block

→ https://docs.nvidia.com/cuda/cuda-runtime-api/group__CUDART__TYPES.html

BASIC CONVENTIONS CONT.

CUDA SDK CONT. — 0_INTRODUCTION/SIMPLEASSERT CONT.

simpleAssert.cu

```
__global__ void testKernel(int N)
{
    int gtid = blockIdx.x*blockDim.x + threadIdx.x;
    assert(gtid < N);
}
int main(int argc, char **argv)
{
    ...
    runTest(argc, argv);
    ...
}
void runTest(int argc, char **argv)
{
    int Nblocks = 2;
    int Nthreads = 32;
    cudaError_t error;
    ...
    findCudaDevice(argc, (const char **)argv);
    dim3 dimGrid(Nblocks);
    dim3 dimBlock(Nthreads);
    testKernel <<< dimGrid, dimBlock >>> (60);
    error = cudaDeviceSynchronize();
}
```

New type (enum)
for error variable

Generic function
to identify GPU

2 threadblocks
with 32 threads

C++ style of initializing variables
dimGrid and dim-
Block

kernel launch

→ https://docs.nvidia.com/cuda/cuda-runtime-api/group__CUDART__TYPES.html

BASIC CONVENTIONS CONT.

CUDA SDK CONT. — 0_INTRODUCTION/SIMPLEASSERT CONT.

simpleAssert.cu

```
__global__ void testKernel(int N)
{
    int gtid = blockIdx.x*blockDim.x + threadIdx.x;
    assert(gtid < N);
}
int main(int argc, char **argv)
{
    ...
    runTest(argc, argv);
    ...
}
void runTest(int argc, char **argv)
{
    int Nblocks = 2;
    int Nthreads = 32;
    cudaError_t error;
    ...
    findCudaDevice(argc, (const char **)argv);
    dim3 dimGrid(Nblocks);
    dim3 dimBlock(Nthreads);
    testKernel <<< dimGrid, dimBlock >>> (60);
    error = cudaDeviceSynchronize();
}
```

New type (enum)
for error variable

Generic function
to identify GPU

Check is whether
gtid is < 60

2 threadblocks
with 32 threads

C++ style of initializing variables
dimGrid and dim-
Block

kernel launch

→ https://docs.nvidia.com/cuda/cuda-runtime-api/group__CUDART__TYPES.html

BASIC CONVENTIONS CONT.

CUDA SDK CONT. — 0_INTRODUCTION/SIMPLEASSERT CONT.

simpleAssert.cu

Device behaviour identical to host;
lots of useful information !

New type (enum)
for error variable

Generic function
to identify GPU

```
__global__ void testKernel(int N)
{
    int gtid = blockIdx.x*blockDim.x + threadIdx.x;
    assert(gtid < N);
}
int main(int argc, char **argv)
{
    ...
    runTest(argc, argv);
    ...
}
void runTest(int argc, char **argv)
{
    int Nblocks = 2;
    int Nthreads = 32;
    cudaError_t error;
    ...
    findCudaDevice(argc, (const char **)argv);
    dim3 dimGrid(Nblocks);
    dim3 dimBlock(Nthreads);
    testKernel <<< dimGrid, dimBlock >>> (60);
    error = cudaDeviceSynchronize();
}
```

Check is whether
gtid is < 60

2 threadblocks
with 32 threads

C++ style of initializing variables
dimGrid and dim-
Block

kernel launch

→ https://docs.nvidia.com/cuda/cuda-runtime-api/group__CUDART__TYPES.html

BASIC CONVENTIONS CONT.

CUDA SDK CONT. — 0_INTRODUCTION/SIMPLEASSERT CONT.

simpleAssert.cu / error handling

```
void runTest(int argc, char **argv)
{
    ...
    printf("Launch kernel to generate assertion failures\n");
    testKernel <<< dimGrid, dimBlock >>> (60);

    //Synchronize (flushes assert output)
    printf("\n- Begin assert output\n\n");
    error = cudaDeviceSynchronize();
    printf("\n- End assert output\n\n");

    //Check for errors
    if (error == cudaErrorAssert) {
        printf("Device assert failed as expected, "
               "CUDA error message is: %s\n\n",
               cudaGetErrorString(error));
    }
    testResult = error == cudaErrorAssert;
}
```

→ https://docs.nvidia.com/cuda/cuda-runtime-api/group__CUDART__TYPES.html

BASIC CONVENTIONS CONT.

CUDA SDK CONT. — 0_INTRODUCTION/SIMPLEASSERT CONT.

simpleAssert.cu / error handling

```
void runTest(int argc, char **argv)
{
    ...
    printf("Launch kernel to generate assertion failures\n");
    testKernel <<< dimGrid, dimBlock >>> (60);

    //Synchronize (flushes assert output)
    printf("\n- Begin assert output\n\n");
    error = cudaDeviceSynchronize(); ←
    printf("\n- End assert output\n\n");

    //Check for errors
    if (error == cudaErrorAssert) {
        printf("Device assert failed as expected, "
               "CUDA error message is: %s\n\n",
               cudaGetErrorString(error));
    }
    testResult = error == cudaErrorAssert;
}
```

Pick up re-
turn value from
CUDA call

→ https://docs.nvidia.com/cuda/cuda-runtime-api/group__CUDART__TYPES.html

BASIC CONVENTIONS CONT.

CUDA SDK CONT. — 0_INTRODUCTION/SIMPLEASSERT CONT.

simpleAssert.cu / error handling

```
void runTest(int argc, char **argv)
{
    ...
    printf("Launch kernel to generate assertion failures\n");
    testKernel <<< dimGrid, dimBlock >>> (60);

    //Synchronize (flushes assert output)
    printf("\n- Begin assert output\n\n");
    error = cudaDeviceSynchronize(); ←
    printf("\n- End assert output\n\n");

    //Check for errors
    if (error == cudaErrorAssert) {
        printf("Device assert failed as expected, "
               "CUDA error message is: %s\n\n",
               cudaGetErrorString(error)); ←
    }
    testResult = error == cudaErrorAssert;
}
```

Pick up re-
turn value from
CUDA call

Get specific
info from re-
turned error

→ https://docs.nvidia.com/cuda/cuda-runtime-api/group__CUDART__TYPES.html

BASIC CONVENTIONS CONT.

CUDA SDK CONT. — 0_INTRODUCTION/SIMPLEASSERT CONT.

simpleAssert.cu / error handling

```
void runTest(int argc, char **argv)
{
    ...
    printf("Launch kernel to generate assertion failures\n");
    testKernel <<< dimGrid, dimBlock >>> (60);

    //Synchronize (flushes assert output)
    printf("\n- Begin assert output\n\n");
    error = cudaDeviceSynchronize(); ←
    printf("\n- End assert output\n\n");

    //Check for errors
    if (error == cudaErrorAssert) {
        printf("Device assert failed as expected, "
               "CUDA error message is: %s\n\n",
               cudaGetErrorString(error)); ←
    }
    testResult = error == cudaErrorAssert;
}
```

Set global variable to true/false

Pick up return value from CUDA call

Get specific info from returned error

→ https://docs.nvidia.com/cuda/cuda-runtime-api/group__CUDART__TYPES.html

BASIC CONVENTIONS CONT.

CUDA SDK CONT. — 0_INTRODUCTION/SIMPLEASSERT CONT.

- assert() is a very simple and convenient way to do low-level debugging of kernel code

→ <https://github.com/nvidia/cuda-samples>

BASIC CONVENTIONS CONT.

CUDA SDK CONT. — 0_INTRODUCTION/SIMPLEASSERT CONT.

- assert() is a very simple and convenient way to do low-level debugging of kernel code
- Returns very detailed information, threadIdx, blockIdx, line number, function name

→ <https://github.com/nvidia/cuda-samples>

BASIC CONVENTIONS CONT.

CUDA SDK CONT. — 0_INTRODUCTION/SIMPLEASSERT CONT.

- assert() is a very simple and convenient way to do low-level debugging of kernel code
- Returns very detailed information, threadIdx, blockIdx, line number, function name
- With printf() — surprisingly — we do also get output written from kernel code sections, however only at full block level terminating correctly

→ <https://github.com/nvidia/cuda-samples>

BASIC CONVENTIONS CONT.

CUDA SDK CONT. — 0_INTRODUCTION/SIMPLEASSERT CONT.

- assert() is a very simple and convenient way to do low-level debugging of kernel code
- Returns very detailed information, threadIdx, blockIdx, line number, function name
- With printf() — surprisingly — we do also get output written from kernel code sections, however only at full block level terminating correctly
- For example add another line after `assert(gtid < N);`
`printf("*** message from thread %d ***\n", gtid);`

→ <https://github.com/nvidia/cuda-samples>

BASIC CONVENTIONS CONT.

CUDA SDK CONT. — 0_INTRODUCTION/SIMPLEASSERT CONT.

```
cuda-zen sh@n3073-004: build$ ./simpleAssert
simpleAssert starting...
OS_System_Type.release = 4.18.0-477.10.1.el8_8.x86_64
OS Info: <#1 SMP Tue May 16 07:35:04 EDT 2023>

GPU Device 0: "Ampere" with compute capability 8.0

Launch kernel to generate assertion failures

-- Begin assert output

*** message from thread 0 ***
*** message from thread 1 ***
*** message from thread 2 ***
.....
*** message from thread 29 ***
*** message from thread 30 ***
*** message from thread 31 ***
simpleAssert.cu:63: void testKernel(int): block: [1,0,0], thread: [28,0,0] Assertion 'gtid < N' failed.
simpleAssert.cu:63: void testKernel(int): block: [1,0,0], thread: [29,0,0] Assertion 'gtid < N' failed.
simpleAssert.cu:63: void testKernel(int): block: [1,0,0], thread: [30,0,0] Assertion 'gtid < N' failed.
simpleAssert.cu:63: void testKernel(int): block: [1,0,0], thread: [31,0,0] Assertion 'gtid < N' failed.

-- End assert output
```

→ <https://github.com/nvidia/cuda-samples>

0__INTRODUCTION/SIMPLEPRINTF

CUDA SDK CONT.

- There is also another simple CUDA example demonstrating regular operation of `printf()` in kernel code sections running on the device
- Compute capability must be at least 2.0
- Otherwise an alternative `cuPrintf()` can be used
- This example is also a good exercise to recall basic builtin variables of the kernel code section, e.g. `threadIdx`, `blockDim` etc.

0_INTRODUCTION/SIMPLEPRINTF CONT.

CUDA SDK CONT.

```
cuda-zen sh@n3073-004: 0_Introduction$ cp -r ./simplePrintf ./my_simplePrintf
cuda-zen sh@n3073-004: 0_Introduction$ cd ./my_simplePrintf
cuda-zen sh@n3073-004: my_simplePrintf$ mkdir build
cuda-zen sh@n3073-004: my_simplePrintf$ cd build
cuda-zen sh@n3073-004: build$ cmake ..
cuda-zen sh@n3073-004: build$ make
cuda-zen sh@n3073-004: build$ ./simplePrintf
GPU Device 0: "Ampere" with compute capability 8.0
```

```
Device 0: "NVIDIA A100-PCIE-40GB" with Compute 8.0 capability
printf() is called. Output:
```

```
[1, 0]:      Value is:10      [0, 0]:      Value is:10
[1, 1]:      Value is:10      [0, 1]:      Value is:10
[1, 2]:      Value is:10      [0, 2]:      Value is:10
[1, 3]:      Value is:10      [0, 3]:      Value is:10
[1, 4]:      Value is:10      [0, 4]:      Value is:10
[1, 5]:      Value is:10      [0, 5]:      Value is:10
[1, 6]:      Value is:10      [0, 6]:      Value is:10
[1, 7]:      Value is:10      [0, 7]:      Value is:10
[2, 0]:      Value is:10      [3, 0]:      Value is:10
[2, 1]:      Value is:10      [3, 1]:      Value is:10
[2, 2]:      Value is:10      [3, 2]:      Value is:10
[2, 3]:      Value is:10      [3, 3]:      Value is:10
[2, 4]:      Value is:10      [3, 4]:      Value is:10
[2, 5]:      Value is:10      [3, 5]:      Value is:10
[2, 6]:      Value is:10      [3, 6]:      Value is:10
[2, 7]:      Value is:10      [3, 7]:      Value is:10
```

0__INTRODUCTION/SIMPLEPRINTF CONT.

CUDA SDK CONT.

```
cuda-zen sh@n3073-004: 0_Introduction$ cp -r ./simplePrintf ./my_simplePrintf
cuda-zen sh@n3073-004: 0_Introduction$ cd ./my_simplePrintf
cuda-zen sh@n3073-004: my_simplePrintf$ mkdir build
cuda-zen sh@n3073-004: my_simplePrintf$ cd build
cuda-zen sh@n3073-004: build$ cmake ..
cuda-zen sh@n3073-004: build$ make
cuda-zen sh@n3073-004: build$ ./simplePrintf
GPU Device 0: "Ampere" with compute capability 8.0

Device 0: "NVIDIA A100-PCIE-40GB" with Compute 8.0 capability
printf() is called. Output:
```

[1, 0]:	Value is:10	[0, 0]:	Value is:10
[1, 1]:	Value is:10	[0, 1]:	Value is:10
[1, 2]:	Value is:10	[0, 2]:	Value is:10
[1, 3]:	Value is:10	[0, 3]:	Value is:10
[1, 4]:	Value is:10	[0, 4]:	Value is:10
[1, 5]:	Value is:10	[0, 5]:	Value is:10
[1, 6]:	Value is:10	[0, 6]:	Value is:10
[1, 7]:	Value is:10	[0, 7]:	Value is:10
[2, 0]:	Value is:10	[3, 0]:	Value is:10
[2, 1]:	Value is:10	[3, 1]:	Value is:10
[2, 2]:	Value is:10	[3, 2]:	Value is:10
[2, 3]:	Value is:10	[3, 3]:	Value is:10
[2, 4]:	Value is:10	[3, 4]:	Value is:10
[2, 5]:	Value is:10	[3, 5]:	Value is:10
[2, 6]:	Value is:10	[3, 6]:	Value is:10
[2, 7]:	Value is:10	[3, 7]:	Value is:10

4 x blocks with
indices from
0-7



0__INTRODUCTION/SIMPLEPRINTF CONT.

CUDA SDK CONT.

```
cuda-zen sh@n3073-004: 0_Introduction$ cp -r ./simplePrintf ./my_simplePrintf
cuda-zen sh@n3073-004: 0_Introduction$ cd ./my_simplePrintf
cuda-zen sh@n3073-004: my_simplePrintf$ mkdir build
cuda-zen sh@n3073-004: my_simplePrintf$ cd build
cuda-zen sh@n3073-004: build$ cmake ..
cuda-zen sh@n3073-004: build$ make
cuda-zen sh@n3073-004: build$ ./simplePrintf
GPU Device 0: "Ampere" with compute capability 8.0

Device 0: "NVIDIA A100-PCIE-40GB" with Compute 8.0 capability
printf() is called. Output:
```

[1, 0]:	Value is:10	[0, 0]:	Value is:10
[1, 1]:	Value is:10	[0, 1]:	Value is:10
[1, 2]:	Value is:10	[0, 2]:	Value is:10
[1, 3]:	Value is:10	[0, 3]:	Value is:10
[1, 4]:	Value is:10	[0, 4]:	Value is:10
[1, 5]:	Value is:10	[0, 5]:	Value is:10
[1, 6]:	Value is:10	[0, 6]:	Value is:10
[1, 7]:	Value is:10	[0, 7]:	Value is:10
[2, 0]:	Value is:10	[3, 0]:	Value is:10
[2, 1]:	Value is:10	[3, 1]:	Value is:10
[2, 2]:	Value is:10	[3, 2]:	Value is:10
[2, 3]:	Value is:10	[3, 3]:	Value is:10
[2, 4]:	Value is:10	[3, 4]:	Value is:10
[2, 5]:	Value is:10	[3, 5]:	Value is:10
[2, 6]:	Value is:10	[3, 6]:	Value is:10
[2, 7]:	Value is:10	[3, 7]:	Value is:10

4 x blocks with
indices from
0-7

stochastic or-
der of blocks
(1st index)

0__INTRODUCTION/SIMPLEPRINTF CONT.

CUDA SDK CONT.

simplePrintf.cu

```
--global__ void testKernel(int val)
{
    int gtid = blockIdx.x*blockDim.x + threadIdx.x;
    printf("[%d, %d]:\t\tValue is:%d\n",
           blockIdx.y*gridDim.x+blockIdx.x,
           threadIdx.z*blockDim.x*blockDim.y+threadIdx.y*blockDim.x+threadIdx.x,
           val);
}

int main(int argc, char **argv)
{
    ...
    dim3 dimGrid(2, 2);
    dim3 dimBlock(2, 2, 2);
    testKernel <<< dimGrid, dimBlock >>> (10);
    error = cudaDeviceSynchronize();
    ...
}
```

→ <https://github.com/nvidia/cuda-samples>

0__INTRODUCTION/SIMPLEPRINTF CONT.

CUDA SDK CONT.

simplePrintf.cu

```
--global__ void testKernel(int val)
{
    int gtid = blockIdx.x*blockDim.x + threadIdx.x;
    printf("[%d, %d]:\t\tValue is:%d\n",
           blockIdx.y*gridDim.x+blockIdx.x,
           threadIdx.z*blockDim.x*blockDim.y+threadIdx.y*blockDim.x+threadIdx.x,
           val);
}

int main(int argc, char **argv)
{
    ...
    dim3 dimGrid(2, 2);
    dim3 dimBlock(2, 2, 2);
    testKernel <<< dimGrid, dimBlock >>> (10);
    error = cudaDeviceSynchronize();
    ...
}
```

4 threadblocks
(2D-grid) with
8 threads (3D-block)

→ <https://github.com/nvidia/cuda-samples>

0__INTRODUCTION/SIMPLEPRINTF CONT.

CUDA SDK CONT.

simplePrintf.cu

```
--global__ void testKernel(int val)
{
    int gtid = blockIdx.x*blockDim.x + threadIdx.x;
    printf("[%d, %d]:\t\tValue is:%d\n",
        blockIdx.y*gridDim.x+blockIdx.x, ←
        threadIdx.z*blockDim.x*blockDim.y+threadIdx.y*blockDim.x+threadIdx.x,
        val);
}

int main(int argc, char **argv)
{
    ...
    dim3 dimGrid(2, 2); ←
    dim3 dimBlock(2, 2, 2);
    testKernel <<< dimGrid, dimBlock >>> (10);
    error = cudaDeviceSynchronize();
    ...
}
```

Linearizes 2D-grid

4 threadblocks
(2D-grid) with
8 threads (3D-block)

→ <https://github.com/nvidia/cuda-samples>

0__INTRODUCTION/SIMPLEPRINTF CONT.

CUDA SDK CONT.

simplePrintf.cu

```
--global__ void testKernel(int val)
{
    int gtid = blockIdx.x*blockDim.x + threadIdx.x;
    printf("[%d, %d]:\t\tValue is:%d\n",
        blockIdx.y*gridDim.x+blockIdx.x, ←
        threadIdx.z*blockDim.x*blockDim.y+threadIdx.y*blockDim.x+threadIdx.x, ←
        val);
}

int main(int argc, char **argv)
{
    ...
    dim3 dimGrid(2, 2); ←
    dim3 dimBlock(2, 2, 2); ←
    testKernel <<< dimGrid, dimBlock >>> (10);
    error = cudaDeviceSynchronize();
    ...
}
```

Linearizes 2D-grid

Again, serial-
ization of 3D-
threadblocks

4 threadblocks
(2D-grid) with
8 threads (3D-
block)

→ <https://github.com/nvidia/cuda-samples>

0__INTRODUCTION/SIMPLEPRINTF CONT.

CUDA SDK CONT.

- Good to see that printf() can be used also in kernel code
- Out of order execution of individual threadblocks in the blockgrid
- Only when the entire threadblock terminates correctly, printf() output will actually show up
- Important for the developmental stage, probably too expensive for production-ready runs

→ <https://github.com/nvidia/cuda-samples>

1 UTILITIES

CUDA SDK CONT.

- 1_Utils contains several CUDA examples that may also be regarded as simple tools to characterize the GPU hardware at hand

→ <https://github.com/nvidia/cuda-samples>

1 UTILITIES

CUDA SDK CONT.

- 1_Utils contains several CUDA examples that may also be regarded as simple tools to characterize the GPU hardware at hand
- deviceQuery has already been presented as a useful standard tool to list all device properties

→ <https://github.com/nvidia/cuda-samples>

1 UTILITIES

CUDA SDK CONT.

- 1_Utils contains several CUDA examples that may also be regarded as simple tools to characterize the GPU hardware at hand
- deviceQuery has already been presented as a useful standard tool to list all device properties
- A recurring issue with CUDA is bandwidth of data transfer — especially because there are so many different variants

→ <https://github.com/nvidia/cuda-samples>

1 UTILITIES

CUDA SDK CONT.

- 1_Utils contains several CUDA examples that may also be regarded as simple tools to characterize the GPU hardware at hand
- deviceQuery has already been presented as a useful standard tool to list all device properties
- A recurring issue with CUDA is bandwidth of data transfer — especially because there are so many different variants
- The bandwidthTest example is nowadays provided as a demo — it helps to get a quick overview of what bandwidth we can expect on the current device

→ <https://github.com/nvidia/cuda-samples>

1 UTILITIES

CUDA SDK CONT.

- 1_Utils contains several CUDA examples that may also be regarded as simple tools to characterize the GPU hardware at hand
- deviceQuery has already been presented as a useful standard tool to list all device properties
- A recurring issue with CUDA is bandwidth of data transfer — especially because there are so many different variants
- The bandwidthTest example is nowadays provided as a demo — it helps to get a quick overview of what bandwidth we can expect on the current device
- There are also several CLI args that may provide guidance for size/type dependence

→ <https://github.com/nvidia/cuda-samples>

1 UTILITIES CONT.

CUDA SDK CONT.

```
cuda-zen sh@n3073-004: 1_Utils$ which nvcc
cuda-zen sh@n3073-004: 1_Utils$ /gpfs/opt/sw/cuda-zen/spack-0.19.0/opt/spack/linux-almalinux8-zen/gcc-12.2.0/cuda-12.9.0-xznhx43d6
xdwudiu7xcutpx6ympzybi/extras/demo_suite/bandwidthTest
```

[CUDA Bandwidth Test] - Starting...

Running on...

Device 0: NVIDIA A100-PCIE-40GB

Quick Mode

Host to Device Bandwidth, 1 Device(s)

PINNED Memory Transfers

Transfer Size (Bytes)	Bandwidth(MB/s)
33554432	26114.0

Device to Host Bandwidth, 1 Device(s)

PINNED Memory Transfers

Transfer Size (Bytes)	Bandwidth(MB/s)
33554432	25440.6

Device to Device Bandwidth, 1 Device(s)

PINNED Memory Transfers

Transfer Size (Bytes)	Bandwidth(MB/s)
33554432	1205414.9

→ <https://github.com/nvidia/cuda-samples>

1 UTILITIES CONT.

CUDA SDK CONT.

```
cuda-zen sh@n3073-004: 1_Utils$ which nvcc
cuda-zen sh@n3073-004: 1_Utils$ /gpfs/opt/sw/cuda-zen/spack-0.19.0/opt/spack/linux-almalinux8-zen/gcc-12.2.0/cuda-12.9.0-xznhx43d6
xdwudiu7xcutpx6ympzybi/extras/demo_suite/bandwidthTest
```

[CUDA Bandwidth Test] - Starting...

Running on...

Device 0: NVIDIA A100-PCIE-40GB

Quick Mode

Host to Device Bandwidth, 1 Device(s)

PINNED Memory Transfers

Transfer Size (Bytes)	Bandwidth(MB/s)
33554432	26114.0

Device to Host Bandwidth, 1 Device(s)

PINNED Memory Transfers

Transfer Size (Bytes)	Bandwidth(MB/s)
33554432	25440.6

Device to Device Bandwidth, 1 Device(s)

PINNED Memory Transfers

Transfer Size (Bytes)	Bandwidth(MB/s)
33554432	1205414.9

32MB (pinned) with
≈25 GB/s

→ <https://github.com/nvidia/cuda-samples>

1 UTILITIES CONT.

CUDA SDK CONT.

```
cuda-zen sh@n3073-004: 1_Utils$ /gpfs/opt/sw/cuda-zen/spack-0.19.0/opt/spack/linux-almalinux8-zen/gcc-12.2.0/cuda-12.9.0-xznhx43d6
xdwudiuv7xcutpx6ympzybi/extras/demo_suite/bandwidthTest --help
cuda-zen sh@n3073-004: 1_Utils$ /gpfs/opt/sw/cuda-zen/spack-0.19.0/opt/spack/linux-almalinux8-zen/gcc-12.2.0/cuda-12.9.0-xznhx43d6
xdwudiuv7xcutpx6ympzybi/extras/demo_suite/bandwidthTest --mode=quick --memory=pageable --htod
[CUDA Bandwidth Test] - Starting...
Running on...

Device 0: NVIDIA A100-PCIE-40GB
Quick Mode

Host to Device Bandwidth, 1 Device(s)
PAGEABLE Memory Transfers
Transfer Size (Bytes)      Bandwidth(MB/s)
33554432                  16050.9

cuda-zen sh@n3073-004: 1_Utils$ /gpfs/opt/sw/cuda-zen/spack-0.19.0/opt/spack/linux-almalinux8-zen/gcc-12.2.0/cuda-12.9.0-xznhx43d6
xdwudiuv7xcutpx6ympzybi/extras/demo_suite/bandwidthTest --mode=quick --memory=pinned --htod
[CUDA Bandwidth Test] - Starting...
Running on...

Device 0: NVIDIA A100-PCIE-40GB
Quick Mode

Host to Device Bandwidth, 1 Device(s)
PINNED Memory Transfers
Transfer Size (Bytes)      Bandwidth(MB/s)
33554432                  26127.6
```

→ <https://github.com/nvidia/cuda-samples>

1 UTILITIES CONT.

CUDA SDK CONT.

```
cuda-zen sh@n3073-004: 1_Utilsities$ /gpfs/opt/sw/cuda-zen/spack-0.19.0/opt/spack/linux-almalinux8-zen/gcc-12.2.0/cuda-12.9.0-xznhx43d6
xdwudiu7xcutpx6ympzybi/extras/demo_suite/bandwidthTest --help
cuda-zen sh@n3073-004: 1_Utilsities$ /gpfs/opt/sw/cuda-zen/spack-0.19.0/opt/spack/linux-almalinux8-zen/gcc-12.2.0/cuda-12.9.0-xznhx43d6
xdwudiu7xcutpx6ympzybi/extras/demo_suite/bandwidthTest --mode=quick --memory=pageable --htod
[CUDA Bandwidth Test] - Starting...
Running on...
```

Device 0: NVIDIA A100-PCIE-40GB
Quick Mode

Host to Device Bandwidth, 1 Device(s)

PAGEABLE Memory Transfers

Transfer Size (Bytes)	Bandwidth(MB/s)
33554432	16050.9

```
cuda-zen sh@n3073-004: 1_Utilsities$ /gpfs/opt/sw/cuda-zen/spack-0.19.0/opt/spack/linux-almalinux8-zen/gcc-12.2.0/cuda-12.9.0-xznhx43d6
xdwudiu7xcutpx6ympzybi/extras/demo_suite/bandwidthTest --mode=quick --memory=pinned --htod
[CUDA Bandwidth Test] - Starting...
Running on...
```

Device 0: NVIDIA A100-PCIE-40GB
Quick Mode

Host to Device Bandwidth, 1 Device(s)

PINNED Memory Transfers

Transfer Size (Bytes)	Bandwidth(MB/s)
33554432	26127.6

PINNED is
preferable to
PAGEABLE

→ <https://github.com/nvidia/cuda-samples>

1 UTILITIES CONT.

CUDA SDK CONT.

```
cuda-zen sh@n3073-004: 1_Utilsities$ /gpfs/opt/sw/cuda-zen/spack-0.19.0/opt/spack/linux-almalinux8-zen/gcc-12.2.0/cuda-12.9.0-xznhx43d6
xdwudiuv7xcutpx6ympzybi/extras/demo_suite/bandwidthTest --help
cuda-zen sh@n3073-004: 1_Utilsities$ /gpfs/opt/sw/cuda-zen/spack-0.19.0/opt/spack/linux-almalinux8-zen/gcc-12.2.0/cuda-12.9.0-xznhx43d6
xdwudiuv7xcutpx6ympzybi/extras/demo_suite/bandwidthTest --mode=quick --memory=pageable --htod
[CUDA Bandwidth Test] - Starting...
Running on...
```

```
Device 0: NVIDIA A100-PCIE-40GB
Quick Mode
```

```
Host to Device Bandwidth, 1 Device(s)
```

```
PAGEABLE Memory Transfers
```

Transfer Size (Bytes)	Bandwidth(MB/s)
33554432	16050.9

```
cuda-zen sh@n3073-004: 1_Utilsities$ /gpfs/opt/sw/cuda-zen/spack-0.19.0/opt/spack/linux-almalinux8-zen/gcc-12.2.0/cuda-12.9.0-xznhx43d6
xdwudiuv7xcutpx6ympzybi/extras/demo_suite/bandwidthTest --mode=quick --memory=pinned --htod
[CUDA Bandwidth Test] - Starting...
```

```
Running on...
```

```
Device 0: NVIDIA A100-PCIE-40GB
Quick Mode
```

```
Host to Device Bandwidth, 1 Device(s)
```

```
PINNED Memory Transfers
```

Transfer Size (Bytes)	Bandwidth(MB/s)
33554432	26127.6

Why so far off
the promised
1555 GB/s ???

PINNED is
preferable to
PAGEABLE

→ <https://github.com/nvidia/cuda-samples>

1 UTILITIES CONT.

CUDA SDK CONT.

bandwidthTest.cu (essentials)

```
float
testDeviceToHostTransfer(unsigned int memSize, memoryMode memMode, bool wc)
{
    StopWatchInterface *timer = NULL;
    float elapsedTimeInMs = 0.0f;
    float bandwidthInGBs = 0.0f;
    unsigned char *h_idata = NULL;
    unsigned char *h_odata = NULL;
    cudaEvent_t start, stop;

    sdkCreateTimer(&timer);
    checkCudaErrors(cudaEventCreate(&start));
    checkCudaErrors(cudaEventCreate(&stop));

    //allocate host memory
    if (PINNED == memMode)
    {
        //pinned memory mode - use special function to get OS-pinned memory
        checkCudaErrors(cudaHostAlloc((void **) &h_idata, memSize, (wc) ? cudaHostAllocWriteCombined : 0));
        checkCudaErrors(cudaHostAlloc((void **) &h_odata, memSize, (wc) ? cudaHostAllocWriteCombined : 0));
    }
    else {
        //pageable memory mode - use malloc
    }
}
```

→ <https://github.com/NVIDIA/cuda-samples/archive/refs/tags/v12.8.zip>

1 UTILITIES CONT.

CUDA SDK CONT.

bandwidthTest.cu (essentials)

```
float
testDeviceToHostTransfer(unsigned int memSize, memoryMode memMode, bool wc)
{
    StopWatchInterface *timer = NULL;
    float elapsedTimeInMs = 0.0f;
    float bandwidthInGBs = 0.0f;
    unsigned char *h_idata = NULL;
    unsigned char *h_odata = NULL;
    cudaEvent_t start, stop;

    sdkCreateTimer(&timer);
    checkCudaErrors(cudaEventCreate(&start));
    checkCudaErrors(cudaEventCreate(&stop));

    //allocate host memory
    if (PINNED == memMode)
    {
        //pinned memory mode - use special function to get OS-pinned memory
        checkCudaErrors(cudaHostAlloc((void **) &h_idata, memSize, (wc) ? cudaHostAllocWriteCombined : 0));
        checkCudaErrors(cudaHostAlloc((void **) &h_odata, memSize, (wc) ? cudaHostAllocWriteCombined : 0));
    }
    else {
        //pageable memory mode - use malloc
    }
}
```

32000000

→ <https://github.com/NVIDIA/cuda-samples/archive/refs/tags/v12.8.zip>

1 UTILITIES CONT.

CUDA SDK CONT.

bandwidthTest.cu (essentials)

```
float
testDeviceToHostTransfer(unsigned int memSize, memoryMode memMode, bool wc)
{
    StopWatchInterface *timer = NULL;
    float elapsedTimeInMs = 0.0f;
    float bandwidthInGBs = 0.0f;
    unsigned char *h_idata = NULL;
    unsigned char *h_odata = NULL;
    cudaEvent_t start, stop;

    sdkCreateTimer(&timer);
    checkCudaErrors(cudaEventCreate(&start));
    checkCudaErrors(cudaEventCreate(&stop));

    //allocate host memory
    if (PINNED == memMode)
    {
        //pinned memory mode - use special function to get OS-pinned memory
        checkCudaErrors(cudaHostAlloc((void **) &h_idata, memSize, (wc) ? cudaHostAllocWriteCombined : 0));
        checkCudaErrors(cudaHostAlloc((void **) &h_odata, memSize, (wc) ? cudaHostAllocWriteCombined : 0));
    }
    else {
        //pageable memory mode - use malloc
    }
}
```

The code snippet shows the implementation of a function to test device-to-host transfer bandwidth. It uses CUDA timers to measure time and calculate bandwidth. The code handles pinned and pageable memory modes. Annotations highlight the variable 'memSize' and its value '32000000', and point to the CUDA types used in the function.

→ <https://github.com/NVIDIA/cuda-samples/archive/refs/tags/v12.8.zip>

1 UTILITIES CONT.

CUDA SDK CONT.

bandwidthTest.cu (essentials)

```
float
testDeviceToHostTransfer(unsigned int memSize, memoryMode memMode, bool wc)
{
    StopWatchInterface *timer = NULL;
    float elapsedTimeInMs = 0.0f;
    float bandwidthInGBs = 0.0f;
    unsigned char *h_idata = NULL;
    unsigned char *h_odata = NULL;
    cudaEvent_t start, stop;

    sdkCreateTimer(&timer);
    checkCudaErrors(cudaEventCreate(&start));
    checkCudaErrors(cudaEventCreate(&stop));

    //allocate host memory
    if (PINNED == memMode)
    {
        //pinned memory mode - use special function to get OS-pinned memory
        checkCudaErrors(cudaHostAlloc((void **) &h_idata, memSize, (wc) ? cudaHostAllocWriteCombined : 0));
        checkCudaErrors(cudaHostAlloc((void **) &h_odata, memSize, (wc) ? cudaHostAllocWriteCombined : 0));
    }
    else {
        //pageable memory mode - use malloc
    }

    32000000
```

The diagram illustrates the components of the `bandwidthTest.cu` code. A green box labeled "32000000" points to the variable `memSize`. Another green box labeled "CUDA types" points to the declaration of `StopWatchInterface *timer`, the event creation code, and the memory allocation code. A third green box labeled "CUDA time measurement" points to the timer creation code.

→ <https://github.com/NVIDIA/cuda-samples/archive/refs/tags/v12.8.zip>

1 UTILITIES CONT.

CUDA SDK CONT.

bandwidthTest.cu (essentials)

```
float
testDeviceToHostTransfer(unsigned int memSize, memoryMode memMode, bool wc)
{
    StopWatchInterface *timer = NULL;           32000000
    float elapsedTimeInMs = 0.0f;
    float bandwidthInGBs = 0.0f;
    unsigned char *h_idata = NULL;
    unsigned char *h_odata = NULL;
    cudaEvent_t start, stop;

    sdkCreateTimer(&timer);
    checkCudaErrors(cudaEventCreate(&start));
    checkCudaErrors(cudaEventCreate(&stop));

    //allocate host memory
    if (PINNED == memMode)
    {
        //pinned memory mode - use special function to get OS-pinned memory
        checkCudaErrors(cudaHostAlloc((void **) &h_idata, memSize, (wc) ? cudaHostAllocWriteCombined : 0));
        checkCudaErrors(cudaHostAlloc((void **) &h_odata, memSize, (wc) ? cudaHostAllocWriteCombined : 0));
    }
    else {
        //pageable memory mode - use malloc
    }

    CUDA types
```

32000000

CUDA time measurement

Error handling via encapsulation

→ <https://github.com/NVIDIA/cuda-samples/archive/refs/tags/v12.8.zip>

1 UTILITIES CONT.

CUDA SDK CONT.

bandwidthTest.cu (essentials)

```
//initialize the memory
for (unsigned int i = 0; i < memSize/sizeof(unsigned char); i++)
{
    h_idata[i] = (unsigned char)(i & 0xff);
}
// allocate device memory
unsigned char *d_idata;
checkCudaErrors(cudaMalloc((void **) &d_idata, memSize));

//initialize the device memory
checkCudaErrors(cudaMemcpy(d_idata, h_idata, memSize,
                           cudaMemcpyHostToDevice));

//copy data from GPU to Host
sdkStartTimer(&timer);
checkCudaErrors(cudaEventRecord(start, 0));
if (PINNED == memMode)
{
    for (unsigned int i = 0; i < MEMCOPY_ITERATIONS; i++)
    {
        checkCudaErrors(cudaMemcpyAsync(h_odata, d_idata, memSize,
                                       cudaMemcpyDeviceToHost, 0));
    }
} else { ... }
```

→ <https://github.com/NVIDIA/cuda-samples/archive/refs/tags/v12.8.zip>

1 UTILITIES CONT.

CUDA SDK CONT.

bandwidthTest.cu (essentials)

```
//initialize the memory
for (unsigned int i = 0; i < memSize/sizeof(unsigned char); i++) ←
{
    h_idata[i] = (unsigned char)(i & 0xff);
}
// allocate device memory
unsigned char *d_idata;
checkCudaErrors(cudaMalloc((void **) &d_idata, memSize));

//initialize the device memory
checkCudaErrors(cudaMemcpy(d_idata, h_idata, memSize,
                           cudaMemcpyHostToDevice));

//copy data from GPU to Host
sdkStartTimer(&timer);
checkCudaErrors(cudaEventRecord(start, 0));
if (PINNED == memMode)
{
    for (unsigned int i = 0; i < MEMCOPY_ITERATIONS; i++)
    {
        checkCudaErrors(cudaMemcpyAsync(h_odata, d_idata, memSize,
                                       cudaMemcpyDeviceToHost, 0));
    }
} else { ... }
```

Loop over 32M items

→ <https://github.com/NVIDIA/cuda-samples/archive/refs/tags/v12.8.zip>

1 UTILITIES CONT.

CUDA SDK CONT.

bandwidthTest.cu (essentials)

```
//initialize the memory
for (unsigned int i = 0; i < memSize/sizeof(unsigned char); i++) ←
{
    h_idata[i] = (unsigned char)(i & 0xff); ←
}
// allocate device memory
unsigned char *d_idata;
checkCudaErrors(cudaMalloc((void **) &d_idata, memSize));

//initialize the device memory
checkCudaErrors(cudaMemcpy(d_idata, h_idata, memSize,
                           cudaMemcpyHostToDevice));

//copy data from GPU to Host
sdkStartTimer(&timer);
checkCudaErrors(cudaEventRecord(start, 0));
if (PINNED == memMode)
{
    for (unsigned int i = 0; i < MEMCOPY_ITERATIONS; i++)
    {
        checkCudaErrors(cudaMemcpyAsync(h_odata, d_idata, memSize,
                                       cudaMemcpyDeviceToHost, 0));
    }
} else { ... }
```

Loop over 32M items

bitwise add like ($i \% 255$)

→ <https://github.com/NVIDIA/cuda-samples/archive/refs/tags/v12.8.zip>

1 UTILITIES CONT.

CUDA SDK CONT.

bandwidthTest.cu (essentials)

```
//initialize the memory
for (unsigned int i = 0; i < memSize/sizeof(unsigned char); i++) ←
{
    h_idata[i] = (unsigned char)(i & 0xff); ←
}
// allocate device memory
unsigned char *d_idata;
checkCudaErrors(cudaMalloc((void **) &d_idata, memSize)); ←
//initialize the device memory
checkCudaErrors(cudaMemcpy(d_idata, h_idata, memSize, ←
                           cudaMemcpyHostToDevice)); ←
                           Memory set up on GPU

//copy data from GPU to Host
sdkStartTimer(&timer);
checkCudaErrors(cudaEventRecord(start, 0));
if (PINNED == memMode)
{
    for (unsigned int i = 0; i < MEMCOPY_ITERATIONS; i++)
    {
        checkCudaErrors(cudaMemcpyAsync(h_odata, d_idata, memSize, ←
                                       cudaMemcpyDeviceToHost, 0));
    }
} else { ... }
```

Loop over 32M items

bitwise add like ($i \% 255$)

Memory set up on GPU

→ <https://github.com/NVIDIA/cuda-samples/archive/refs/tags/v12.8.zip>

1 UTILITIES CONT.

CUDA SDK CONT.

bandwidthTest.cu (essentials)

```
//initialize the memory
for (unsigned int i = 0; i < memSize/sizeof(unsigned char); i++) ←
{
    h_idata[i] = (unsigned char)(i & 0xff); ←
}
// allocate device memory
unsigned char *d_idata;
checkCudaErrors(cudaMalloc((void **) &d_idata, memSize)); ←
//initialize the device memory
checkCudaErrors(cudaMemcpy(d_idata, h_idata, memSize, ←
                           cudaMemcpyHostToDevice)); ←
//copy data from GPU to Host
sdkStartTimer(&timer);
checkCudaErrors(cudaEventRecord(start, 0)); ←
if (PINNED == memMode)
{
    for (unsigned int i = 0; i < MEMCOPY_ITERATIONS; i++)
    {
        checkCudaErrors(cudaMemcpyAsync(h_odata, d_idata, memSize, ←
                                       cudaMemcpyDeviceToHost, 0));
    }
} else { ... }
```

Loop over 32M items

bitwise add like ($i \% 255$)

Memory set up on GPU

Timing begin

→ <https://github.com/NVIDIA/cuda-samples/archive/refs/tags/v12.8.zip>

1 UTILITIES CONT.

CUDA SDK CONT.

bandwidthTest.cu (essentials)

```
//initialize the memory
for (unsigned int i = 0; i < memSize/sizeof(unsigned char); i++) ←
{
    h_idata[i] = (unsigned char)(i & 0xff); ←
}
// allocate device memory
unsigned char *d_idata;
checkCudaErrors(cudaMalloc((void **) &d_idata, memSize)); ←
//initialize the device memory
checkCudaErrors(cudaMemcpy(d_idata, h_idata, memSize, ←
                           cudaMemcpyHostToDevice)); ←
//copy data from GPU to Host
sdkStartTimer(&timer);
checkCudaErrors(cudaEventRecord(start, 0)); ←
if (PINNED == memMode)
{
    for (unsigned int i = 0; i < MEMCOPY_ITERATIONS; i++) ←
    {
        checkCudaErrors(cudaMemcpyAsync(h_odata, d_idata, memSize, ←
                                       cudaMemcpyDeviceToHost, 0));
    }
} else { ... }
```

Loop over 32M items

bitwise add like $(i \% 255)$

Memory set up on GPU

Timing begin

100

→ <https://github.com/NVIDIA/cuda-samples/archive/refs/tags/v12.8.zip>

1 UTILITIES CONT.

CUDA SDK CONT.

bandwidthTest.cu (essentials)

```
checkCudaErrors(cudaEventRecord(stop, 0));

// make sure GPU has finished copying
checkCudaErrors(cudaDeviceSynchronize());
//get the total elapsed time in ms
sdkStopTimer(&timer);
checkCudaErrors(cudaEventElapsedTime(&elapsedTimeInMs, start, stop));

//calculate bandwidth in GB/s
double time_s = elapsedTimeInMs / 1e3;
bandwidthInGBs = (memSize * (float)MEMCOPY_ITERATIONS) / (double)1e9;
bandwidthInGBs = bandwidthInGBs / time_s;
//clean up memory
checkCudaErrors(cudaEventDestroy(stop));
checkCudaErrors(cudaEventDestroy(start));
sdkDeleteTimer(&timer);

... freeing allocated memory

return bandwidthInGBs;
}
```

→ <https://github.com/NVIDIA/cuda-samples/archive/refs/tags/v12.8.zip>

1 UTILITIES CONT.

CUDA SDK CONT.

bandwidthTest.cu (essentials)

```
checkCudaErrors(cudaEventRecord(stop, 0)); ←
// make sure GPU has finished copying
checkCudaErrors(cudaDeviceSynchronize());
//get the total elapsed time in ms
sdkStopTimer(&timer);
checkCudaErrors(cudaEventElapsedTime(&elapsedTimeInMs, start, stop));

//calculate bandwidth in GB/s
double time_s = elapsedTimeInMs / 1e3;
bandwidthInGBs = (memSize * (float)MEMCOPY_ITERATIONS) / (double)1e9;
bandwidthInGBs = bandwidthInGBs / time_s;
//clean up memory
checkCudaErrors(cudaEventDestroy(stop));
checkCudaErrors(cudaEventDestroy(start));
sdkDeleteTimer(&timer);

... freeing allocated memory

return bandwidthInGBs;
}
```

Timing end

→ <https://github.com/NVIDIA/cuda-samples/archive/refs/tags/v12.8.zip>

1 UTILITIES CONT.

CUDA SDK CONT.

bandwidthTest.cu (essentials)

```
checkCudaErrors(cudaEventRecord(stop, 0)); ← Timing end
// make sure GPU has finished copying
checkCudaErrors(cudaDeviceSynchronize());
//get the total elapsed time in ms
sdkStopTimer(&timer);
checkCudaErrors(cudaEventElapsedTime(&elapsedTimeInMs, start, stop)); ← Exe time in ms

//calculate bandwidth in GB/s
double time_s = elapsedTimeInMs / 1e3;
bandwidthInGBs = (memSize * (float)MEMCOPY_ITERATIONS) / (double)1e9;
bandwidthInGBs = bandwidthInGBs / time_s;
//clean up memory
checkCudaErrors(cudaEventDestroy(stop));
checkCudaErrors(cudaEventDestroy(start));
sdkDeleteTimer(&timer);

... freeing allocated memory

return bandwidthInGBs;
}
```

→ <https://github.com/NVIDIA/cuda-samples/archive/refs/tags/v12.8.zip>

1 UTILITIES CONT.

CUDA SDK CONT.

bandwidthTest.cu (essentials)

```
checkCudaErrors(cudaEventRecord(stop, 0));  
// make sure GPU has finished copying  
checkCudaErrors(cudaDeviceSynchronize());  
//get the total elapsed time in ms  
sdkStopTimer(&timer);  
checkCudaErrors(cudaEventElapsedTime(&elapsedTimeInMs, start, stop));  
  
//calculate bandwidth in GB/s  
double time_s = elapsedTimeInMs / 1e3;  
bandwidthInGBs = (memSize * (float)MEMCOPY_ITERATIONS) / (double)1e9;  
bandwidthInGBs = bandwidthInGBs / time_s;  
//clean up memory  
checkCudaErrors(cudaEventDestroy(stop));  
checkCudaErrors(cudaEventDestroy(start));  
sdkDeleteTimer(&timer);  
  
... freeing allocated memory  
  
return bandwidthInGBs;  
}
```

Timing end

Exe time in ms

Convert and compute bw

→ <https://github.com/NVIDIA/cuda-samples/archive/refs/tags/v12.8.zip>

1 UTILITIES CONT.

CUDA SDK CONT.

- Data transfer between host and device is the slowest link involved in GPU computing
- Needs to be carefully designed/minimized case-by-case
- Peak bandwidth disparity between device memory \leftrightarrow GPU cores (1555 GB/s on A100) and host memory \leftrightarrow device memory (25 GB/s PCIe Gen4)
- GPU receives pinned memory only, which is a temporary translation of pageable host memory
- That's why the directly allocated memory in pinned form is transferred faster
- Another optimization strategy is to overlap memory transfer with computing

→ <https://devblogs.nvidia.com/how-optimize-data-transfers-cuda-cc>

0_INTRODUCTION/SIMPLESTREAMS

CUDA SDK CONT.

- Efficiency comes with concurrently executing functions on all sorts of processing elements including CPU- and GPU-cores

→ <https://devblogs.nvidia.com/gpu-pro-tip-cuda-7-streams-simplify-concurrency>

CUDA 4 DUMMIES — OCT 22-23, 2025

0_INTRODUCTION/SIMPLESTREAMS

CUDA SDK CONT.

- Efficiency comes with concurrently executing functions on all sorts of processing elements including CPU- and GPU-cores
- CUDA applications manage concurrency with streams

→ <https://devblogs.nvidia.com/gpu-pro-tip-cuda-7-streams-simplify-concurrency>

0_INTRODUCTION/SIMPLESTREAMS

CUDA SDK CONT.

- Efficiency comes with concurrently executing functions on all sorts of processing elements including CPU- and GPU-cores
- CUDA applications manage concurrency with streams
- A stream is a sequence of commands executed in order

→ <https://devblogs.nvidia.com/gpu-pro-tip-cuda-7-streams-simplify-concurrency>

0_INTRODUCTION/SIMPLESTREAMS

CUDA SDK CONT.

- Efficiency comes with concurrently executing functions on all sorts of processing elements including CPU- and GPU-cores
- CUDA applications manage concurrency with streams
- A stream is a sequence of commands executed in order
- Several streams may execute their respective sequence of commands concurrently/asynchronously

0_INTRODUCTION/SIMPLESTREAMS

CUDA SDK CONT.

- Efficiency comes with concurrently executing functions on all sorts of processing elements including CPU- and GPU-cores
- CUDA applications manage concurrency with streams
- A stream is a sequence of commands executed in order
- Several streams may execute their respective sequence of commands concurrently/asynchronously
- With CUDA 7 control over more than one (default stream) was introduced, so that multiple host threads can now have their own associated default stream for launching kernels

→ <https://devblogs.nvidia.com/gpu-pro-tip-cuda-7-streams-simplify-concurrency>

0_INTRODUCTION/SIMPLESTREAMS

CUDA SDK CONT.

- Efficiency comes with concurrently executing functions on all sorts of processing elements including CPU- and GPU-cores
- CUDA applications manage concurrency with streams
- A stream is a sequence of commands executed in order
- Several streams may execute their respective sequence of commands concurrently/asynchronously
- With CUDA 7 control over more than one (default stream) was introduced, so that multiple host threads can now have their own associated default stream for launching kernels
- Asynchronous commands in CUDA return control to the calling host thread before the device has finished the requested task (non-blocking), e.g. kernel launches, memory copies performed by functions with the Async suffix, etc.

→ <https://devblogs.nvidia.com/gpu-pro-tip-cuda-7-streams-simplify-concurrency>

0__INTRODUCTION/SIMPLESTREAMS CONT.

CUDA SDK CONT.

CUDA streams

```
KrnlDmmmy <<< numBlocks, threadsPerBlock, numBytes >>> (); // default stream  
KrnlDmmmy <<< numBlocks, threadsPerBlock, numBytes, 0 >>> (); // stream 0
```

- <https://devblogs.nvidia.com/gpu-pro-tip-cuda-7-streams-simplify-concurrency>
- <https://stackoverflow.com/questions/27162408/shared-memory-and-streams-when-launching-kernel>

0__INTRODUCTION/SIMPLESTREAMS CONT.

CUDA SDK CONT.

CUDA streams

Specifying a stream
for a kernel launch (or
memcpy) is optional

```
KrnlDmmy <<< numBlocks, threadsPerBlock, numBytes >>> (); // default stream  
KrnlDmmy <<< numBlocks, threadsPerBlock, numBytes, 0 >>> (); // stream 0
```

- <https://devblogs.nvidia.com/gpu-pro-tip-cuda-7-streams-simplify-concurrency>
- <https://stackoverflow.com/questions/27162408/shared-memory-and-streams-when-launching-kernel>

0_INTRODUCTION/SIMPLESTREAMS CONT.

CUDA SDK CONT.

CUDA streams

Specifying a stream
for a kernel launch (or
memcpy) is optional

```
KrnlDmmmy <<< numBlocks, threadsPerBlock, numBytes >>> (); // default stream  
KrnlDmmmy <<< numBlocks, threadsPerBlock, numBytes, 0 >>> (); // stream 0
```

Stream 0 is the
default stream

- <https://devblogs.nvidia.com/gpu-pro-tip-cuda-7-streams-simplify-concurrency>
- <https://stackoverflow.com/questions/27162408/shared-memory-and-streams-when-launching-kernel>

0_INTRODUCTION/SIMPLESTREAMS CONT.

CUDA SDK CONT.

CUDA streams

Size of dynamically allocated shared memory

Specifying a stream for a kernel launch (or memcpy) is optional

```
KrnlDmmmy <<< numBlocks, threadsPerBlock, numBytes >>> (); // default stream  
KrnlDmmmy <<< numBlocks, threadsPerBlock, numBytes, 0 >>> (); // stream 0
```

Stream 0 is the default stream

- <https://devblogs.nvidia.com/gpu-pro-tip-cuda-7-streams-simplify-concurrency>
- <https://stackoverflow.com/questions/27162408/shared-memory-and-streams-when-launching-kernel>

0_INTRODUCTION/SIMPLESTREAMS CONT.

CUDA SDK CONT.

CUDA streams

Size of dynamically allocated shared memory

Specifying a stream for a kernel launch (or memcpy) is optional

```
KrnlDmmmy <<< numBlocks, threadsPerBlock, numBytes >>> (); // default stream  
KrnlDmmmy <<< numBlocks, threadsPerBlock, numBytes, 0 >>> (); // stream 0
```

Since CUDA 7 per-(host)thread default streams may be used

Stream 0 is the default stream

→ <https://devblogs.nvidia.com/gpu-pro-tip-cuda-7-streams-simplify-concurrency>

→ <https://stackoverflow.com/questions/27162408/shared-memory-and-streams-when-launching-kernel>

0_INTRODUCTION/SIMPLESTREAMS CONT.

CUDA SDK CONT.

CUDA streams

Size of dynamically allocated shared memory

Specifying a stream for a kernel launch (or memcpy) is optional

```
KrnlDmmmy <<< numBlocks, threadsPerBlock, numBytes >>> (); // default stream  
KrnlDmmmy <<< numBlocks, threadsPerBlock, numBytes, 0 >>> (); // stream 0
```

Since CUDA 7 per-(host)thread default streams may be used

Stream 0 is the default stream

Considered at compile time, nvcc --default-stream per-thread

→ <https://devblogs.nvidia.com/gpu-pro-tip-cuda-7-streams-simplify-concurrency>

→ <https://stackoverflow.com/questions/27162408/shared-memory-and-streams-when-launching-kernel>

0__INTRODUCTION/SIMPLESTREAMS CONT.

CUDA SDK CONT.

```
const int N = 1048576;

__global__ void kernel(float *x, int n)
{
    int tid = blockIdx.x * blockDim.x + threadIdx.x;
    for (int i = tid; i < n; i += blockDim.x) {
        x[i] = sqrt(pow(3.14159,i));
    }
}
int main()
{
    const int num_streams = 8;
    cudaStream_t streams[num_streams];
    float *data[num_streams];
    for (int i = 0; i < num_streams; i++) {
        cudaStreamCreate(&streams[i]);
        cudaMalloc(&data[i], N * sizeof(float));
        // launch one worker kernel per stream
        kernel <<< 1, 64, 0, streams[i]  >>> (data[i], N);
        // launch a dummy kernel on the default stream
        kernel <<< 1, 1 >>> (0, 0);
    }
    cudaDeviceReset();
    return 0;
}
```

Thread-specific run through array
x[] with stride blockDim.x

→ https://tinyurl.com/cudafordummies/ii/t/stream_test.cu

0_INTRODUCTION/SIMPLESTREAMS CONT.

CUDA SDK CONT.

```
const int N = 1048576;

__global__ void kernel(float *x, int n)
{
    int tid = blockIdx.x * blockDim.x + threadIdx.x;
    for (int i = tid; i <n; i += blockDim.x) {
        x[i] = sqrt(pow(3.14159,i));
    }
}
int main()
{
    const int num_streams = 8;
    cudaStream_t streams[num_streams];
    float *data[num_streams];
    for (int i = 0; i <num_streams; i++) {
        cudaStreamCreate(&streams[i]);
        cudaMalloc(&data[i], N * sizeof(float));
        // launch one worker kernel per stream
        kernel <<< 1, 64, 0, streams[i]  >>> (data[i], N);
        // launch a dummy kernel on the default stream
        kernel <<< 1, 1 >>> (0, 0);
    }
    cudaDeviceReset();
    return 0;
}
```

Thread-specific run through array
x[] with stride blockDim.x

Special type declaration

0_INTRODUCTION/SIMPLESTREAMS CONT.

CUDA SDK CONT.

```
const int N = 1048576;

__global__ void kernel(float *x, int n)
{
    int tid = blockIdx.x * blockDim.x + threadIdx.x;
    for (int i = tid; i < n; i += blockDim.x) {
        x[i] = sqrt(pow(3.14159,i));
    }
}

int main()
{
    const int num_streams = 8;
    cudaStream_t streams[num_streams];
    float *data[num_streams];
    for (int i = 0; i < num_streams; i++) {
        cudaStreamCreate(&streams[i]);
        cudaMalloc(&data[i], N * sizeof(float));
        // launch one worker kernel per stream
        kernel <<< 1, 64, 0, streams[i] >>> (data[i], N);
        // launch a dummy kernel on the default stream
        kernel <<< 1, 1 >>> (0, 0);
    }
    cudaDeviceReset();
    return 0;
}
```

Thread-specific run through array x[] with stride blockDim.x

Special type declaration

Stream creation and specific memory allocation

0_INTRODUCTION/SIMPLESTREAMS CONT.

CUDA SDK CONT.

```
const int N = 1048576;
```

```
__global__ void kernel(float *x, int n)
{
    int tid = blockIdx.x * blockDim.x + threadIdx.x;
    for (int i = tid; i < n; i += blockDim.x) {
        x[i] = sqrt(pow(3.14159,i));
    }
}
```

```
int main()
{
    const int num_streams = 8;
    cudaStream_t streams[num_streams];
    float *data[num_streams];
    for (int i = 0; i < num_streams; i++) {
        cudaStreamCreate(&streams[i]);
        cudaMalloc(&data[i], N * sizeof(float));
        // launch one worker kernel per stream
        kernel <<< 1, 64, 0, streams[i] >>> (data[i], N);
        // launch a dummy kernel on the default stream
        kernel <<< 1, 1 >>> (0, 0);
    }
    cudaDeviceReset();
    return 0;
}
```

Thread-specific run through array x[] with stride blockDim.x

Special type declaration

Stream creation and specific memory allocation

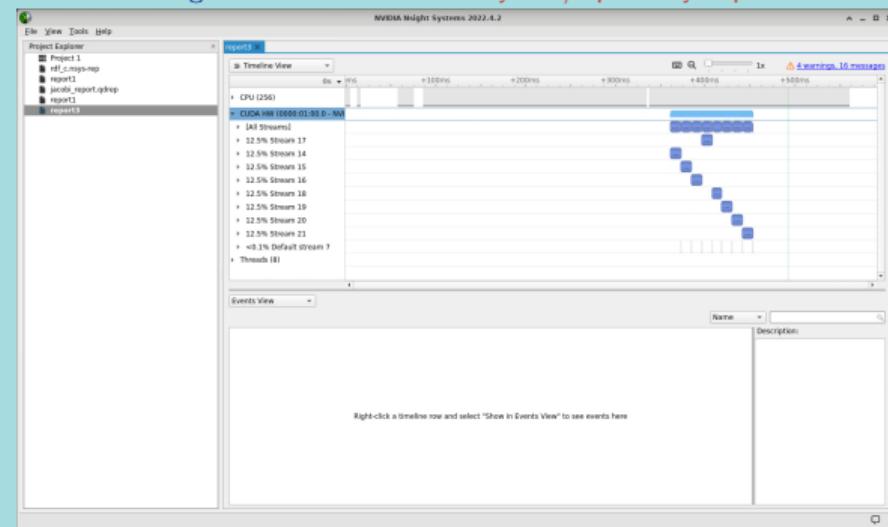
Kernel launch via streams

→ https://tinyurl.com/cudafordummies/ii/t/stream_test.cu

0_INTRODUCTION/SIMPLESTREAMS CONT.

CUDA SDK CONT.

```
cuda-zen sh@n3073-004: my_simpleStreams$ cp .../stream_test.cu ./simpleStreams.cu
cuda-zen sh@n3073-004: my_simpleStreams$ mkdir build
cuda-zen sh@n3073-004: my_simpleStreams$ cd build
cuda-zen sh@n3073-004: build$ cmake ..
cuda-zen sh@n3073-004: build$ make
cuda-zen sh@n3073-004: build$ nsys nvprof ./simpleStreams
cuda-zen sh@gui3068-009: build$ nsys-ui ./report1.nsys-rep
```

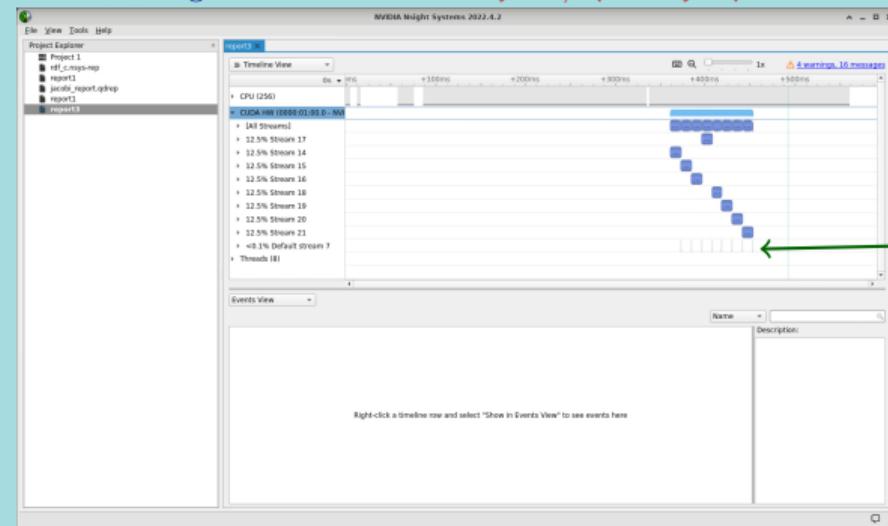


→ <https://devblogs.nvidia.com/gpu-pro-tip-cuda-7-streams-simplify-concurrency>

0_INTRODUCTION/SIMPLESTREAMS CONT.

CUDA SDK CONT.

```
cuda-zen sh@n3073-004: my_simpleStreams$ cp .../stream_test.cu ./simpleStreams.cu
cuda-zen sh@n3073-004: my_simpleStreams$ mkdir build
cuda-zen sh@n3073-004: my_simpleStreams$ cd build
cuda-zen sh@n3073-004: build$ cmake ..
cuda-zen sh@n3073-004: build$ make
cuda-zen sh@n3073-004: build$ nsys nvprof ./simpleStreams
cuda-zen sh@gui3068-009: build$ nsys-ui ./report1.nsys-rep
```



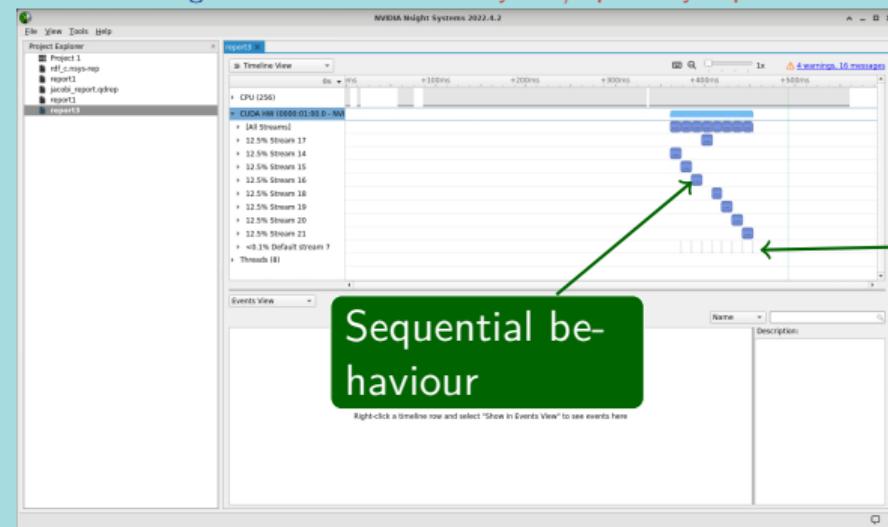
Interleaved
dummy kernel
sent to the default
stream → no con-
currency

→ <https://devblogs.nvidia.com/gpu-pro-tip-cuda-7-streams-simplify-concurrency>

0_INTRODUCTION/SIMPLESTREAMS CONT.

CUDA SDK CONT.

```
cuda-zen sh@n3073-004: my_simpleStreams$ cp .../stream_test.cu ./simpleStreams.cu
cuda-zen sh@n3073-004: my_simpleStreams$ mkdir build
cuda-zen sh@n3073-004: my_simpleStreams$ cd build
cuda-zen sh@n3073-004: build$ cmake ..
cuda-zen sh@n3073-004: build$ make
cuda-zen sh@n3073-004: build$ nsys nvprof ./simpleStreams
cuda-zen sh@gui3068-009: build$ nsys-ui ./report1.nsys-rep
```



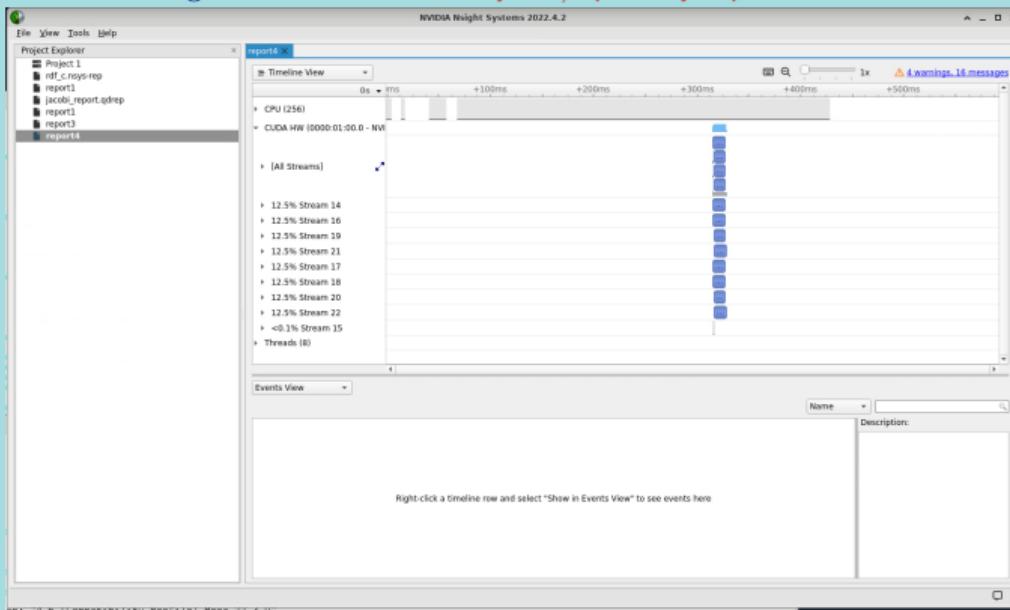
Interleaved dummy kernel sent to the default stream → no concurrency

→ <https://devblogs.nvidia.com/gpu-pro-tip-cuda-7-streams-simplify-concurrency>

0_INTRODUCTION/SIMPLESTREAMS CONT.

CUDA SDK CONT.

```
cuda-zen sh@n3073-004: build$ vi CMakeCache.txt # CMAKE_CUDA_FLAGS:STRING=--default-stream per-thread
cuda-zen sh@n3073-004: build$ make
cuda-zen sh@n3073-004: build$ nsys nvprof ./simpleStreams
cuda-zen sh@gui3068-009: build$ nsys-ui ./report2.nsys-rep
```

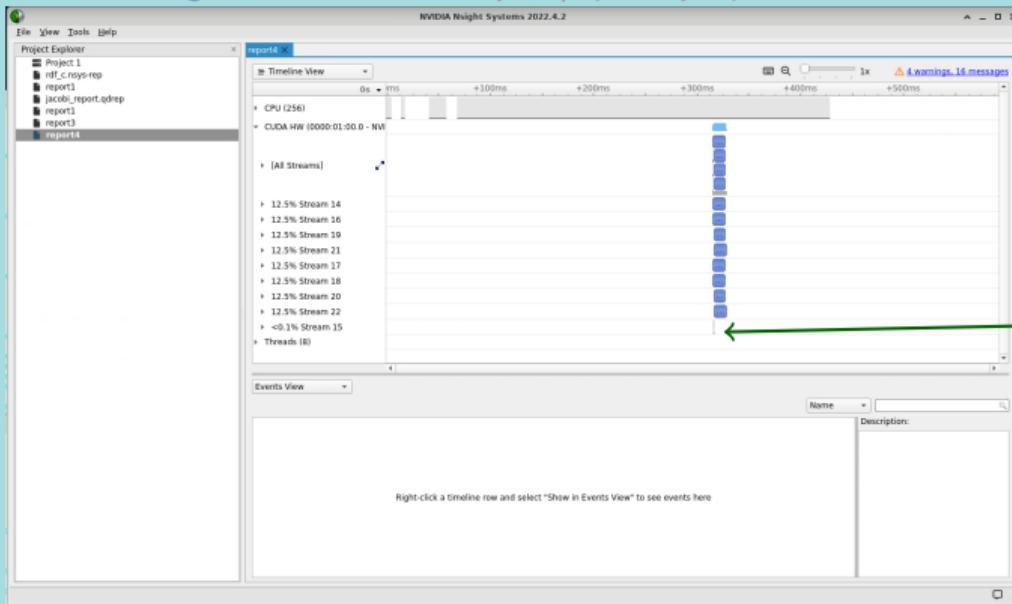


→ <https://devblogs.nvidia.com/gpu-pro-tip-cuda-7-streams-simplify-concurrency>

0_INTRODUCTION/SIMPLESTREAMS CONT.

CUDA SDK CONT.

```
cuda-zen sh@n3073-004: build$ vi CMakeCache.txt # CMAKE_CUDA_FLAGS:STRING=--default-stream per-thread
cuda-zen sh@n3073-004: build$ make
cuda-zen sh@n3073-004: build$ nsys nvprof ./simpleStreams
cuda-zen sh@gui3068-009: build$ nsys-ui ./report2.nsys-rep
```



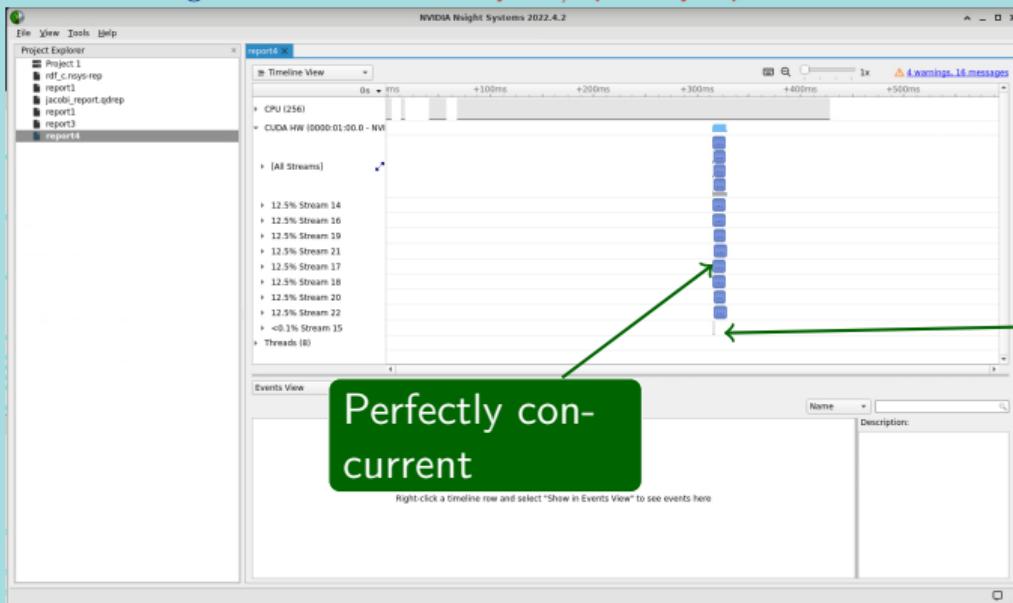
Dummy kernel
(sent to the de-
fault stream) in
parallel

→ <https://devblogs.nvidia.com/gpu-pro-tip-cuda-7-streams-simplify-concurrency>

0_INTRODUCTION/SIMPLESTREAMS CONT.

CUDA SDK CONT.

```
cuda-zen sh@n3073-004: build$ vi CMakeCache.txt # CMAKE_CUDA_FLAGS:STRING=--default-stream per-thread
cuda-zen sh@n3073-004: build$ make
cuda-zen sh@n3073-004: build$ nsys nvprof ./simpleStreams
cuda-zen sh@gui3068-009: build$ nsys-ui ./report2.nsys-rep
```



Dummy kernel
(sent to the de-
fault stream) in
parallel

→ <https://devblogs.nvidia.com/gpu-pro-tip-cuda-7-streams-simplify-concurrency>

0__INTRODUCTION/SIMPLESTREAMS CONT.

CUDA SDK CONT.

```
#include <omp.h>
const int N = 1048576;
__global__ void kernel(float *x, int n)
{
    int tid = blockIdx.x * blockDim.x + threadIdx.x;
    for (int i = tid; i <n; i += blockDim.x) {
        x[i] = sqrt(pow(3.14159,i));
    }
}
int main()
{
    const int num_streams = 8;
    cudaStream_t streams[num_streams];
    float *data[num_streams];
    omp_set_num_threads(num_streams);
    #pragma omp parallel for
    for (int i = 0; i <num_streams; i++) {
        cudaStreamCreate(&streams[i]);
        cudaMalloc(&data[i], N * sizeof(float));
        // launch one worker kernel per stream
        kernel <<< 1, 64, 0, streams[i]  >>> (data[i], N);
    }
    cudaDeviceReset();
    return 0;
}
```

Individual host-threads
on separate CPU cores
with associated stream

→ https://tinyurl.com/cudafordummies/ii/t/stream_test_v5.cu

0__INTRODUCTION/SIMPLESTREAMS CONT.

CUDA SDK CONT.

```
#include <omp.h>
const int N = 1048576;
__global__ void kernel(float *x, int n)
{
    int tid = blockIdx.x * blockDim.x + threadIdx.x;
    for (int i = tid; i <n; i += blockDim.x) {
        x[i] = sqrt(pow(3.14159,i));
    }
}
int main()
{
    const int num_streams = 8;
    cudaStream_t streams[num_streams];
    float *data[num_streams];
    omp_set_num_threads(num_streams);
    #pragma omp parallel for
    for (int i = 0; i <num_streams; i++) {
        cudaStreamCreate(&streams[i]);
        cudaMalloc(&data[i], N * sizeof(float));
        // launch one worker kernel per stream
        kernel <<< 1, 64, 0, streams[i] >>> (data[i], N);
    }
    cudaDeviceReset();
    return 0;
}
```

Individual host-threads
on separate CPU cores
with associated stream

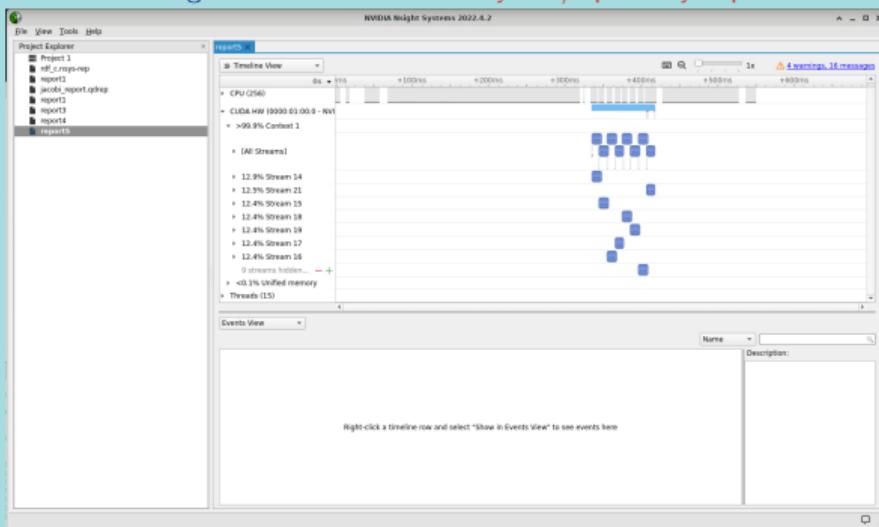
Simplest way of ex-
ploiting all available
CPU/GPU resources

→ https://tinyurl.com/cudafordummies/ii/t/stream_test_v5.cu

0_INTRODUCTION/SIMPLESTREAMS CONT.

CUDA SDK CONT.

```
cuda-zen sh@n3073-004: my_simpleStreams$ cp .../stream_test_v5.cu ./simpleStreams.cu
cuda-zen sh@n3073-004: my_simpleStreams$ cd build
cuda-zen sh@n3073-004: build$ vi CMakeCache.txt # CMAKE_CUDA_FLAGS:STRING==--default-stream per-thread -ccbin g++ -m64 -Xcompiler -fopenmp
cuda-zen sh@n3073-004: build$ make
cuda-zen sh@n3073-004: build$ nvprof ./simpleStreams
cuda-zen sh@gui3068-009: build$ nsys-ui ./report3.nsys-report
```

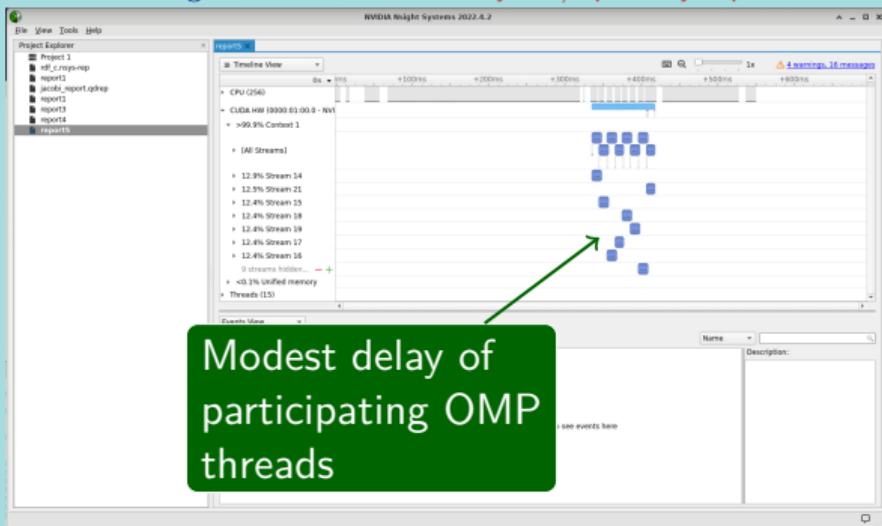


→ <https://devblogs.nvidia.com/gpu-pro-tip-cuda-7-streams-simplify-concurrency>

0_INTRODUCTION/SIMPLESTREAMS CONT.

CUDA SDK CONT.

```
cuda-zen sh@n3073-004: my_simpleStreams$ cp .../stream_test_v5.cu ./simpleStreams.cu
cuda-zen sh@n3073-004: my_simpleStreams$ cd build
cuda-zen sh@n3073-004: build$ vi CMakeCache.txt # CMAKE_CUDA_FLAGS:STRING==--default-stream per-thread -ccbin g++ -m64 -Xcompiler -fopenmp
cuda-zen sh@n3073-004: build$ make
cuda-zen sh@n3073-004: build$ nsys nvprof ./simpleStreams
cuda-zen sh@gui3068-009: build$ nsys-ui ./report3.nsys-report
```

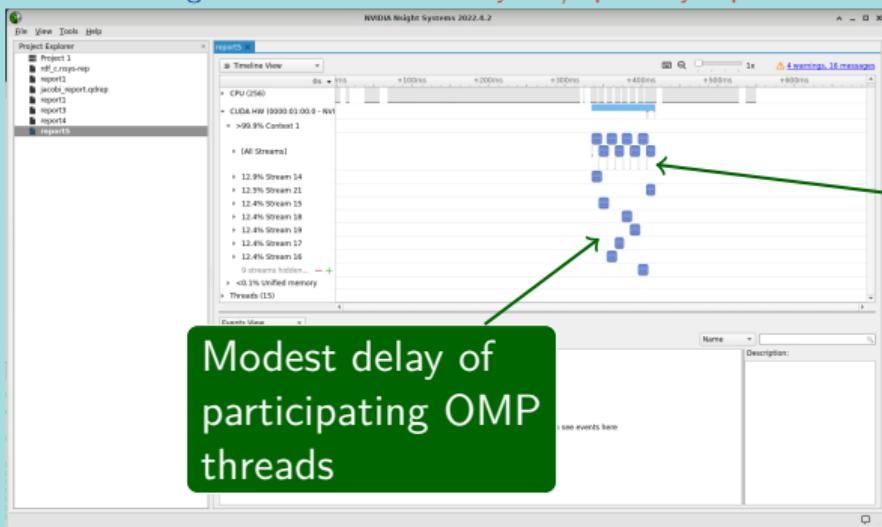


→ <https://devblogs.nvidia.com/gpu-pro-tip-cuda-7-streams-simplify-concurrency>

0_INTRODUCTION/SIMPLESTREAMS CONT.

CUDA SDK CONT.

```
cuda-zen sh@n3073-004: my_simpleStreams$ cp .../stream_test_v5.cu ./simpleStreams.cu
cuda-zen sh@n3073-004: my_simpleStreams$ cd build
cuda-zen sh@n3073-004: build$ vi CMakeCache.txt # CMAKE_CUDA_FLAGS:STRING=--default-stream per-thread -ccbin g++ -m64 -Xcompiler -fopenmp
cuda-zen sh@n3073-004: build$ make
cuda-zen sh@n3073-004: build$ nsys nvprof ./simpleStreams
cuda-zen sh@gui3068-009: build$ nsys-ui ./report3.nsys-report
```



Dummy kernel
perfectly concur-
rent !

→ <https://devblogs.nvidia.com/gpu-pro-tip-cuda-7-streams-simplify-concurrency>

0__INTRODUCTION/SIMPLESTREAMS CONT.

CUDA SDK CONT.

- Streams enable a lot of flexibility in CUDA workloads
- Only the legacy default stream can pose problems
- Compiler flag --default-stream per-thread needs to be applied to all *.cu units involved
- `cudaDeviceSynchronize()` continues to synchronize everything on the device
- Individual streams may be synchronized via `cudaStreamSynchronize()`
- Ruling out interference by the default stream completely may be achieved with non-blocking streams, i.e. by passing the flag `cudaStreamNonBlocking` to `cudaStreamCreate()`

→ <https://devblogs.nvidia.com/gpu-pro-tip-cuda-7-streams-simplify-concurrency>

4_CUDA_LIBRARIES/RANDOMFOG

CUDA SDK CONT.

- CUDA is for graphics cards, so there are a lot of graphics examples too
- Example 4_CUDA_Libraries/randomFog
- Random number generation (200k) with CURAND
- Spherical polar coordinates are used (radius, rho, theta) normalized and presented as uniform distribution on the sphere
- Several options to display the data set

→ ~/cuda-samples/Samples/4_CUDA_Libraries/my_randomFog/build/randomFog

5_DOMAIN_SPECIFIC/NBODY

CUDA SDK CONT.

- nbody is a CUDA demo of a gravitational n-body simulation
- Rather efficient scaling (strong) with multiple GPUs
- OpenGL rendering
- Command line args like -numbodies=10000 or -fp64 or -fullscreen or -cpu

→ [~/cuda-samples/Samples/5_Domain_Specific/my_nbody/build/nbody](#)

4_CUDA_LIBRARIES/OCEANFFT

CUDA SDK CONT.

- oceanFFT is a graphical demo of an ocean surface
- Height field is computed with the help of the CUFFT library (CUDA Fast Fourier Transform)
- OpenGL rendering
- 'w' — toggle wireframe

→ ~/cuda-samples/Samples/4_CUDA_Libraries/my_oceanFFT/build/oceanFFT

TAKE HOME MESSAGES

- CUDA SDK — a rich playground for beginners interested in learning the basics of GPU computing

TAKE HOME MESSAGES

- CUDA SDK — a rich playground for beginners interested in learning the basics of GPU computing
- Organized into different subject areas

TAKE HOME MESSAGES

- CUDA SDK — a rich playground for beginners interested in learning the basics of GPU computing
- Organized into different subject areas
- An exploratory space (memory management, error-handling...)

TAKE HOME MESSAGES

- CUDA SDK — a rich playground for beginners interested in learning the basics of GPU computing
- Organized into different subject areas
- An exploratory space (memory management, error-handling...)
- PCIe bandwidth remains a critical limitation in GPU computing

TAKE HOME MESSAGES

- CUDA SDK — a rich playground for beginners interested in learning the basics of GPU computing
- Organized into different subject areas
- An exploratory space (memory management, error-handling...)
- PCIe bandwidth remains a critical limitation in GPU computing
- CUDA streams bring in another level of flexibility, especially when run concurrently (perhaps from individual OpenMP threads on the host)

TAKE HOME MESSAGES

- CUDA SDK — a rich playground for beginners interested in learning the basics of GPU computing
- Organized into different subject areas
- An exploratory space (memory management, error-handling...)
- PCIe bandwidth remains a critical limitation in GPU computing
- CUDA streams bring in another level of flexibility, especially when run concurrently (perhaps from individual OpenMP threads on the host)
- Graphical demos — nice to have them too !