

CUDA SDK — LIBRARIES, NUMERICAL ACCURACY

Siegfried Höfinger

ASC Research Center, TU Wien

October 20, 2025

→ <https://tinyurl.com/cudafordummies/ii/l4/notes-l4.pdf>

OUTLINE

TO THE FORTRAN-MINDED

CUBLAS

CUSOLVER

MORE LIBRARIES

NUMERICAL ACCURACY & PERFORMANCE

5_DOMAIN_SPECIFIC/MARCHING CUBES

PROFILING CUDA CODE

HIP

TAKE HOME MESSAGES

TO THE FORTRAN-MINDED

CUDA SDK

- Fortran is the language of choice in many scientific code development projects

→ <https://developer.nvidia.com/hpc-sdk>

CUDA 4 DUMMIES — OCT 22-23, 2025

TO THE FORTRAN-MINDED

CUDA SDK

- Fortran is the language of choice in many scientific code development projects
- Professional coverage from the NVIDIA hpc SDK — formerly portland group compilers — now an entire suite of compilers, libraries and developmental tools

→ <https://developer.nvidia.com/hpc-sdk>

TO THE FORTRAN-MINDED

CUDA SDK

- Fortran is the language of choice in many scientific code development projects
- Professional coverage from the NVIDIA hpc SDK — formerly portland group compilers — now an entire suite of compilers, libraries and developmental tools
- NVFORTRAN, NVC++, NVC and NVCC compilers

→ <https://developer.nvidia.com/hpc-sdk>

TO THE FORTRAN-MINDED

CUDA SDK

- Fortran is the language of choice in many scientific code development projects
- Professional coverage from the NVIDIA hpc SDK — formerly portland group compilers — now an entire suite of compilers, libraries and developmental tools
- NVFORTRAN, NVC++, NVC and NVCC compilers
- Simply download from <https://developer.nvidia.com/nvidia-hpc-sdk-251-downloads>

→ <https://developer.nvidia.com/hpc-sdk>

TO THE FORTRAN-MINDED

CUDA SDK

- Fortran is the language of choice in many scientific code development projects
- Professional coverage from the NVIDIA hpc SDK — formerly portland group compilers — now an entire suite of compilers, libraries and developmental tools
- NVFORTRAN, NVC++, NVC and NVCC compilers
- Simply download from <https://developer.nvidia.com/nvidia-hpc-sdk-251-downloads>
- Follow guidelines in
<https://docs.nvidia.com/hpc-sdk/archive/25.1/hpc-sdk-install-guide/index.html>

→ <https://developer.nvidia.com/hpc-sdk>

TO THE FORTRAN-MINDED

CUDA SDK

- Fortran is the language of choice in many scientific code development projects
- Professional coverage from the NVIDIA hpc SDK — formerly portland group compilers — now an entire suite of compilers, libraries and developmental tools
- NVFORTRAN, NVC++, NVC and NVCC compilers
- Simply download from <https://developer.nvidia.com/nvidia-hpc-sdk-251-downloads>
- Follow guidelines in
<https://docs.nvidia.com/hpc-sdk/archive/25.1/hpc-sdk-install-guide/index.html>
- Install and load the module to get ready

→ <https://developer.nvidia.com/hpc-sdk>

TO THE FORTRAN-MINDED

CUDA SDK

- Fortran is the language of choice in many scientific code development projects
- Professional coverage from the NVIDIA hpc SDK — formerly portland group compilers — now an entire suite of compilers, libraries and developmental tools
- NVFORTRAN, NVC++, NVC and NVCC compilers
- Simply download from <https://developer.nvidia.com/nvidia-hpc-sdk-251-downloads>
- Follow guidelines in
<https://docs.nvidia.com/hpc-sdk/archive/25.1/hpc-sdk-install-guide/index.html>
- Install and load the module to get ready
- ```
cuda-zen sh@n3073-004:~$ module purge
cuda-zen sh@n3073-004:~$ module load nvhpc/25.1-gcc-9.5.0-efr6qch
```

→ <https://developer.nvidia.com/hpc-sdk>

# TO THE FORTRAN-MINDED CONT.

## CUDA SDK CONT.

### vecadd.cuf

```
! host
program t1
use cudafor
use myvecadd
integer, parameter :: n = 100
integer :: i
real, allocatable, device :: da(:, db(:, dc(:)
real :: ha(n), hb(n), hc(n)
istat = cudaSetDevice(0)
allocate(da(n))
...
da = ha
...
call vecadd<<< 1, n >>>(da, db, dc)
ha = da
...
deallocate(da)
...
end program t1
```

```
! kernel
module myvecadd
contains
attributes(global) &
subroutine vecadd(a, b, c)
 integer :: i
 real :: a(*), b(*), c(*)

 i = (blockidx%o-1) * blockDim%o + threadIdx%o
 c(i) = a(i) + b(i)
end subroutine vecadd
end module myvecadd
```

→ <https://tinyurl.com/cudafordummies/ii/l4/vecadd.cuf>

# TO THE FORTRAN-MINDED CONT.

## CUDA SDK CONT.

### vecadd.cuf

```
! host
program t1
use cudafor
use myvecadd
integer, parameter :: n = 100
integer :: i
real, allocatable, device :: da(:, db(:, dc(:)
real :: ha(n), hb(n), hc(n)
istat = cudaSetDevice(0) ←
allocate(da(n))
...
da = ha
...
call vecadd<<< 1, n >>>(da, db, dc)
ha = da
...
deallocate(da)
...
end program t1
```

Select GPU 0

```
! kernel
module myvecadd
contains
attributes(global) &
subroutine vecadd(a, b, c)
integer :: i
real :: a(*), b(*), c(*)

i = (blockidx%o-1) * blockDim%o + threadIdx%o
c(i) = a(i) + b(i)
end subroutine vecadd
end module myvecadd
```

→ <https://tinyurl.com/cudafordummie/i14/vecadd.cuf>

# TO THE FORTRAN-MINDED CONT.

## CUDA SDK CONT.

### vecadd.cuf

```
! host
program t1
use cudafor
use myvecadd
integer, parameter :: n = 100
integer :: i
real, allocatable, device :: da(:), db(:), dc(:)
real :: ha(n), hb(n), hc(n)
istat = cudaSetDevice(0) ←
allocate(da(n))
...
da = ha ←
...
call vecadd<<< 1, n >>>(da, db, dc) ←
ha = da ←
...
deallocate(da)
...
end program t1
```

Select GPU 0

Simple  
htod/dtoh  
copies

```
! kernel
module myvecadd
contains
attributes(global) &
subroutine vecadd(a, b, c)
integer :: i
real :: a(*), b(*), c(*)

i = (blockidx%o-1) * blockDim%o + threadIdx%o
c(i) = a(i) + b(i)
end subroutine vecadd
end module myvecadd
```

→ <https://tinyurl.com/cudafordummies/ii/l4/vecadd.cuf>

# TO THE FORTRAN-MINDED CONT.

## CUDA SDK CONT.

### vecadd.cuf

```
! host
program t1
use cudafor
use myvecadd
integer, parameter :: n = 100
integer :: i
real, allocatable, device :: da(:), db(:), dc(:)
real :: ha(n), hb(n), hc(n)
istat = cudaSetDevice(0)
allocate(da(n))
...
da = ha
...
call vecadd<<< 1, n >>>(da, db, dc)
ha = da
...
deallocate(da)
...
end program t1
```

Select GPU 0

Simple  
htod/dtoh  
copies

```
! kernel
module myvecadd
contains
attributes(global) &
subroutine vecadd(a, b, c)
integer :: i
real :: a(*), b(*), c(*)
i = (blockidx%<x-1) * blockdim%<x + threadidx%<x
c(i) = a(i) + b(i)
end subroutine vecadd
end module myvecadd
```

Slightly different syntax

→ <https://tinyurl.com/cudafordummies/ii/l4/vecadd.cuf>

# TO THE FORTRAN-MINDED CONT.

## CUDA SDK CONT.

### vecadd.cuf

```
! host
program t1
use cudafor
use myvecadd
integer, parameter :: n = 100
integer :: i
real, allocatable, device :: da(:), db(:), dc(:)
real :: ha(n), hb(n), hc(n)
istat = cudaSetDevice(0)
allocate(da(n))
...
da = ha
...
call vecadd<<< 1, n >>>(da, db, dc)
ha = da
...
deallocate(da)
...
end program t1
```

Select GPU 0

Simple  
htod/dtoh  
copies

```
! kernel
module myvecadd
contains
attributes(global) &
subroutine vecadd(a, b, c)
integer :: i
real :: a(*), b(*), c(*)

i = (blockidx%>-1) * blockDim%> + threadIdx%>
c(i) = a(i) + b(i)
end subroutine vecadd
end module myvecadd
```

Slightly different syntax

Fortran counting  
1,2,3...

→ <https://tinyurl.com/cudafordummies/ii/14/vecadd.cuf>

# TO THE FORTRAN-MINDED CONT.

## CUDA SDK CONT.

```
cuda-zen sh@n3073-004:~$ nvfortran vecadd.cuf
cuda-zen sh@n3073-004:~$./a.out

 1 1.000000 99.00000 100.0000
 2 2.000000 98.00000 100.0000
 3 3.000000 97.00000 100.0000
 4 4.000000 96.00000 100.0000
 5 5.000000 95.00000 100.0000

 98 98.00000 2.000000 100.0000
 99 99.00000 1.000000 100.0000
 100 100.0000 0.000000 100.0000
```

→ <https://tinyurl.com/cudafordummies/ii/14/vecadd.cuf>

## TO THE FORTRAN-MINDED CONT.

### CUDA SDK CONT.

- Frequently — only a single Fortran subroutine/function needs to be ported

→ <https://www.pgroup.com/resources/docs/18.4/pdf/pgi18cudaforug.pdf>

CUDA 4 DUMMIES — OCT 22-23, 2025

## TO THE FORTRAN-MINDED CONT.

### CUDA SDK CONT.

- Frequently — only a single Fortran subroutine/function needs to be ported
- Basic feasibility studies can already be done with the help of short C routines called from Fortran

→ <https://www.pgroup.com/resources/docs/18.4/pdf/pgi18cudaforug.pdf>

## TO THE FORTRAN-MINDED CONT.

### CUDA SDK CONT.

- Frequently — only a single Fortran subroutine/function needs to be ported
- Basic feasibility studies can already be done with the help of short C routines called from Fortran
- The usual Fortran↔C interfacing rules apply

→ <https://www.pgroup.com/resources/docs/18.4/pdf/pgi18cudaforug.pdf>

## TO THE FORTRAN-MINDED CONT.

### CUDA SDK CONT.

- Frequently — only a single Fortran subroutine/function needs to be ported
- Basic feasibility studies can already be done with the help of short C routines called from Fortran
- The usual Fortran↔C interfacing rules apply
- Memory management is more complicated, would prefer `cudaMallocManaged()`

# TO THE FORTRAN-MINDED CONT.

## CUDA SDK CONT.

### fvcadd.f

```
PROGRAM FVCADD
IMPLICIT NONE
INTEGER N
PARAMETER (N = 100)
INTEGER I
REAL A(N), B(N), C(N)

DO I=1, N
 A(I) = REAL(I)
 B(I) = REAL(N) - REAL(I)
 C(I) = REAL(0)
ENDDO

CALL NTMDTR(A, B, C, N)

DO I=1, N
 WRITE(6, '(I6F12.6)') I, C(I)
ENDDO

END
```

### ntmdtr.cu

```
__global__ void VecAdd(float *A, float *B, float *C)
{
 int i = threadIdx.x;
 C[i] = A[i] + B[i];
}

extern "C" void ntmdtr_(float *A, float *B, float *C, int *N)
{
 dim3 numBlocks, threadsPerBlock;
 float *AD, *BD, *CD;
 threadsPerBlock.x = *N;
 numBlocks.x = 1;
 cudaMalloc((void **) &AD, (*N) * sizeof(float));
 cudaMalloc((void **) &BD, (*N) * sizeof(float));
 cudaMalloc((void **) &CD, (*N) * sizeof(float));
 cudaMemcpy(AD, A, (*N) * sizeof(float), cudaMemcpyHostToDevice);
 cudaMemcpy(BD, B, (*N) * sizeof(float), cudaMemcpyHostToDevice);
 cudaMemcpy(CD, C, (*N) * sizeof(float), cudaMemcpyHostToDevice);
 VecAdd <<< numBlocks, threadsPerBlock >>> (AD, BD, CD);
 cudaDeviceSynchronize();
 cudaMemcpy(C, CD, (*N) * sizeof(float), cudaMemcpyDeviceToHost);
 cudaFree(AD);
 cudaFree(BD);
 cudaFree(CD);
 return;
}
```

# TO THE FORTRAN-MINDED CONT.

## CUDA SDK CONT.

### fvcadd.f

```
PROGRAM FVCADD
IMPLICIT NONE
INTEGER N
PARAMETER (N = 100)
INTEGER I
REAL A(N), B(N), C(N)

DO I=1, N
 A(I) = REAL(I)
 B(I) = REAL(N) - REAL(I)
 C(I) = REAL(0)
ENDDO

CALL NTMDTR(A, B, C, N)

DO I=1, N
 WRITE(6, '(I6F12.6)') I, C(I)
ENDDO

END
```

Intermediator in C

### ntmdtr.cu

```
__global__ void VecAdd(float *A, float *B, float *C)
{
 int i = threadIdx.x;
 C[i] = A[i] + B[i];
}

extern "C" void ntmdtr_(float *A, float *B, float *C, int *N)
{
 dim3 numBlocks, threadsPerBlock;
 float *AD, *BD, *CD;
 threadsPerBlock.x = *N;
 numBlocks.x = 1;
 cudaMalloc((void **) &AD, (*N) * sizeof(float));
 cudaMalloc((void **) &BD, (*N) * sizeof(float));
 cudaMalloc((void **) &CD, (*N) * sizeof(float));
 cudaMemcpy(AD, A, (*N) * sizeof(float), cudaMemcpyHostToDevice);
 cudaMemcpy(BD, B, (*N) * sizeof(float), cudaMemcpyHostToDevice);
 cudaMemcpy(CD, C, (*N) * sizeof(float), cudaMemcpyHostToDevice);
 VecAdd <<< numBlocks, threadsPerBlock >>> (AD, BD, CD);
 cudaDeviceSynchronize();
 cudaMemcpy(C, CD, (*N) * sizeof(float), cudaMemcpyDeviceToHost);
 cudaFree(AD);
 cudaFree(BD);
 cudaFree(CD);
 return;
}
```

# TO THE FORTRAN-MINDED CONT.

## CUDA SDK CONT.

fvcadd.f

```
PROGRAM FVCADD
IMPLICIT NONE
INTEGER N
PARAMETER (N = 100)
INTEGER I
REAL A(N), B(N), C(N)
```

```
DO I=1, N
 A(I) = REAL(I)
 B(I) = REAL(N) - REAL(I)
 C(I) = REAL(0)
ENDDO
```

```
CALL NTMDTR(A, B, C, N)
```

```
DO I=1, N
 WRITE(6, '(I6F12.6)') I, C(I)
ENDDO
```

```
END
```

Intermediator in C

ntmdtr.cu

```
__global__ void VecAdd(float *A, float *B, float *C)
{
 int i = threadIdx.x;
 C[i] = A[i] + B[i];
}

extern "C" void ntmdtr_(float *A, float *B, float *C, int *N)
{
 dim3 numBlocks, threadsPerBlock;
 float *AD, *BD, *CD;
 threadsPerBlock.x = *N;
 numBlocks.x = 1;
 cudaMalloc((void **) &AD, (*N) * sizeof(float));
 cudaMalloc((void **) &BD, (*N) * sizeof(float));
 cudaMalloc((void **) &CD, (*N) * sizeof(float));
 cudaMemcpy(AD, A, (*N) * sizeof(float), cudaMemcpyHostToDevice);
 cudaMemcpy(BD, B, (*N) * sizeof(float), cudaMemcpyHostToDevice);
 cudaMemcpy(CD, C, (*N) * sizeof(float), cudaMemcpyHostToDevice);
 VecAdd <<< numBlocks, threadsPerBlock >>> (AD, BD, CD);
 cudaDeviceSynchronize();
 cudaMemcpy(C, CD, (*N) * sizeof(float), cudaMemcpyDeviceToHost);
 cudaFree(AD);
 cudaFree(BD);
 cudaFree(CD);
 return;
}
```

Trailing underscore !

# TO THE FORTRAN-MINDED CONT.

CUDA SDK CONT.

**fvcadd.f**

```
PROGRAM FVCADD
IMPLICIT NONE
INTEGER N
PARAMETER (N = 100)
INTEGER I
REAL A(N), B(N), C(N)
```

```
DO I=1, N
 A(I) = REAL(I)
 B(I) = REAL(N) - REAL(I)
 C(I) = REAL(0)
ENDDO
```

```
CALL NTMDTR(A, B, C, N)
```

```
DO I=1, N
 WRITE(6, '(I6F12.6)') I, C(I)
ENDDO
```

```
END
```

Intermediator in C

**ntmdtr.cu**

```
global__ void VecAdd(float *A, float *B, float *C)
{
 int i = threadIdx.x;
 C[i] = A[i] + B[i];
}

extern "C" void ntmdtr_(float *A, float *B, float *C, int *N)
{
 dim3 numBlocks, threadsPerBlock;
 float *AD, *BD, *CD;
 threadsPerBlock.x = *N;
 numBlocks.x = 1;
 cudaMalloc((void **) &AD, (*N) * sizeof(float));
 cudaMalloc((void **) &BD, (*N) * sizeof(float));
 cudaMalloc((void **) &CD, (*N) * sizeof(float));
 cudaMemcpy(AD, A, (*N) * sizeof(float), cudaMemcpyHostToDevice);
 cudaMemcpy(BD, B, (*N) * sizeof(float), cudaMemcpyHostToDevice);
 cudaMemcpy(CD, C, (*N) * sizeof(float), cudaMemcpyHostToDevice);
 VecAdd <<< numBlocks, threadsPerBlock >>> (AD, BD, CD);
 cudaDeviceSynchronize();
 cudaMemcpy(C, CD, (*N) * sizeof(float), cudaMemcpyDeviceToHost);
 cudaFree(AD);
 cudaFree(BD);
 cudaFree(CD);
 return;
}
```

Trailing underscore !

Call by Reference

# TO THE FORTRAN-MINDED CONT.

## CUDA SDK CONT.

fvcadd.f

```
PROGRAM FVCADD
IMPLICIT NONE
INTEGER N
PARAMETER (N = 100)
INTEGER I
REAL A(N), B(N), C(N)
```

```
DO I=1, N
 A(I) = REAL(I)
 B(I) = REAL(N) - REAL(I)
 C(I) = REAL(0)
ENDDO
```

```
CALL NTMDTR(A, B, C, N)
```

```
DO I=1, N
 WRITE(6, '(I6F12.6)') I, C(I)
ENDDO
```

```
END
```

Intermediator in C

ntmdtr.cu

```
global__ void VecAdd(float *A, float *B, float *C)
{
 int i = threadIdx.x;
 C[i] = A[i] + B[i];
}

extern "C" void ntmdtr_(float *A, float *B, float *C, int *N)
{
 dim3 numBlocks, threadsPerBlock;
 float *AD, *BD, *CD;
 threadsPerBlock.x = *N;
 numBlocks.x = 1;
 cudaMalloc((void **) &AD, (*N) * sizeof(float));
 cudaMalloc((void **) &BD, (*N) * sizeof(float));
 cudaMalloc((void **) &CD, (*N) * sizeof(float));
 cudaMemcpy(AD, A, (*N) * sizeof(float), cudaMemcpyHostToDevice);
 cudaMemcpy(BD, B, (*N) * sizeof(float), cudaMemcpyHostToDevice);
 cudaMemcpy(CD, C, (*N) * sizeof(float), cudaMemcpyHostToDevice);
 VecAdd <<< numBlocks, threadsPerBlock >>> (AD, BD, CD);
 cudaDeviceSynchronize();
 cudaMemcpy(C, CD, (*N) * sizeof(float), cudaMemcpyDeviceToHost);
 cudaFree(AD);
 cudaFree(BD);
 cudaFree(CD);
 return;
}
```

Trailing underscore !

Call by Reference

Extra data transfer ↔

# TO THE FORTRAN-MINDED CONT.

## CUDA SDK CONT.

```
cuda-zen sh@n3073-004:~$ ls
fvcadd.f ntmdtr.cu

cuda-zen sh@n3073-004:~$ gfortran -c fvcadd.f
cuda-zen sh@n3073-004:~$ nvcc -c ntmdtr.cu
cuda-zen sh@n3073-004:~$ nvcc fvcadd.o ntmdtr.o -lcudart -lgfortran
cuda-zen sh@n3073-004:~$./a.out

 1 100.000000
 2 100.000000
 3 100.000000
 4 100.000000
 5 100.000000
 6 100.000000
 7 100.000000
 ...
98 100.000000
99 100.000000
100 100.000000
```

- <https://www.olcf.ornl.gov/tutorials/compiling-mixed-gpu-and-cpu-code>
- <https://tinyurl.com/cudafordummies/ii/l4/fvcadd.f>
- <https://tinyurl.com/cudafordummies/ii/l4/ntmdtr.cu>

# TO THE FORTRAN-MINDED CONT.

## CUDA SDK CONT.

```
cuda-zen sh@n3073-004:~$ ls
fvcadd.f ntmdtr.cu

cuda-zen sh@n3073-004:~$ gfortran -c fvcadd.f
cuda-zen sh@n3073-004:~$ nvcc -c ntmdtr.cu
cuda-zen sh@n3073-004:~$ nvcc fvcadd.o ntmdtr.o -lcudart -lgfortran
cuda-zen sh@n3073-004:~$./a.out

1 100.000000
2 100.000000
3 100.000000
4 100.000000
5 100.000000
6 100.000000
7 100.000000
....
98 100.000000
99 100.000000
100 100.000000
```

Name mangling 2b considered !

- <https://www.olcf.ornl.gov/tutorials/compiling-mixed-gpu-and-cpu-code>
- <https://tinyurl.com/cudafordummies/ii/l4/fvcadd.f>
- <https://tinyurl.com/cudafordummies/ii/l4/ntmdtr.cu>

# TO THE FORTRAN-MINDED CONT.

## CUDA SDK CONT.

```
cuda-zen sh@n3073-004:~$ ls
fvcadd.f ntmdtr.cu

cuda-zen sh@n3073-004:~$ gfortran -c fvcadd.f
cuda-zen sh@n3073-004:~$ nvcc -c ntmdtr.cu
cuda-zen sh@n3073-004:~$ nvcc fvcadd.o ntmdtr.o -lcudart -lgfortran
cuda-zen sh@n3073-004:~$./a.out

1 100.000000
2 100.000000
3 100.000000
4 100.000000
5 100.000000
6 100.000000
7 100.000000
....
98 100.000000
99 100.000000
100 100.000000
```

Name mangling 2b considered !

C++ like compiler !

- <https://www.olcf.ornl.gov/tutorials/compiling-mixed-gpu-and-cpu-code>
- <https://tinyurl.com/cudafordummies/ii/l4/fvcadd.f>
- <https://tinyurl.com/cudafordummies/ii/l4/ntmdtr.cu>

# TO THE FORTRAN-MINDED CONT.

## CUDA SDK CONT.

- Type analogues,  
int  $\leftrightarrow$  INTEGER  
float  $\leftrightarrow$  REAL  
double  $\leftrightarrow$  DOUBLE PRECISION
- Name mangling: compiling Fortran code will alter names of called subroutines/function by adding a trailing underscore, \_ Therefore C-functions will need to be aware of this and add this underscore ahead of time
- Fortran uses column-major order for storing multidimensional arrays while C uses row-major order
- Linking may need the addition of -lcudart -lgfortran etc

→ [http://www.computationalmathematics.org/topics/files/calling\\_cuda\\_from\\_fortran.html](http://www.computationalmathematics.org/topics/files/calling_cuda_from_fortran.html)

# CUBLAS

## CUDA SDK CONT.

- BLAS/LAPACK are among the most widely used libraries in scientific computing

→ <http://www.netlib.org/blas/>

# CUBLAS

## CUDA SDK CONT.

- BLAS/LAPACK are among the most widely used libraries in scientific computing
- General tasks in linear algebra with systematic naming scheme and a focus on high performance computing

→ <http://www.netlib.org/blas/>

# CUBLAS

## CUDA SDK CONT.

- BLAS/LAPACK are among the most widely used libraries in scientific computing
- General tasks in linear algebra with systematic naming scheme and a focus on high performance computing
- Standard API for many prominent implementations (MKL, OpenBLAS, ATLAS...)

→ <http://www.netlib.org/blas/>

# CUBLAS

## CUDA SDK CONT.

- BLAS/LAPACK are among the most widely used libraries in scientific computing
- General tasks in linear algebra with systematic naming scheme and a focus on high performance computing
- Standard API for many prominent implementations (MKL, OpenBLAS, ATLAS...)
- Originally written in Fortran

→ <http://www.netlib.org/blas/>

## CUBLAS CONT.

## CUDA SDK CONT.

S real, single precision

D real, double precision

C complex, single precision

Z complex, double precision

**DGEMM(...)**

## CUBLAS CONT.

## CUDA SDK CONT.

S real, single precision

D real, double precision

C complex, single precision

Z complex, double precision

DGEMM(...)

# CUBLAS CONT.

## CUDA SDK CONT.

SY symmetric HE Hermitian HP Hermitian packed TP triangular packed

SP symmetric packed GB general band HB Hermitian band TB triangular band

SB symmetric band GE general TR triangular

S real, single precision

D real, double precision

C complex, single precision

Z complex, double precision



## CUBLAS CONT.

CUDA SDK CONT.

**SY** symmetric      **HE** Hermitian      **HP** Hermitian packed    **TP** triangular packed  
**SP** symmetric packed    **GB** general band    **HB** Hermitian band    **TB** triangular band  
**SB** symmetric band     **GE** general        **TR** triangular

S real, single precision

## D real, double precision

## C complex, single precision

## Z complex, double precision

## RK rank-k matrix update

## R2K rank-2k matrix update

**U      unconjugate vector**

M modified Givens rot



MV matrix · vector

SV syst eq. 1 vec unk

## R rank-1 matrix update

## R2 rank-2 matrix update

**MM** matrix · matrix

SM syst eq. matrix unk

### C conjugate vector

## G Givens rotation

## MG mod. Givens rot constr

# CUBLAS CONT.

## CUDA SDK CONT.

```
...
DGEMM(TRANSA, TRANSB, M, N, K, ALPHA, A, LDA, B, LDB, BETA, C, LDC)
...

...
Atype = Btype = 'T';
Adim = Bdim = Cdim = n;
alpha = (double) 1;
beta = (double) 0;
info = dgemm(&Atype, &Btype, &Adim, &Bdim, &Cdim, &alpha,
 &A[0], &Adim, &B[0], &Bdim, &beta, &C[0], &Cdim);
if (info != 0) {
 printf("blas error in dgemm returning %d\n", info);
 exit(99);
}
...

stat = cublasDgemm (handle, dev_Atype, dev_Btype, Adim, Bdim, Cdim, &alpha,
 dev_ptr_A, Adim, dev_ptr_B, Bdim, &beta, dev_ptr_C, Cdim);
if (stat != CUBLAS_STATUS_SUCCESS) {
 printf("cublas error in dgemm \n");
 exit(99);
}
```

→ <http://www.netlib.org/blas/>

# CUBLAS CONT.

## CUDA SDK CONT.

```
...
DGEMM(TRANSA, TRANSB, M, N, K, ALPHA, A, LDA, B, LDB, BETA, C, LDC)
...
Atype = Btype = 'T';
Adim = Bdim = Cdim = n;
alpha = (double) 1;
beta = (double) 0;
info = dgemm(&Atype, &Btype, &Adim, &Bdim, &Cdim, &alpha,
 &A[0], &Adim, &B[0], &Bdim, &beta, &C[0], &Cdim);
if (info != 0) {
 printf("blas error in dgemm returning %d\n", info);
 exit(99);
}
...

stat = cublasDgemm (handle, dev_Atype, dev_Btype, Adim, Bdim, Cdim, &alpha,
 dev_ptr_A, Adim, dev_ptr_B, Bdim, &beta, dev_ptr_C, Cdim);
if (stat != CUBLAS_STATUS_SUCCESS) {
 printf("cublas error in dgemm \n");
 exit(99);
}
```

gfortran ... -lblas

→ <http://www.netlib.org/blas/>

# CUBLAS CONT.

## CUDA SDK CONT.

```
...
DGEMM(TRANSA, TRANSB, M, N, K, ALPHA, A, LDA, B, LDB, BETA, C, LDC)
...
```

gfortran ... -lblas

```
...
Atype = Btype = 'T';
Adim = Bdim = Cdim = n;
alpha = (double) 1;
beta = (double) 0;
info = dgemm(&Atype, &Btype, &Adim, &Bdim, &Cdim, &alpha,
 &A[0], &Adim, &B[0], &Bdim, &beta, &C[0], &Cdim);
```

```
if (info != 0) {
 printf("blas error in dgemm returning %d\n", info);
 exit(99);
}
```

gcc -Ddgemm=dgemm\_ ... -lblas

```
...
stat = cublasDgemm (handle, dev_Atype, dev_Btype, Adim, Bdim, Cdim, &alpha,
 dev_ptr_A, Adim, dev_ptr_B, Bdim, &beta, dev_ptr_C, Cdim);
```

```
if (stat != CUBLAS_STATUS_SUCCESS) {
 printf("cublas error in dgemm \n");
 exit(99);
}
```

→ <http://www.netlib.org/blas/>

# CUBLAS CONT.

## CUDA SDK CONT.

```
...
DGEMM(TRANSA, TRANSB, M, N, K, ALPHA, A, LDA, B, LDB, BETA, C, LDC)
...
```

gfortran ... -lblas

```
...
Atype = Btype = 'T';
Adim = Bdim = Cdim = n;
alpha = (double) 1;
beta = (double) 0;
info = dgemm(&Atype, &Btype, &Adim, &Bdim, &Cdim, &alpha,
 &A[0], &Adim, &B[0], &Bdim, &beta, &C[0], &Cdim);
```

```
if (info != 0) {
 printf("blas error in dgemm returning %d\n", info);
 exit(99);
}
```

gcc -Ddgemm=dgemm\_ ... -lblas

```
...
stat = cublasDgemm (handle, dev_Atype, dev_Btype, Adim, Bdim, Cdim, &alpha,
 dev_ptr_A, Adim, dev_ptr_B, Bdim, &beta, dev_ptr_C, Cdim);
```

```
if (stat != CUBLAS_STATUS_SUCCESS) {
 printf("cublas error in dgemm \n");
 exit(99);
}
```

nvcc ... -lcublas

→ <http://www.netlib.org/blas/>

# CUBLAS CONT.

## CUDA SDK CONT.

col-wise  
linearized  
matrices;  
 $\forall$  pointers

```
...
DGEMM(TRANS, TRANSB, M, N, K, ALPHA, A, LDA, B, LDB, BETA, C, LDC)
...
```

gfortran ... -lblas

```
...
Atype = Btype = 'T';
Adim = Bdim = Cdim = n;
alpha = (double) 1;
beta = (double) 0;
info = dgemm(&Atype, &Btype, &Adim, &Bdim, &Cdim, &alpha,
 &A[0], &Adim, &B[0], &Bdim, &beta, &C[0], &Cdim);
if (info != 0) {
 printf("blas error in dgemm returning %d\n", info);
 exit(99);
}
...
```

gcc -Ddgemm=dgemm\_ ... -lblas

```
...
stat = cublasDgemm (handle, dev_Atype, dev_Btype, Adim, Bdim, Cdim, &alpha,
 dev_ptr_A, Adim, dev_ptr_B, Bdim, &beta, dev_ptr_C, Cdim);
if (stat != CUBLAS_STATUS_SUCCESS) {
 printf("cublas error in dgemm \n");
 exit(99);
}
```

nvcc ... -lcublas

→ <http://www.netlib.org/blas/>

# CUBLAS CONT.

## CUDA SDK CONT.

col-wise  
linearized  
matrices;  
 $\forall$  pointers

```
...
DGEMM(TRANS, TRANSB, M, N, K, ALPHA, A, LDA, B, LDB, BETA, C, LDC)
...
```

gfortran ... -lblas

```
...
Atype = Btype = 'T';
Adim = Bdim = Cdim = n;
alpha = (double) 1;
beta = (double) 0;
info = dgemm(&Atype, &Btype, &Adim, &Bdim, &Cdim, &alpha,
 &A[0], &Adim, &B[0], &Bdim, &beta, &C[0], &Cdim);
if (info != 0) {
 printf("blas error in dgemm returning %d\n", info);
 exit(99);
}
...
```

gcc -Ddgemm=dgemm\_ ... -lblas

separate  
cublas  
arg/types

```
...
stat = cublasDgemm (handle, dev_Atype, dev_Btype, Adim, Bdim, Cdim, &alpha,
 dev_ptr_A, Adim, dev_ptr_B, Bdim, &beta, dev_ptr_C, Cdim);
if (stat != CUBLAS_STATUS_SUCCESS) {
 printf("cublas error in dgemm \n");
 exit(99);
}
```

nvcc ... -lcublas

→ <http://www.netlib.org/blas/>

## CUBLAS CONT.

### CUDA SDK CONT.

Porting code to CUBLAS requires a more-or-less standard procedure of the following steps:

1. Initiate the CUBLAS context (cublasCreate())
2. Allocate device memory using cudaMalloc() !
3. Transfer content of host arrays to the device (cublasSetMatrix())
4. Call a specific CUBLAS routine, e.g. cublasDgemm()
5. Transfer back the result from device memory to host (cublasGetMatrix())
6. Free device memory
7. Destroy CUBLAS context

→ <https://docs.nvidia.com/cuda/cublas/index.html>

→ <https://devtalk.nvidia.com/default/topic/1047981/b/t/post/5318441>

# CUBLAS CONT.

## CUDA SDK CONT.

### Example cublasDgemm()

```
#include <stdio.h>
#include <stdlib.h>
#include <cuda_runtime.h>
#include "cublas_v2.h"

int main(int argc, char **argv)
{
 int N, i, j;
 double alpha, beta, *A, *Adev, *B, *Bdev, *C, *Cdev;
 cublasStatus_t stat;
 cublasHandle_t handle;
 cublasOperation_t Atype, Btype;

 // memory allocation and parameter set up
 N = 5;
 alpha = (double) 1;
 beta = (double) 0;
 Atype = CUBLAS_OP_N;
 Btype = CUBLAS_OP_N;
 A = (double *) malloc(N * N * sizeof(double));
 B = (double *) malloc(N * N * sizeof(double));
 C = (double *) malloc(N * N * sizeof(double));

 // fill matrices A[] and B[] with data
```

→ <https://docs.nvidia.com/cuda/cublas/index.html>

# CUBLAS CONT.

## CUDA SDK CONT.

### Example cublasDgemm()

```
#include <stdio.h>
#include <stdlib.h>
#include <cuda_runtime.h>
#include "cublas_v2.h"

int main(int argc, char **argv)
{
 int N, i, j;
 double alpha, beta, *A, *Adev, *B, *Bdev, *C, *Cdev;
 cublasStatus_t stat;
 cublasHandle_t handle;
 cublasOperation_t Atype, Btype;

 // memory allocation and parameter set up
 N = 5;
 alpha = (double) 1;
 beta = (double) 0;
 Atype = CUBLAS_OP_N;
 Btype = CUBLAS_OP_N;
 A = (double *) malloc(N * N * sizeof(double));
 B = (double *) malloc(N * N * sizeof(double));
 C = (double *) malloc(N * N * sizeof(double));

 // fill matrices A[] and B[] with data
```

go for the new  
CUBLAS library API

→ <https://docs.nvidia.com/cuda/cublas/index.html>

# CUBLAS CONT.

## CUDA SDK CONT.

### Example cublasDgemm()

```
#include <stdio.h>
#include <stdlib.h>
#include <cuda_runtime.h>
#include "cublas_v2.h"

int main(int argc, char **argv)
{
 int N, i, j;
 double alpha, beta, *A, *Adev, *B, *Bdev, *C, *Cdev;
 cublasStatus_t stat;
 cublasHandle_t handle;
 cublasOperation_t Atype, Btype;

 // memory allocation and parameter set up
 N = 5;
 alpha = (double) 1;
 beta = (double) 0;
 Atype = CUBLAS_OP_N;
 Btype = CUBLAS_OP_N;
 A = (double *) malloc(N * N * sizeof(double));
 B = (double *) malloc(N * N * sizeof(double));
 C = (double *) malloc(N * N * sizeof(double));

 // fill matrices A[] and B[] with data
```

go for the new CUBLAS library API

will establish CUBLAS context

→ <https://docs.nvidia.com/cuda/cublas/index.html>

# CUBLAS CONT.

## CUDA SDK CONT.

### Example cublasDgemm()

```
#include <stdio.h>
#include <stdlib.h>
#include <cuda_runtime.h>
#include "cublas_v2.h"

int main(int argc, char **argv)
{
 int N, i, j;
 double alpha, beta, *A, *Adev, *B, *Bdev, *C, *Cdev;
 cublasStatus_t stat;
 cublasHandle_t handle;
 cublasOperation_t Atype, Btype;

 // memory allocation and parameter set up
 N = 5;
 alpha = (double) 1;
 beta = (double) 0;
 Atype = CUBLAS_OP_N;
 Btype = CUBLAS_OP_N;
 A = (double *) malloc(N * N * sizeof(double));
 B = (double *) malloc(N * N * sizeof(double));
 C = (double *) malloc(N * N * sizeof(double));

 // fill matrices A[] and B[] with data
```

go for the new  
CUBLAS library API

will establish CUBLAS context

CUBLAS way of  
TRANS/A=B='T'∨'N'

→ <https://docs.nvidia.com/cuda/cublas/index.html>

# CUBLAS CONT.

## CUDA SDK CONT.

### Example cublasDgemm()

```
#include <stdio.h>
#include <stdlib.h>
#include <cuda_runtime.h>
#include "cublas_v2.h"

int main(int argc, char **argv)
{
 int N, i, j;
 double alpha, beta, *A, *Adev, *B, *Bdev, *C, *Cdev;
 cublasStatus_t stat;
 cublasHandle_t handle;
 cublasOperation_t Atype, Btype;

 // memory allocation and parameter set up
 N = 5;
 alpha = (double) 1;
 beta = (double) 0;
 Atype = CUBLAS_OP_N;
 Btype = CUBLAS_OP_N;
 A = (double *) malloc(N * N * sizeof(double));
 B = (double *) malloc(N * N * sizeof(double));
 C = (double *) malloc(N * N * sizeof(double));

 // fill matrices A[] and B[] with data
```

go for the new  
CUBLAS library API

will establish CUBLAS context

CUBLAS way of  
TRANS/A=B='T'∨'N'

normal memory  
allocation

→ <https://docs.nvidia.com/cuda/cublas/index.html>

# CUBLAS CONT.

## CUDA SDK CONT.

### Example cublasDgemm() cont.

```
// initiate the CUBLAS context
cublasCreate(&handle);
// allocate device memory
cudaMalloc(&Adev, N * N * sizeof(double));
cudaMalloc(&Bdev, N * N * sizeof(double));
cudaMalloc(&Cdev, N * N * sizeof(double));
// copy contents of arrays A[], B[] into device memory
cublasSetMatrix(N, N, sizeof(double), &A[0], N, Adev, N);
cublasSetMatrix(N, N, sizeof(double), &B[0], N, Bdev, N);
// carry out CUBLAS operation
cublasDgemm(handle, Atype, Btype, N, N, N, &alpha, Adev, N, Bdev, N, &beta, Cdev, N);
// retrieve results from device memory
cublasGetMatrix(N, N, sizeof(double), Cdev, N, &C[0], N);
// do something with the result
... = C[]
// make clean
cudaFree(Cdev); cudaFree(Bdev); cudaFree(Adev);
cublasDestroy(handle);
// free host memory and return
free(C); free(B); free(A);
return(0);
}
```

→ <https://docs.nvidia.com/cuda/cublas/index.html>

# CUBLAS CONT.

## CUDA SDK CONT.

### Example `cublasDgemm()` cont.

```
// initiate the CUBLAS context
cublasCreate(&handle); ← Step 1
// allocate device memory
cudaMalloc(&Adev, N * N * sizeof(double));
cudaMalloc(&Bdev, N * N * sizeof(double));
cudaMalloc(&Cdev, N * N * sizeof(double));
// copy contents of arrays A[], B[] into device memory
cublasSetMatrix(N, N, sizeof(double), &A[0], N, Adev, N);
cublasSetMatrix(N, N, sizeof(double), &B[0], N, Bdev, N);
// carry out CUBLAS operation
cublasDgemm(handle, Atype, Btype, N, N, N, &alpha, Adev, N, Bdev, N, &beta, Cdev, N);
// retrieve results from device memory
cublasGetMatrix(N, N, sizeof(double), Cdev, N, &C[0], N);
// do something with the result
... = C[]
// make clean
cudaFree(Cdev); cudaFree(Bdev); cudaFree(Adev);
cublasDestroy(handle);
// free host memory and return
free(C); free(B); free(A);
return(0);
}
```

→ <https://docs.nvidia.com/cuda/cublas/index.html>

# CUBLAS CONT.

## CUDA SDK CONT.

### Example `cublasDgemm()` cont.

```
// initiate the CUBLAS context
cublasCreate(&handle); ← Step 1
// allocate device memory
cudaMalloc(&Adev, N * N * sizeof(double));
cudaMalloc(&Bdev, N * N * sizeof(double));
cudaMalloc(&Cdev, N * N * sizeof(double)); ← Step 2
// copy contents of arrays A[], B[] into device memory
cublasSetMatrix(N, N, sizeof(double), &A[0], N, Adev, N);
cublasSetMatrix(N, N, sizeof(double), &B[0], N, Bdev, N);
// carry out CUBLAS operation
cublasDgemm(handle, Atype, Btype, N, N, N, &alpha, Adev, N, Bdev, N, &beta, Cdev, N);
// retrieve results from device memory
cublasGetMatrix(N, N, sizeof(double), Cdev, N, &C[0], N);
// do something with the result
... = C[]
// make clean
cudaFree(Cdev); cudaFree(Bdev); cudaFree(Adev);
cublasDestroy(handle);
// free host memory and return
free(C); free(B); free(A);
return(0);
}
```

→ <https://docs.nvidia.com/cuda/cublas/index.html>

# CUBLAS CONT.

## CUDA SDK CONT.

### Example `cublasDgemm()` cont.

```
// initiate the CUBLAS context
cublasCreate(&handle); ← Step 1
// allocate device memory
cudaMalloc(&Adev, N * N * sizeof(double));
cudaMalloc(&Bdev, N * N * sizeof(double));
cudaMalloc(&Cdev, N * N * sizeof(double)); ← Step 2
// copy contents of arrays A[], B[] into device memory
cublasSetMatrix(N, N, sizeof(double), &A[0], N, Adev, N);
cublasSetMatrix(N, N, sizeof(double), &B[0], N, Bdev, N); ← Step 3
// carry out CUBLAS operation
cublasDgemm(handle, Atype, Btype, N, N, N, &alpha, Adev, N, Bdev, N, &beta, Cdev, N);
// retrieve results from device memory
cublasGetMatrix(N, N, sizeof(double), Cdev, N, &C[0], N);
// do something with the result
... = C[]
// make clean
cudaFree(Cdev); cudaFree(Bdev); cudaFree(Adev);
cublasDestroy(handle);
// free host memory and return
free(C); free(B); free(A);
return(0);
}
```

→ <https://docs.nvidia.com/cuda/cublas/index.html>

# CUBLAS CONT.

## CUDA SDK CONT.

### Example `cublasDgemm()` cont.

```
// initiate the CUBLAS context
cublasCreate(&handle); Step 1
// allocate device memory
cudaMalloc(&Adev, N * N * sizeof(double));
cudaMalloc(&Bdev, N * N * sizeof(double));
cudaMalloc(&Cdev, N * N * sizeof(double)); Step 2
// copy contents of arrays A[], B[] into device memory
cublasSetMatrix(N, N, sizeof(double), &A[0], N, Adev, N);
cublasSetMatrix(N, N, sizeof(double), &B[0], N, Bdev, N); Step 3
// carry out CUBLAS operation
cublasDgemm(handle, Atype, Btype, N, N, N, &alpha, Adev, N, Bdev, N, &beta, Cdev, N); Step 4
// retrieve results from device memory
cublasGetMatrix(N, N, sizeof(double), Cdev, N, &C[0], N);
// do something with the result
... = C[]
// make clean
cudaFree(Cdev); cudaFree(Bdev); cudaFree(Adev);
cublasDestroy(handle);
// free host memory and return
free(C); free(B); free(A);
return(0);
}
```

→ <https://docs.nvidia.com/cuda/cublas/index.html>

# CUBLAS CONT.

## CUDA SDK CONT.

### Example `cublasDgemm()` cont.

```
// initiate the CUBLAS context
cublasCreate(&handle); Step 1
// allocate device memory
cudaMalloc(&Adev, N * N * sizeof(double));
cudaMalloc(&Bdev, N * N * sizeof(double));
cudaMalloc(&Cdev, N * N * sizeof(double)); Step 2
// copy contents of arrays A[], B[] into device memory
cublasSetMatrix(N, N, sizeof(double), &A[0], N, Adev, N);
cublasSetMatrix(N, N, sizeof(double), &B[0], N, Bdev, N); Step 3
// carry out CUBLAS operation
cublasDgemm(handle, Atype, Btype, N, N, N, &alpha, Adev, N, Bdev, N, &beta, Cdev, N); Step 4
// retrieve results from device memory
cublasGetMatrix(N, N, sizeof(double), Cdev, N, &C[0], N);
// do something with the result
... = C[]
// make clean
cudaFree(Cdev); cudaFree(Bdev); cudaFree(Adev);
cublasDestroy(handle);
// free host memory and return
free(C); free(B); free(A);
return(0);
}
```

→ <https://docs.nvidia.com/cuda/cublas/index.html>

# CUBLAS CONT.

## CUDA SDK CONT.

### Example `cublasDgemm()` cont.

```
// initiate the CUBLAS context
cublasCreate(&handle); Step 1
// allocate device memory
cudaMalloc(&Adev, N * N * sizeof(double));
cudaMalloc(&Bdev, N * N * sizeof(double));
cudaMalloc(&Cdev, N * N * sizeof(double)); Step 2
// copy contents of arrays A[], B[] into device memory
cublasSetMatrix(N, N, sizeof(double), &A[0], N, Adev, N);
cublasSetMatrix(N, N, sizeof(double), &B[0], N, Bdev, N); Step 3
// carry out CUBLAS operation
cublasDgemm(handle, Atype, Btype, N, N, N, &alpha, Adev, N, Bdev, N, &beta, Cdev, N); Step 4
// retrieve results from device memory
cublasGetMatrix(N, N, sizeof(double), Cdev, N, &C[0], N); Step 5
// do something with the result
... = C[]
// make clean
cudaFree(Cdev); cudaFree(Bdev); cudaFree(Adev);
cublasDestroy(handle); Step 6
// free host memory and return
free(C); free(B); free(A);
return(0);
}
```

→ <https://docs.nvidia.com/cuda/cublas/index.html>

# CUBLAS CONT.

## CUDA SDK CONT.

### Example `cublasDgemm()` cont.

```
// initiate the CUBLAS context
cublasCreate(&handle); Step 1
// allocate device memory
cudaMalloc(&Adev, N * N * sizeof(double));
cudaMalloc(&Bdev, N * N * sizeof(double));
cudaMalloc(&Cdev, N * N * sizeof(double)); Step 2
// copy contents of arrays A[], B[] into device memory
cublasSetMatrix(N, N, sizeof(double), &A[0], N, Adev, N);
cublasSetMatrix(N, N, sizeof(double), &B[0], N, Bdev, N); Step 3
// carry out CUBLAS operation
cublasDgemm(handle, Atype, Btype, N, N, N, &alpha, Adev, N, Bdev, N, &beta, Cdev, N); Step 4
// retrieve results from device memory
cublasGetMatrix(N, N, sizeof(double), Cdev, N, &C[0], N); Step 5
// do something with the result
... = C[]
// make clean
cudaFree(Cdev); cudaFree(Bdev); cudaFree(Adev);
cublasDestroy(handle); Step 6
// free host memory and return
free(C); free(B); free(A);
return(0); Step 7
```

→ <https://docs.nvidia.com/cuda/cublas/index.html>

## CUBLAS CONT.

### CUDA SDK CONT.

Words of caution:

- Always start with a small explicit example to check/confirm each individual step
- Formats of matrices are crucial ! 1D-representation, column-wise linearized;
- (double) on consumer grade cards is slower by  $\approx 1/64$  than (float)
- Two new APIs, CUBLASXT (for CUDA  $\geq 6.0$ ) and cuBLASLt (for CUDA 10.1)
- Pity that we can't use `cudaMallocManaged()` at this point

→ <https://docs.nvidia.com/cuda/cublas/index.html>

→ <https://devtalk.nvidia.com/default/topic/1047981/b/t/post/5318441>

# CU Solver

CUDA SDK CONT.

Next level of problems:

$$\left( \begin{array}{ccccc} 1.96 & -6.49 & -0.47 & -7.20 & -0.65 \\ -6.49 & 3.80 & -6.39 & 1.50 & -6.34 \\ -0.47 & -6.39 & 4.17 & -1.51 & 2.67 \\ -7.20 & 1.50 & -1.51 & 5.70 & 1.80 \\ -0.65 & -6.34 & 2.67 & 1.80 & -7.10 \end{array} \right) \underbrace{\phantom{\left( \begin{array}{ccccc} 1.96 & -6.49 & -0.47 & -7.20 & -0.65 \\ -6.49 & 3.80 & -6.39 & 1.50 & -6.34 \\ -0.47 & -6.39 & 4.17 & -1.51 & 2.67 \\ -7.20 & 1.50 & -1.51 & 5.70 & 1.80 \\ -0.65 & -6.34 & 2.67 & 1.80 & -7.10 \end{array} \right)}}_{A}$$

- For example, what are the eigenvalues and corresponding eigenvectors of A ?
- $A x = \lambda x$
- A typical LAPACK problem — usually tightly coupled to BLAS
- Not on the GPU ! CUBLAS is BLAS only !
- However, there are cuSolver and MAGMA

→ <https://docs.nvidia.com/cuda/cusolver>

→ [https://icl.cs.utk.edu/projectsfiles/magma/doxygen/group\\_\\_cublas\\_\\_const.html](https://icl.cs.utk.edu/projectsfiles/magma/doxygen/group__cublas__const.html)

## CUSOLVER CONT.

### CUDA SDK CONT.

Making use of CUSOLVER is schematically very similar to using CUBLAS:

1. Initiate the CUSOLVER context (`cusolverDnCreate()`)
2. Allocate device memory using `cudaMalloc()` !
3. Transfer content of host arrays to the device (`cudaMemcpy()`)
4. Semi-automatically set up working space required by CUSOLVER routine, e.g. WORK, LWORK in LAPACK jargon
5. Call a specific CUSOLVER routine, e.g. `cusolverDnDsyevd()`
6. Transfer back the result from device memory to host (`cudaMemcpy()`)
7. Free device memory
8. Destroy CUSOLVER context

→ [https://docs.nvidia.com/cuda/cusolver/index.html#eig\\_examples](https://docs.nvidia.com/cuda/cusolver/index.html#eig_examples)

## CUSOLVER CONT.

### CUDA SDK CONT.

#### Example cusolverDnDsyevd()

```
#include <stdio.h>
#include <stdlib.h>
#include <assert.h>
#include <cuda_runtime.h>
#include "cusolverDn.h"

int main(int argc, char **argv)
{
 int i, j, lwork, info_gpu, *devInfo;
 double *d_A, *d_W, *d_work;
 cusolverDnHandle_t cusolverH;
 cusolverStatus_t cusolver_status;
 cusolverEigMode_t jobz;
 cublasFillMode_t uplo;
 cudaError_t cudaStat;
 const int m = 5;
 const int lda = m;
 double W[m];
 double V[lda*m];
 double A[lda*m] = { 1.96, -6.49, -0.47, -7.20, -0.65,
 -6.49, 3.80, -6.39, 1.50, -6.34,
 -0.47, -6.39, 4.17, -1.51, 2.67,
 -7.20, 1.50, -1.51, 5.70, 1.80,
 -0.65, -6.34, 2.67, 1.80, -7.10};
```

# CUSOLVER CONT.

## CUDA SDK CONT.

### Example cusolverDnDsyevd()

```
#include <stdio.h>
#include <stdlib.h>
#include <assert.h>
#include <cuda_runtime.h>
#include "cusolverDn.h"
```

```
int main(int argc, char **argv)
{
 int i, j, lwork, info_gpu, *devInfo;
 double *d_A, *d_W, *d_work;
 cusolverDnHandle_t cusolverH;
 cusolverStatus_t cusolver_status;
 cusolverEigMode_t jobz;
 cublasFillMode_t uplo;
 cudaError_t cudaStat;
 const int m = 5;
 const int lda = m;
 double W[m];
 double V[lda*m];
 double A[lda*m] = { 1.96, -6.49, -0.47, -7.20, -0.65,
 -6.49, 3.80, -6.39, 1.50, -6.34,
 -0.47, -6.39, 4.17, -1.51, 2.67,
 -7.20, 1.50, -1.51, 5.70, 1.80,
 -0.65, -6.34, 2.67, 1.80, -7.10};
```

CUSOLVER header  
for 'dense' subset

# CUSOLVER CONT.

## CUDA SDK CONT.

### Example cusolverDnDsyevd()

```
#include <stdio.h>
#include <stdlib.h>
#include <assert.h>
#include <cuda_runtime.h>
#include "cusolverDn.h"
```

```
int main(int argc, char **argv)
{
 int i, j, lwork, info_gpu, *devInfo;
 double *d_A, *d_W, *d_work;
 cusolverDnHandle_t cusolverH;
 cusolverStatus_t cusolver_status;
 cusolverEigMode_t jobz;
 cublasFillMode_t uplo;
 cudaError_t cudaStat;
 const int m = 5;
 const int lda = m;
 double W[m];
 double V[lda*m];
 double A[lda*m] = { 1.96, -6.49, -0.47, -7.20, -0.65,
 -6.49, 3.80, -6.39, 1.50, -6.34,
 -0.47, -6.39, 4.17, -1.51, 2.67,
 -7.20, 1.50, -1.51, 5.70, 1.80,
 -0.65, -6.34, 2.67, 1.80, -7.10};
```

CUSOLVER header  
for 'dense' subset

4 CUSOLVER context

# CUSOLVER CONT.

## CUDA SDK CONT.

### Example cusolverDnDsyevd()

```
#include <stdio.h>
#include <stdlib.h>
#include <assert.h>
#include <cuda_runtime.h>
#include "cusolverDn.h"
```

CUSOLVER header  
for 'dense' subset

```
int main(int argc, char **argv)
{
 int i, j, lwork, info_gpu, *devInfo;
 double *d_A, *d_W, *d_work;
 cusolverDnHandle_t cusolverH;
 cusolverStatus_t cusolver_status;
 cusolverEigMode_t jobz;
 cublasFillMode_t uplo;
 cudaError_t cudaStat;
 const int m = 5;
 const int lda = m;
 double W[m];
 double V[lda*m];
 double A[lda*m] = { 1.96, -6.49, -0.47, -7.20, -0.65,
 -6.49, 3.80, -6.39, 1.50, -6.34,
 -0.47, -6.39, 4.17, -1.51, 2.67,
 -7.20, 1.50, -1.51, 5.70, 1.80,
 -0.65, -6.34, 2.67, 1.80, -7.10};
```

4 CUSOLVER context

CUSOLVER flags 4 LAPACK's

# CUSOLVER CONT.

## CUDA SDK CONT.

### Example cusolverDnDsyevd() cont.

```
jobz = CUSOLVER_EIG_MODE_VECTOR;
uplo = CUBLAS_FILL_MODE_LOWER;
cusolver_status = cusolverDnCreate(&cusolverH);
assert(CUSOLVER_STATUS_SUCCESS == cusolver_status);
cudaStat = cudaMalloc ((void**)&d_A, sizeof(double) * lda * m);
cudaStat = cudaMalloc ((void**)&d_W, sizeof(double) * m);
cudaStat = cudaMalloc ((void**)&devInfo, sizeof(int));
cudaStat = cudaMemcpy(d_A, A, sizeof(double) * lda * m, cudaMemcpyHostToDevice);

cusolver_status = cusolverDnDsyevd_bufferSize(cusolverH, jobz, uplo, m, d_A, lda, d_W, &lwork);
cudaStat = cudaMalloc((void**)&d_work, sizeof(double)*lwork);
cusolver_status = cusolverDnDsyevd(cusolverH, jobz, uplo, m, d_A, lda, d_W, d_work, lwork, devInfo);
cudaStat = cudaDeviceSynchronize();

cudaStat = cudaMemcpy(W, d_W, sizeof(double)*m, cudaMemcpyDeviceToHost);
cudaStat = cudaMemcpy(V, d_A, sizeof(double)*lda*m, cudaMemcpyDeviceToHost);
cudaStat = cudaMemcpy(&info_gpu, devInfo, sizeof(int), cudaMemcpyDeviceToHost);
... do something with W[] and V[] ...
cudaFree(d_A); cudaFree(d_W); cudaFree(devInfo); cudaFree(d_work);
cusolverDnDestroy(cusolverH);
cudaDeviceReset();

return(0);
}
```

# CUSOLVER CONT.

## CUDA SDK CONT.

### Example cusolverDnDsyevd() cont.

```
jobz = CUSOLVER_EIG_MODE_VECTOR;
uplo = CUBLAS_FILL_MODE_LOWER;
cusolver_status = cusolverDnCreate(&cusolverH); ← establish CUSOLVER context
assert(CUSOLVER_STATUS_SUCCESS == cusolver_status);
cudaStat = cudaMalloc ((void**)&d_A, sizeof(double) * lda * m);
cudaStat = cudaMalloc ((void**)&d_W, sizeof(double) * m);
cudaStat = cudaMalloc ((void**)&devInfo, sizeof(int));
cudaStat = cudaMemcpy(d_A, A, sizeof(double) * lda * m, cudaMemcpyHostToDevice);

cusolver_status = cusolverDnDsyevd_bufferSize(cusolverH, jobz, uplo, m, d_A, lda, d_W, &lwork);
cudaStat = cudaMalloc((void**)&d_work, sizeof(double)*lwork);
cusolver_status = cusolverDnDsyevd(cusolverH, jobz, uplo, m, d_A, lda, d_W, d_work, lwork, devInfo);
cudaStat = cudaDeviceSynchronize();

cudaStat = cudaMemcpy(W, d_W, sizeof(double)*m, cudaMemcpyDeviceToHost);
cudaStat = cudaMemcpy(V, d_A, sizeof(double)*lda*m, cudaMemcpyDeviceToHost);
cudaStat = cudaMemcpy(&info_gpu, devInfo, sizeof(int), cudaMemcpyDeviceToHost);
... do something with W[] and V[] ...
cudaFree(d_A); cudaFree(d_W); cudaFree(devInfo); cudaFree(d_work);
cusolverDnDestroy(cusolverH);
cudaDeviceReset();

return(0);
}
```

# CUSOLVER CONT.

## CUDA SDK CONT.

### Example cusolverDnDsyevd() cont.

```
jobz = CUSOLVER_EIG_MODE_VECTOR;
uplo = CUBLAS_FILL_MODE_LOWER;
cusolver_status = cusolverDnCreate(&cusolverH); ← establish CUSOLVER context
assert(CUSOLVER_STATUS_SUCCESS == cusolver_status); ← error check
cudaStat = cudaMalloc ((void**)&d_A, sizeof(double) * lda * m);
cudaStat = cudaMalloc ((void**)&d_W, sizeof(double) * m);
cudaStat = cudaMalloc ((void**)&devInfo, sizeof(int));
cudaStat = cudaMemcpy(d_A, A, sizeof(double) * lda * m, cudaMemcpyHostToDevice);

cusolver_status = cusolverDnDsyevd_bufferSize(cusolverH, jobz, uplo, m, d_A, lda, d_W, &lwork);
cudaStat = cudaMalloc((void**)&d_work, sizeof(double)*lwork);
cusolver_status = cusolverDnDsyevd(cusolverH, jobz, uplo, m, d_A, lda, d_W, d_work, lwork, devInfo);
cudaStat = cudaDeviceSynchronize();

cudaStat = cudaMemcpy(W, d_W, sizeof(double)*m, cudaMemcpyDeviceToHost);
cudaStat = cudaMemcpy(V, d_A, sizeof(double)*lda*m, cudaMemcpyDeviceToHost);
cudaStat = cudaMemcpy(&info_gpu, devInfo, sizeof(int), cudaMemcpyDeviceToHost);
... do something with W[] and V[] ...
cudaFree(d_A); cudaFree(d_W); cudaFree(devInfo); cudaFree(d_work);
cusolverDnDestroy(cusolverH);
cudaDeviceReset();

return(0);
}
```

# CUSOLVER CONT.

## CUDA SDK CONT.

### Example cusolverDnDsyevd() cont.

```
jobz = CUSOLVER_EIG_MODE_VECTOR;
uplo = CUBLAS_FILL_MODE_LOWER;
cusolver_status = cusolverDnCreate(&cusolverH); ← establish CUSOLVER context
assert(CUSOLVER_STATUS_SUCCESS == cusolver_status); ← error check
cudaStat = cudaMalloc ((void**)&d_A, sizeof(double) * lda * m);
cudaStat = cudaMalloc ((void**)&d_W, sizeof(double) * m);
cudaStat = cudaMalloc ((void**)&devInfo, sizeof(int)); ← standard procedure
cudaStat = cudaMemcpy(d_A, A, sizeof(double) * lda * m, cudaMemcpyHostToDevice);

cusolver_status = cusolverDnDsyevd_bufferSize(cusolverH, jobz, uplo, m, d_A, lda, d_W, &lwork);
cudaStat = cudaMalloc((void**)&d_work, sizeof(double)*lwork);
cusolver_status = cusolverDnDsyevd(cusolverH, jobz, uplo, m, d_A, lda, d_W, d_work, lwork, devInfo);
cudaStat = cudaDeviceSynchronize();

cudaStat = cudaMemcpy(W, d_W, sizeof(double)*m, cudaMemcpyDeviceToHost);
cudaStat = cudaMemcpy(V, d_A, sizeof(double)*lda*m, cudaMemcpyDeviceToHost);
cudaStat = cudaMemcpy(&info_gpu, devInfo, sizeof(int), cudaMemcpyDeviceToHost);
... do something with W[] and V[] ...
cudaFree(d_A); cudaFree(d_W); cudaFree(devInfo); cudaFree(d_work);
cusolverDnDestroy(cusolverH);
cudaDeviceReset();

return(0);
}
```

# CUSOLVER CONT.

## CUDA SDK CONT.

### Example cusolverDnDsyevd() cont.

```
adjust jobz = CUSOLVER_EIG_MODE_VECTOR;
work uplo = CUBLAS_FILL_MODE_LOWER;
ar- cusolver_status = cusolverDnCreate(&cusolverH); ← establish CUSOLVER context
rays assert(CUSOLVER_STATUS_SUCCESS == cusolver_status); ← error check
 cudaStat = cudaMalloc ((void**)&d_A, sizeof(double) * lda * m);
 cudaStat = cudaMalloc ((void**)&d_W, sizeof(double) * m);
 cudaStat = cudaMalloc ((void**)&devInfo, sizeof(int));
 cudaStat = cudaMemcpy(d_A, A, sizeof(double) * lda * m, cudaMemcpyHostToDevice);

 ↓ cusolver_status = cusolverDnDsyevd_bufferSize(cusolverH, jobz, uplo, m, d_A, lda, d_W, &lwork);
 cudaStat = cudaMalloc((void**)&d_work, sizeof(double)*lwork);
 cusolver_status = cusolverDnDsyevd(cusolverH, jobz, uplo, m, d_A, lda, d_W, d_work, lwork, devInfo);
 cudaStat = cudaDeviceSynchronize();

 cudaStat = cudaMemcpy(W, d_W, sizeof(double)*m, cudaMemcpyDeviceToHost);
 cudaStat = cudaMemcpy(V, d_A, sizeof(double)*lda*m, cudaMemcpyDeviceToHost);
 cudaStat = cudaMemcpy(&info_gpu, devInfo, sizeof(int), cudaMemcpyDeviceToHost);
 ... do something with W[] and V[] ...
 cudaFree(d_A); cudaFree(d_W); cudaFree(devInfo); cudaFree(d_work);
 cusolverDnDestroy(cusolverH);
 cudaDeviceReset();

 return(0);
}
```

# CUSOLVER CONT.

## CUDA SDK CONT.

### Example cusolverDnDsyevd() cont.

```
adjust jobz = CUSOLVER_EIG_MODE_VECTOR;
work uplo = CUBLAS_FILL_MODE_LOWER;
ar- cusolver_status = cusolverDnCreate(&cusolverH); ← establish CUSOLVER context
rays assert(CUSOLVER_STATUS_SUCCESS == cusolver_status); ← error check
 cudaStat = cudaMalloc ((void**)&d_A, sizeof(double) * lda * m);
 cudaStat = cudaMalloc ((void**)&d_W, sizeof(double) * m);
 cudaStat = cudaMalloc ((void**)&devInfo, sizeof(int));
 cudaStat = cudaMemcpy(d_A, A, sizeof(double) * lda * m, cudaMemcpyHostToDevice);

 ↓
call cusolver_status = cusolverDnDsyevd_bufferSize(cusolverH, jobz, uplo, m, d_A, lda, d_W, &lwork);
solver cudaStat = cudaMalloc((void**)&d_work, sizeof(double)*lwork);
 cusolver_status = cusolverDnDsyevd(cusolverH, jobz, uplo, m, d_A, lda, d_W, d_work, lwork, devInfo);
 cudaStat = cudaDeviceSynchronize();

 cudaStat = cudaMemcpy(W, d_W, sizeof(double)*m, cudaMemcpyDeviceToHost);
 cudaStat = cudaMemcpy(V, d_A, sizeof(double)*lda*m, cudaMemcpyDeviceToHost);
 cudaStat = cudaMemcpy(&info_gpu, devInfo, sizeof(int), cudaMemcpyDeviceToHost);
... do something with W[] and V[] ...
 cudaFree(d_A); cudaFree(d_W); cudaFree(devInfo); cudaFree(d_work);
 cusolverDnDestroy(cusolverH);
 cudaDeviceReset();

 return(0);
}
```

# CUSOLVER CONT.

## CUDA SDK CONT.

### Example cusolverDnDsyevd() cont.

```
adjust jobz = CUSOLVER_EIG_MODE_VECTOR;
work uplo = CUBLAS_FILL_MODE_LOWER;
ar- cusolver_status = cusolverDnCreate(&cusolverH); ← establish CUSOLVER context
rays assert(CUSOLVER_STATUS_SUCCESS == cusolver_status); ← error check
 cudaStat = cudaMalloc ((void**)&d_A, sizeof(double) * lda * m);
 cudaStat = cudaMalloc ((void**)&d_W, sizeof(double) * m);
 cudaStat = cudaMalloc ((void**)&devInfo, sizeof(int));
 cudaStat = cudaMemcpy(d_A, A, sizeof(double) * lda * m, cudaMemcpyHostToDevice);

 ↓
call cusolver_status = cusolverDnDsyevd_bufferSize(cusolverH, jobz, uplo, m, d_A, lda, d_W, &lwork);
solver cudaStat = cudaMalloc((void**)&d_work, sizeof(double)*lwork);
 cusolver_status = cusolverDnDsyevd(cusolverH, jobz, uplo, m, d_A, lda, d_W, d_work, lwork, devInfo);
 cudaStat = cudaDeviceSynchronize();

 cudaStat = cudaMemcpy(W, d_W, sizeof(double)*m, cudaMemcpyDeviceToHost);
 cudaStat = cudaMemcpy(V, d_A, sizeof(double)*lda*m, cudaMemcpyDeviceToHost);
 cudaStat = cudaMemcpy(&info_gpu, devInfo, sizeof(int), cudaMemcpyDeviceToHost);
... do something with W[] and V[] ...
 cudaFree(d_A); cudaFree(d_W); cudaFree(devInfo); cudaFree(d_work);
 cusolverDnDestroy(cusolverH);
 cudaDeviceReset();

 return(0);
}
```

Diagram annotations:

- establish CUSOLVER context**: Points to the line `cusolver_status = cusolverDnCreate(&cusolverH);`.
- error check**: Points to the line `assert(CUSOLVER_STATUS_SUCCESS == cusolver_status);`.
- standard procedure**: Points to the line `cudaStat = cudaMemcpy(d_A, A, sizeof(double) * lda * m, cudaMemcpyHostToDevice);`.
- call solver**: Points to the call to `cusolverDnDsyevd`.
- back copy results**: Points to the line `cudaStat = cudaMemcpy(W, d_W, sizeof(double)*m, cudaMemcpyDeviceToHost);`.

# CUSOLVER CONT.

## CUDA SDK CONT.

### Example cusolverDnDsyevd() cont.

```
adjust jobz = CUSOLVER_EIG_MODE_VECTOR;
work uplo = CUBLAS_FILL_MODE_LOWER;
ar- cusolver_status = cusolverDnCreate(&cusolverH); ← establish CUSOLVER context
rays assert(CUSOLVER_STATUS_SUCCESS == cusolver_status); ← error check
 cudaStat = cudaMalloc ((void**)&d_A, sizeof(double) * lda * m);
 cudaStat = cudaMalloc ((void**)&d_W, sizeof(double) * m);
 cudaStat = cudaMalloc ((void**)&devInfo, sizeof(int));
 cudaStat = cudaMemcpy(d_A, A, sizeof(double) * lda * m, cudaMemcpyHostToDevice); ← standard procedure

call cusolver_status = cusolverDnDsyevd_bufferSize(cusolverH, jobz, uplo, m, d_A, lda, d_W, &lwork);
solver cudaStat = cudaMalloc((void**)&d_work, sizeof(double)*lwork);
 cusolver_status = cusolverDnDsyevd(cusolverH, jobz, uplo, m, d_A, lda, d_W, d_work, lwork, devInfo);
 cudaStat = cudaDeviceSynchronize(); ← call solver

 cudaStat = cudaMemcpy(W, d_W, sizeof(double)*m, cudaMemcpyDeviceToHost);
 cudaStat = cudaMemcpy(V, d_A, sizeof(double)*lda*m, cudaMemcpyDeviceToHost);
 cudaStat = cudaMemcpy(&info_gpu, devInfo, sizeof(int), cudaMemcpyDeviceToHost);
... do something with W[] and V[] ...
 cudaFree(d_A); cudaFree(d_W); cudaFree(devInfo); cudaFree(d_work);
 cusolverDnDestroy(cusolverH);
 cudaDeviceReset(); ← back copy results

 return(0); ← finalize
 }
```

- Very straightforward procedure offering significant speed-up for large scale problems at very minor porting effort
- Again, (float)  $\leftrightarrow$  (double) gap in terms of performance,  $\approx 1/64$  on consumer cards !
- A related problem of considerable interest may be addressed via `cusolverEigType_t` — generalized symmetric-definite eigenvalue problem (section 2.2.1.4)  
$$A x = \lambda B x$$
- Sparse problems covered too — `cuSolverSP`

# MORE LIBRARIES

## CUDA SDK CONT.

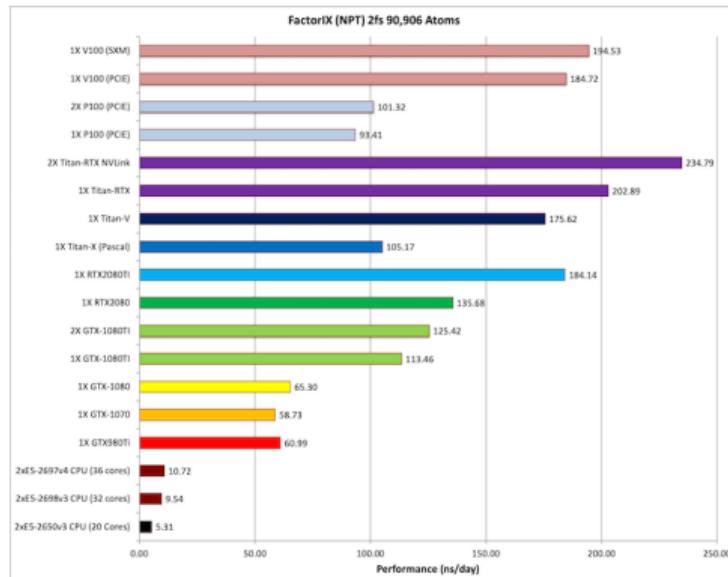
- CUFFT — CUDA Fast Fourier Transform
- CURAND — CUDA random number generation library
- CUDNN & TENSORRT — Deep Learning, training & inference
- CUSPARSE — Basic linear algebra for sparse matrices
- NVBLAS — Another replacement for BLAS calls with little effort of porting (re-linking or LD\_PRELOAD)
- NVGRAPH — Big Data analytics via graph problems

→ <https://docs.nvidia.com/cuda/cufft/index.html>  
→ <https://docs.nvidia.com/cuda/curand/index.html>  
→ <https://docs.nvidia.com/deeplearning/sdk/index.html>

# NUMERICAL ACCURACY & PERFORMANCE

## CUDA SDK CONT.

How come that scientific apps, e.g. AMBER's pmemd.cuda are doing so well on consumer grade GPUs ?



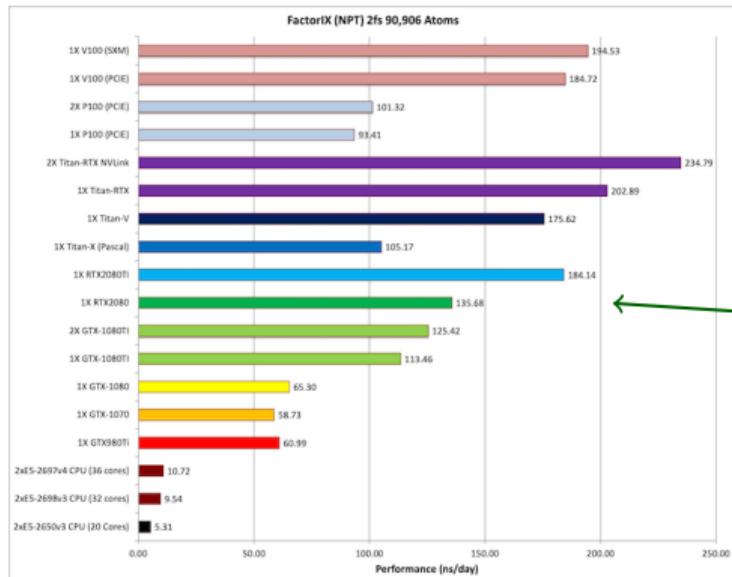
→ <http://ambermd.org/GPUPerformance.php>

→ <https://www.sciencedirect.com/science/article/pii/S0010465512003098>

# NUMERICAL ACCURACY & PERFORMANCE

CUDA SDK CONT.

How come that scientific apps, e.g. AMBER's pmemd.cuda are doing so well on consumer grade GPUs ?



Combine several 32bit types (float, int) to approximate double

→ <http://ambermd.org/GPUPerformance.php>

→ <https://www.sciencedirect.com/science/article/pii/S0010465512003098>

# NUMERICAL ACCURACY & PERFORMANCE CONT.

## CUDA SDK CONT.

Dekker (1971) and Knuth (1969)

### Quasi-double precision method

- Combine two floating-point numbers to express one variable

$$\begin{array}{ll} x + y = z + zz & x = 10.2, y = 0.345 \\ \text{(where } |z| / |zz| < 2^{-24} \text{)} & \text{(three significant digits)} \\ z = \text{fl}(x + y) & z = 10.5 \\ w = \text{fl}(z - x) & w = 0.3 \\ zz = \text{fl}(y - w) & zz = 0.045 \\ & (z + zz = 10.545, \text{ no cancellation}) \end{array}$$

x, y, z, zz : FP32 variable

fl(operation) : FP32 arithmetic operation

- More than twice operations are needed for similar accuracy to double-precision  
⇒ Only works for consumer GPUs

(slide courtesy of Prof Tetsu Narumi, UEC, Tokyo)

→ <https://link.springer.com/article/10.1007/BF01397083>

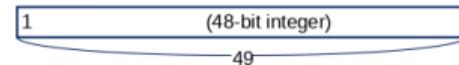
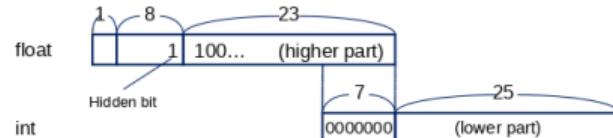
→ [https://en.wikipedia.org/wiki/The\\_Art\\_of\\_Computer\\_Programming](https://en.wikipedia.org/wiki/The_Art_of_Computer_Programming)

# NUMERICAL ACCURACY & PERFORMANCE CONT.

## CUDA SDK CONT.

### Combine floating- and fixed-point

- and INT32 for one variable
- Initialized with a large number beforehand
- Similar to 48-bit integer



Tetsu Narumi et al., International Journal of Computational Methods, vol. 8, No. 3, pp. 561-581 (2011)

8

(slide courtesy of Prof Tetsu Narumi, UEC, Tokyo)

→ doi:10.1142/S0219876211002708

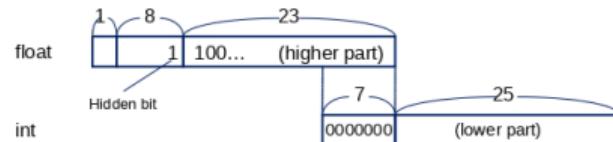
→ doi:10.1016/j.cpc.2012.09.022

# NUMERICAL ACCURACY & PERFORMANCE CONT.

## CUDA SDK CONT.

### Combine floating- and fixed-point

- and INT32 for one variable
- Initialized with a large number beforehand
- Similar to 48-bit integer



Tetsu Narumi et al., International Journal of Computational Methods, vol. 8, No. 3, pp. 561-581 (2011)

All elementary operations, +,  
-, \*

(slide courtesy of Prof Tetsu Narumi, UEC, Tokyo)

→ doi:10.1142/S0219876211002708

→ doi:10.1016/j.cpc.2012.09.022

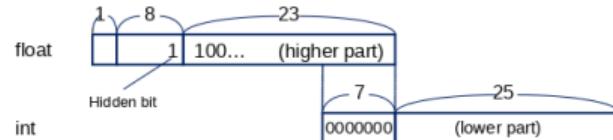
# NUMERICAL ACCURACY & PERFORMANCE CONT.

## CUDA SDK CONT.

Similar strategy in AMBER's pmemd.cuda

### Combine floating- and fixed-point

- and INT32 for one variable
- Initialized with a large number beforehand
- Similar to 48-bit integer



Tetsu Narumi et al., International Journal of Computational Methods, vol. 8, No. 3, pp. 561-581 (2011)

All elementary operations, +, -, \*,

(slide courtesy of Prof Tetsu Narumi, UEC, Tokyo)

→ doi:10.1142/S0219876211002708

→ doi:10.1016/j.cpc.2012.09.022

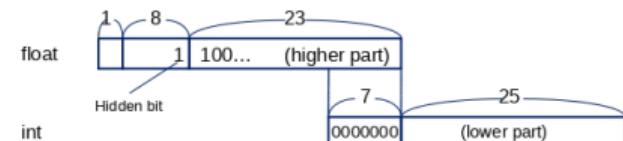
# NUMERICAL ACCURACY & PERFORMANCE CONT.

## CUDA SDK CONT.

Similar strategy in AMBER's pmemd.cuda

### Combine floating- and fixed-point

- INT32 and INT32 for one variable
- Initialized with a large number beforehand
- Similar to 48-bit integer



Tetsu Narumi et al., International Journal of Computational Methods, vol. 8, No. 3, pp. 561-581 (2011)

Also see Mandelbrot example in SDK,  
5\_Domain\_Specific

All elementary operations, +, -, \*, /

(slide courtesy of Prof Tetsu Narumi, UEC, Tokyo)

→ doi:10.1142/S0219876211002708

→ doi:10.1016/j.cpc.2012.09.022

## 5\_DOMAIN\_SPECIFIC/MARCHINGCUBES

CUDA SDK CONT.

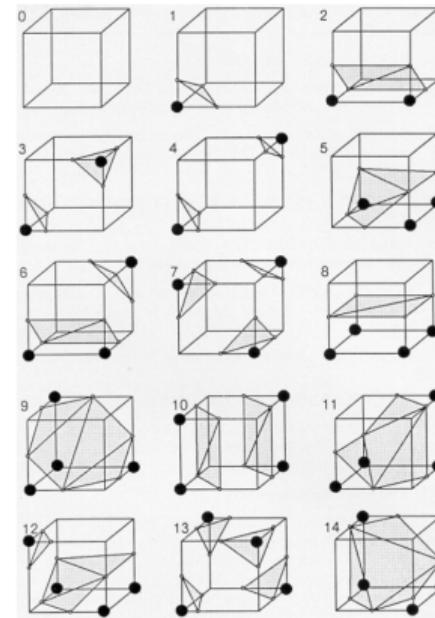
Practical example, of frequent use in biophysical chemistry

- Extracts an isosurface from a volumetric dataset using the 'marching cubes algorithm'
- OpenGL interoperation
- Chosen example is the electron density of C<sub>60</sub>

→ [./run\\_marchingCubes.sh](#)

## 5\_DOMAIN\_SPECIFIC/MARCHINGCUBES CONT.

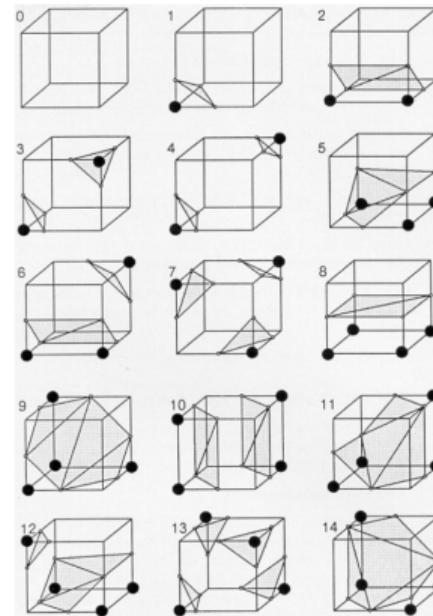
CUDA SDK CONT.



→ <https://dl.acm.org/citation.cfm?doid=37402.37422>

## 5\_DOMAIN\_SPECIFIC/MARCHINGCUBES CONT.

CUDA SDK CONT.

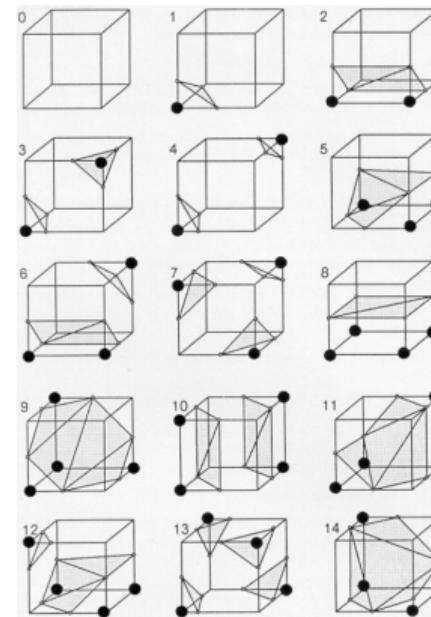


Assign vertices 0 if below iso-value, 1 if above

→ <https://dl.acm.org/citation.cfm?doid=37402.37422>

## 5\_DOMAIN\_SPECIFIC/MARCHINGCUBES CONT.

CUDA SDK CONT.



Neighbouring  
vertices of 0,  
1 assignment  
determine iso-  
surface edges

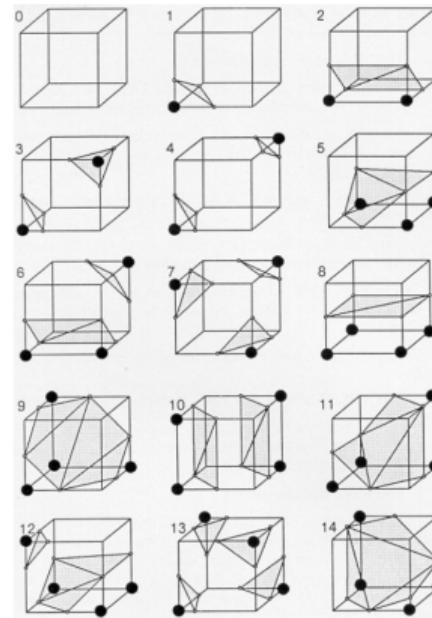
Assign ver-  
tices 0 if below  
iso-value, 1 if  
above

→ <https://dl.acm.org/citation.cfm?doid=37402.37422>

## 5\_DOMAIN\_SPECIFIC/MARCHINGCUBES CONT.

CUDA SDK CONT.

Neighbouring vertices of 0, 1 assignment determine iso-surface edges



14 elementary cases, the rest (256) from symmetry

Assign vertices 0 if below iso-value, 1 if above

→ <https://dl.acm.org/citation.cfm?doid=37402.37422>

## 5\_DOMAIN\_SPECIFIC/MARCHINGCUBES CONT.

CUDA SDK CONT.

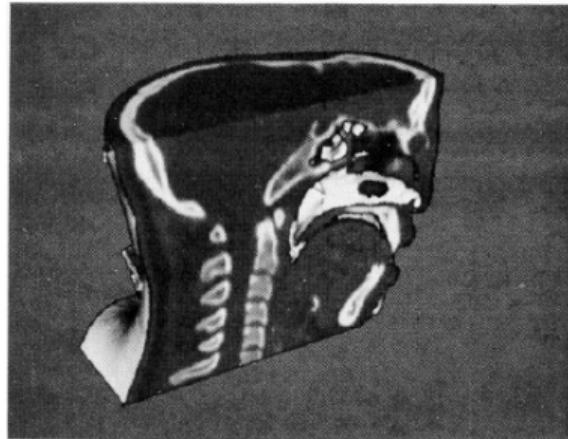


Figure 11. Sagittal Cut with Texture Mapping.

CT data, shows the slice data in relation to the constructed

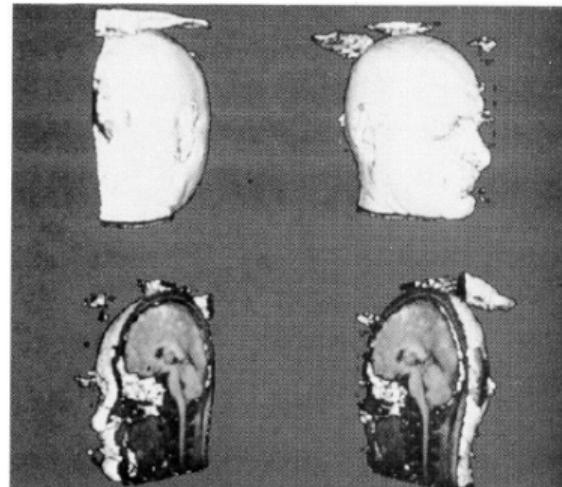


Figure 12. Rotated Sequence of Cut MR Brain.

Interesting for medical  
images (CT,MR)

→ <https://dl.acm.org/citation.cfm?doid=37402.37422>

# PROFILING CUDA CODE

## CUDA SDK CONT.

The main profiling tool in CUDA nowadays is NVIDIA Nsight Compute

- nsys nvprof (quick and easy)
- nsys/nsys-ui (for system traces and timeline analysis – streams)
- ncu/ncu-ui (for detailed kernel optimization)

→ <https://docs.nvidia.com/nsight-compute/NsightComputeCli/index.html>

# PROFILING CUDA CODE

## CUDA SDK CONT.

The main profiling tool in CUDA nowadays is NVIDIA Nsight Compute

- nsys nvprof (quick and easy)
- nsys/nsys-ui (for system traces and timeline analysis – streams)
- ncu/ncu-ui (for detailed kernel optimization)

```
cuda-zen sh@n3073-004:~$ nvcc ./mmmm_example_1.cu
cuda-zen sh@n3073-004:~$ ncu -f -set full -o profile ./a.out
cuda-zen sh@gui3068-009:~$ ncu-ui ./profile.ncu-rep
```

→ <https://docs.nvidia.com/nsight-compute/NsightComputeCli/index.html>

# PROFILING CUDA CODE

## CUDA SDK CONT.

The main profiling tool in CUDA nowadays is NVIDIA Nsight Compute

- nsys nvprof (quick and easy)
- nsys/nsys-ui (for system traces and timeline analysis – streams)
- ncu/ncu-ui (for detailed kernel optimization)

### 1. Compilation

```
cuda-zen sh@n3073-004:~$ nvcc ./mmmm_example_1.cu
cuda-zen sh@n3073-004:~$ ncu -f -set full -o profile ./a.out
cuda-zen sh@gui3068-009:~$ ncu-ui ./profile.ncu-rep
```

→ <https://docs.nvidia.com/nsight-compute/NsightComputeCli/index.html>

# PROFILING CUDA CODE

## CUDA SDK CONT.

The main profiling tool in CUDA nowadays is NVIDIA Nsight Compute

- nsys nvprof (quick and easy)
- nsys/nsys-ui (for system traces and timeline analysis – streams)
- ncu/ncu-ui (for detailed kernel optimization)

### 1. Compilation

```
cuda-zen sh@n3073-004:~$ nvcc ./mmmm_example_1.cu
cuda-zen sh@n3073-004:~$ ncu -f -set full -o profile ./a.out
cuda-zen sh@gui3068-009: $ ncu-ui ./profile.ncu-rep
```

### 2. Write profile

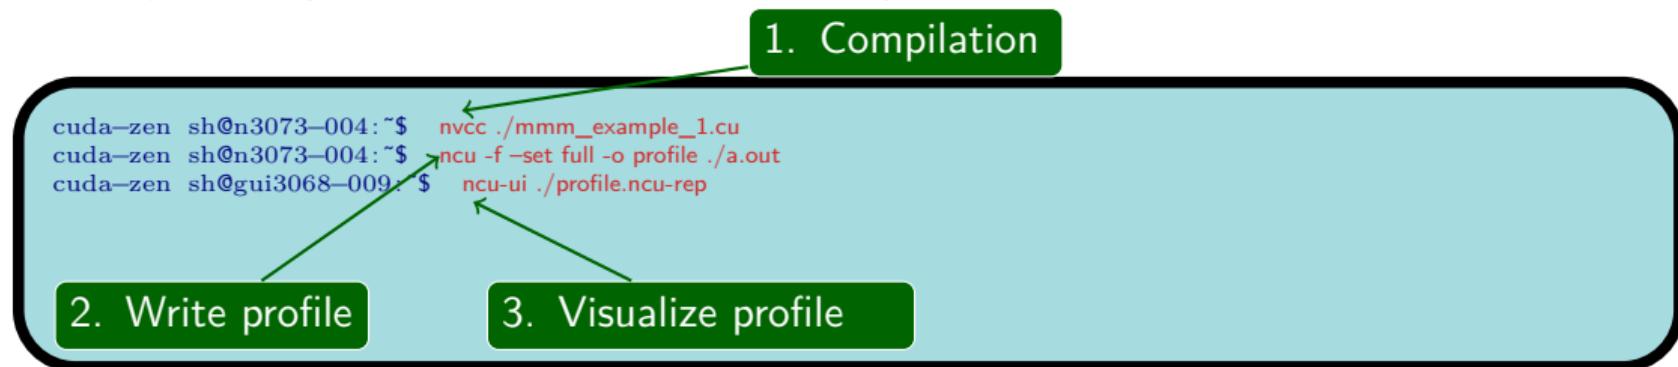
→ <https://docs.nvidia.com/nsight-compute/NsightComputeCli/index.html>

# PROFILING CUDA CODE

## CUDA SDK CONT.

The main profiling tool in CUDA nowadays is NVIDIA Nsight Compute

- nsys nvprof (quick and easy)
- nsys/nsys-ui (for system traces and timeline analysis – streams)
- ncu/ncu-ui (for detailed kernel optimization)

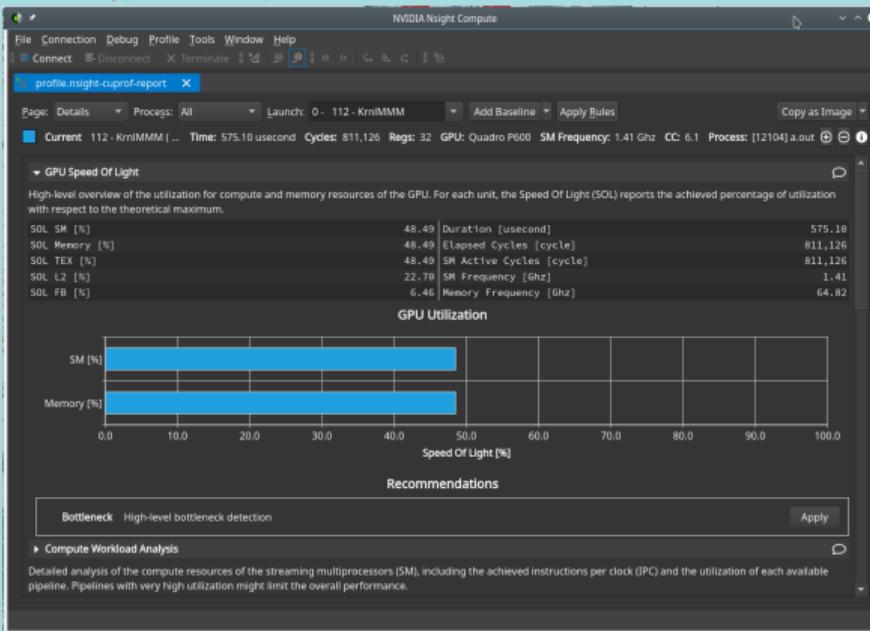


→ <https://docs.nvidia.com/nsight-compute/NsightComputeCli/index.html>

# PROFILING CUDA CODE CONT.

## CUDA SDK CONT.

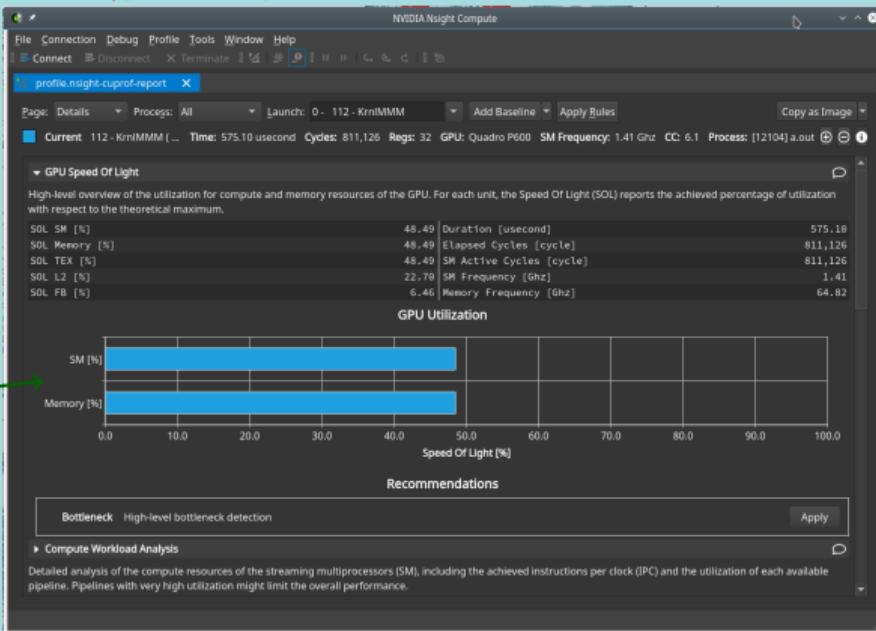
```
cuda-zen sh@gui3068-009:~$ ncu-ui ./profile.ncu-report
```



# PROFILING CUDA CODE CONT.

## CUDA SDK CONT.

```
cuda-zen sh@gui3068-009:~$ ncu-ui ./profile.ncu-report
```

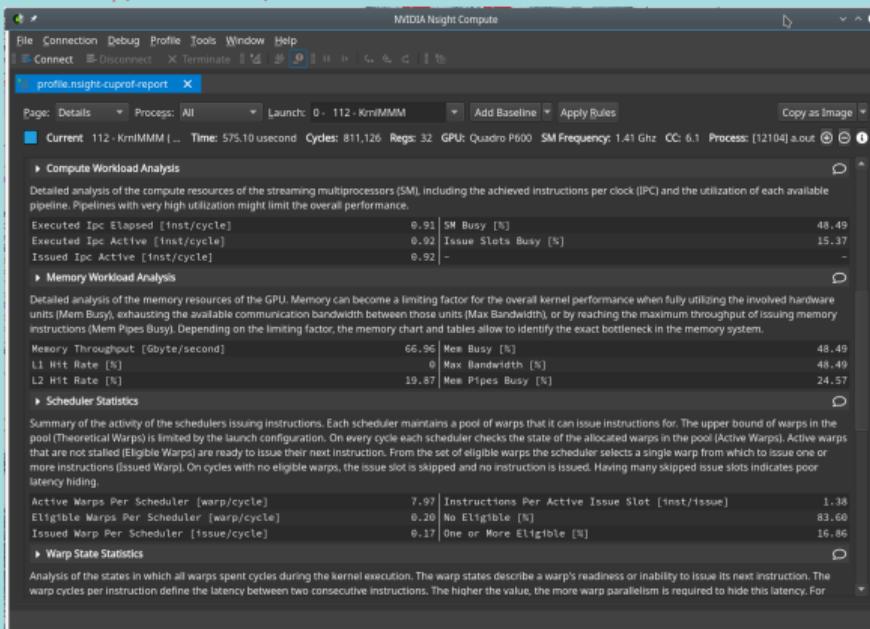


GPU features evaluated w.r.t theoretical max (SOL)

# PROFILING CUDA CODE CONT.

## CUDA SDK CONT.

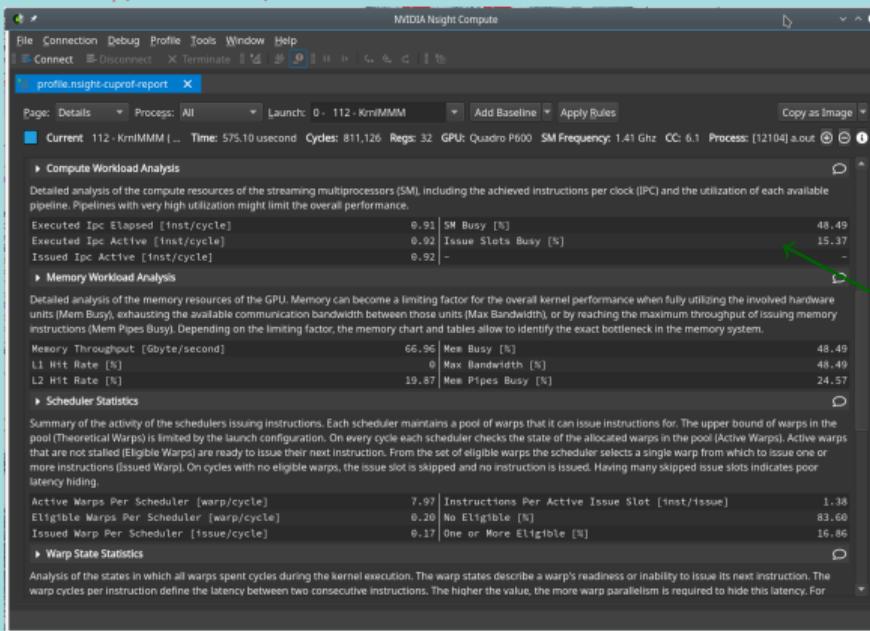
```
cuda-zen sh@gui3068-009:~$ ncu-ui ./profile.ncu-report
```



# PROFILING CUDA CODE CONT.

## CUDA SDK CONT.

```
cuda-zen sh@gui3068-009:~$ ncu-ui ./profile.ncu-report
```

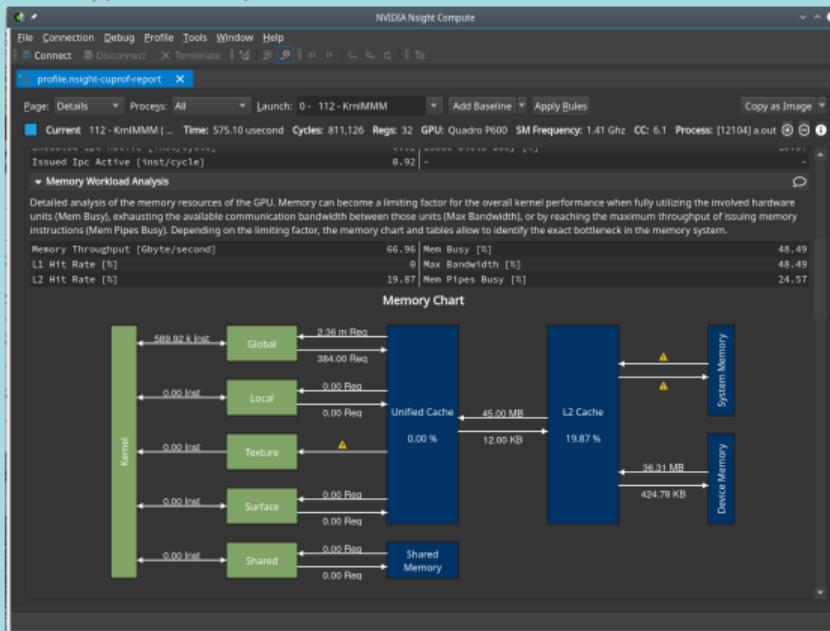


Different categories' evaluation

# PROFILING CUDA CODE CONT.

## CUDA SDK CONT.

```
cuda-zen sh@gui3068-009:~$ ncu-ui ./profile.ncu-report
```

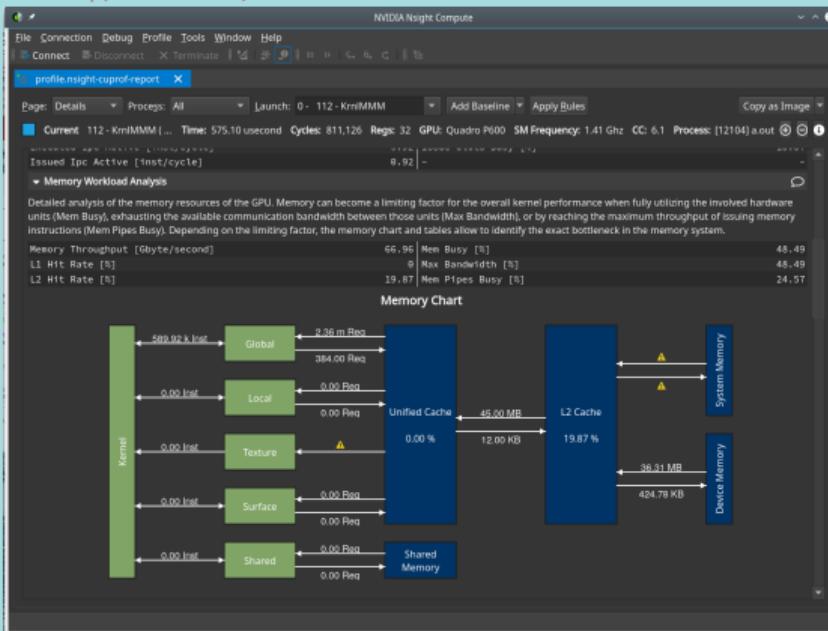


# PROFILING CUDA CODE CONT.

## CUDA SDK CONT.

cuda-zen sh@gui3068-009:~\$

ncu-ui ./profile.ncu-report



- The Heterogeneous Interface for Portability (HIP) is AMD's dedicated GPU programming environment (e.g. to program MI250X/MI300X devices)
- Very similar to CUDA
- HIP code will also run on NVIDIA platforms
- HIP forms another ecosystem with tools, libraries, etc
- Every basic CUDA construct has a direct counterpart in HIP

# HIP CONT.

## HIP EXAMPLE

### Vector Addition with HIP

```
// HIP kernel. Each thread takes care of one element of c;
__global__ void vecAdd(double *a, double *b, double *c, int n)
{
 // Get our global thread ID
 int id = (blockIdx.x * blockDim.x) + threadIdx.x;
 // Make sure we do not go beyond bounds
 if (id < n)
 c[id] = a[id] + b[id];
}

int main()
{
 ...
 // Execute the kernel with n threads
 hipLaunchKernelGGL(vecAdd, dim3(gridSize), dim3(blockSize), 0, 0, d_a, d_b, d_c, n);
 hipDeviceSynchronize();
}
```

# HIP CONT.

## HIP EXAMPLE

### Vector Addition with HIP

```
// HIP kernel. Each thread takes care of one element of c;
__global__ void vecAdd(double *a, double *b, double *c, int n)
{
 // Get our global thread ID
 int id = (blockIdx.x * blockDim.x) + threadIdx.x;
 // Make sure we do not go beyond bounds
 if (id < n)
 c[id] = a[id] + b[id];
}

int main()
{
 ...
 // Execute the kernel with n threads
 hipLaunchKernelGGL(vecAdd, dim3(gridSize), dim3(blockSize), 0, 0, d_a, d_b, d_c, n);
 hipDeviceSynchronize();
}
```

same kernel declaration specifiers

→ <https://developer.amd.com/wp-content/resources/R0Cm%20Learning%20Centre/chapter3/HIP-Coding-3.pdf>

# HIP CONT.

## HIP EXAMPLE

### Vector Addition with HIP

```
// HIP kernel. Each thread takes care of one element of c;
__global__ void vecAdd(double *a, double *b, double *c, int n)
{
 // Get our global thread ID
 int id = (blockIdx.x * blockDim.x) + threadIdx.x;
 // Make sure we do not go beyond bounds
 if (id < n)
 c[id] = a[id] + b[id];
}

int main()
{
 ...
 // Execute the kernel with n threads
 hipLaunchKernelGGL(vecAdd, dim3(gridSize), dim3(blockSize), 0, 0, d_a, d_b, d_c, n);
 hipDeviceSynchronize();
}
```

same kernel declaration specifiers

same built-in variables,  
e.g. threadIdx.x=0,1,2...

# HIP CONT.

## HIP EXAMPLE

### Vector Addition with HIP

```
// HIP kernel. Each thread takes care of one element of c;
__global__ void vecAdd(double *a, double *b, double *c, int n)
{
 // Get our global thread ID
 int id = (blockIdx.x * blockDim.x) + threadIdx.x;
 // Make sure we do not go beyond bounds
 if (id < n)
 c[id] = a[id] + b[id];
}

int main()
{
 ...
 // Execute the kernel with n threads
 hipLaunchKernelGGL(vecAdd, dim3(gridSize), dim3(blockSize), 0, 0, d_a, d_b, d_c, n);
 hipDeviceSynchronize();
}
```

same kernel declaration specifiers

same built-in variables,  
e.g. threadIdx.x=0,1,2...

different kernel execution configuration

# HIP CONT.

## COMPILE AND RUN AND MONITOR

```
(zen3) [sh@some - mi250x ~]$ hipcc vadd_hip.cpp
(zen3) [sh@some - mi250x ~]$./a.out

0 100.000000
1 100.000000
2 100.000000
3 100.000000
4 100.000000
5 100.000000
6 100.000000
7 100.000000
8 100.000000
9 100.000000
10 100.000000
11 100.000000
...
99 100.000000
```

- <https://developer.amd.com/wp-content/resources/R0Cm%20Learning%20Centre/chapter3/HIP-Coding-3.pdf>
- [https://rocmdocs.amd.com/en/latest/Programming\\_Guides/HIP-porting-guide.html](https://rocmdocs.amd.com/en/latest/Programming_Guides/HIP-porting-guide.html)
- <https://forums.extremehw.net/topic/782-solvedimage-upload-broken-on-firefox>

# HIP CONT.

## COMPILE AND RUN AND MONITOR

```
(zen3) [sh@some - mi250x ~]$ hipcc vadd_hip.cpp
(zen3) [sh@some - mi250x ~]$./a.out

0 100.000000
1 100.000000
2 100.000000
3 100.000000
4 100.000000
5 100.000000
6 100.000000
7 100.000000
8 100.000000
9 100.000000
10 100.000000
11 100.000000
...
99 100.000000
```

```
tictoc@TickTockArch $ rocm-smi

=====
=====ROCM System Management Interface=====
=====
GPU Temp AvgPwr SCLK MCLK Fan Perf PwrCap VRAM% GPU%
0 42.0c 294.0W 1925Mhz 1000Mhz 0.0% high 450.0W 17% 100%
1 50.0c 219.0W 1772Mhz 1000Mhz 0.0% high 220.0W 4% 100%
2 45.0c 273.0W 1925Mhz 1000Mhz 0.0% high 300.0W 4% 100%
3 55.0c 163.0W 1700Mhz 1000Mhz 0.0% high 170.0W 4% 100%
=====
=====End of ROCm SMI Log =====
_
tictoc@TickTockArch $ _
```

- <https://developer.amd.com/wp-content/resources/ROCM%20Learning%20Centre/chapter3/HIP-Coding-3.pdf>
- [https://rocmdocs.amd.com/en/latest/Programming\\_Guides/HIP-porting-guide.html](https://rocmdocs.amd.com/en/latest/Programming_Guides/HIP-porting-guide.html)
- <https://forums.extremehw.net/topic/782-solvedimage-upload-broken-on-firefox>

## TAKE HOME MESSAGES

- CUDA can be easily used in Fortran programs

## TAKE HOME MESSAGES

- CUDA can be easily used in Fortran programs
- BLAS calls may be substituted with calls to CUBLAS potentially offering significant acceleration

## TAKE HOME MESSAGES

- CUDA can be easily used in Fortran programs
- BLAS calls may be substituted with calls to CUBLAS potentially offering significant acceleration
- Basic LAPACK calls may be replaced with CUSOLVER routines

## TAKE HOME MESSAGES

- CUDA can be easily used in Fortran programs
- BLAS calls may be substituted with calls to CUBLAS potentially offering significant acceleration
- Basic LAPACK calls may be replaced with CUSOLVER routines
- Sparse problems are supported with specific libraries, other basic tasks as well, CUFFT, CURAND, CUDNN

## TAKE HOME MESSAGES

- CUDA can be easily used in Fortran programs
- BLAS calls may be substituted with calls to CUBLAS potentially offering significant acceleration
- Basic LAPACK calls may be replaced with CUSOLVER routines
- Sparse problems are supported with specific libraries, other basic tasks as well, CUFFT, CURAND, CUDNN
- Taking advantage of fast FP32 operations with quasi-FP64 accuracy on consumer grade cards requires a couple of algorithmic changes

## TAKE HOME MESSAGES

- CUDA can be easily used in Fortran programs
- BLAS calls may be substituted with calls to CUBLAS potentially offering significant acceleration
- Basic LAPACK calls may be replaced with CUSOLVER routines
- Sparse problems are supported with specific libraries, other basic tasks as well, CUFFT, CURAND, CUDNN
- Taking advantage of fast FP32 operations with quasi-FP64 accuracy on consumer grade cards requires a couple of algorithmic changes
- NVIDIA's NSIGHT compute is an advanced profiling tool for in-depth kernel optimization

## TAKE HOME MESSAGES

- CUDA can be easily used in Fortran programs
- BLAS calls may be substituted with calls to CUBLAS potentially offering significant acceleration
- Basic LAPACK calls may be replaced with CUSOLVER routines
- Sparse problems are supported with specific libraries, other basic tasks as well, CUFFT, CURAND, CUDNN
- Taking advantage of fast FP32 operations with quasi-FP64 accuracy on consumer grade cards requires a couple of algorithmic changes
- NVIDIA's NSIGHT compute is an advanced profiling tool for in-depth kernel optimization
- AMD's HIP is very similar to CUDA and kernel code can be directly interchanged