

TUTORIAL 2: EXPLORING CUDA

Siegfried Höfinger

ASC Research Center, Vienna University of Technology

October 19, 2025

→ <https://tinyurl.com/cudafordummies/ii/t/notes-t2.pdf>

OUTLINE

SECOND STEPS WITH CUDA

SECOND STEPS WITH CUDA

CUDA SDK

Exercise

- Q1)** *Using again assert() in the kernel code of the SDK sample 0_Introduction/simplePrintf, how could we quickly identify whether the shape of the used threadblock is 1D/2D/3D ?*

10 min

SECOND STEPS WITH CUDA CONT.

CUDA SDK CONT.

A1) A useful threadblock must at least consist of a single thread, hence have $\text{threadIdx.x/y/z} = 1/0/0$. If the threadblock is 1D, then for all other threads threadIdx.y/z will remain 0, which can be probed from `assert(threadIdx.z == 0)` and `assert(threadIdx.y == 0)` (see below variant for download and own experiments).

n.b. first change to private dir (or create it from scratch)

`cd wherever_the_SDK_may_be/Samples/0_Introduction/my_simplePrintf`

then edit `simplePrintf.cu` therein (or download the variant and move/copy it to `simplePrintf.cu`) then compile and run the modified version

`cd build`

`cmake ..`

`make`

`./simplePrintf`

→ https://tinyurl.com/cudafordummies/ii/t/simplePrintf_v4.cu

SECOND STEPS WITH CUDA CONT.

CUDA SDK CONT.

Exercise

- Q2)** Verify the maximum available bandwidth of ≈ 25 GB/s for host-to-device transfers on the A100 architecture using the demo sample bandwidthTest. Examine what could be optionally tested (calling .../wherever/it/is/bandwidthTest --help). Is the confirmed limit value a reasonable estimate ?

10 min

SECOND STEPS WITH CUDA CONT.

CUDA SDK CONT.

- A2)** *Use the command "which nvcc" to get an idea of the install location of the sample program bandwidthTest, locate it and run it. Playing around with various command line arguments does actually confirm the limit of 25 GB/s and this is probably all we can expect from Gen4 PCIe.*

SECOND STEPS WITH CUDA CONT.

CUDA SDK CONT.

Exercise

- Q3)** *Use the low-level profiling toolchain,
nsys nvprof .../wherever/it/is/bandwidthTes --htod
and examine the output. Based on this, can we get confirmation of the
bandwidth computed by the demo sample ?*

15 min

SECOND STEPS WITH CUDA CONT.

CUDA SDK CONT.

A3) *From the profile we get confirmation of the amount of transferred data per copy to be 33.554 MB, i.e. 0.033554 GB. The corresponding time for a single transfer read from the profile is ≈ 1607320 ns, i.e. 0.001607320 s, so the resulting bandwidth then is, $0.033554 / 0.001607320 = 20.9$ GB/s which is pretty close to what the demo sample reported...*

SECOND STEPS WITH CUDA CONT.

CUDA LIBRARIES: CUSOLVER

Exercise

Q4) *Solve the eigenvalue/eigenvector problem of the following matrix*

$$\mathbf{M} = \begin{pmatrix} 1.96 & -6.49 & -0.47 & -7.20 & -0.65 \\ -6.49 & 3.80 & -6.39 & 1.50 & -6.34 \\ -0.47 & -6.39 & 4.17 & -1.51 & 2.67 \\ -7.20 & 1.50 & -1.51 & 5.70 & 1.80 \\ -0.65 & -6.34 & 2.67 & 1.80 & -7.10 \end{pmatrix}$$

using cusolverDN

20 min

SECOND STEPS WITH CUDA CONT.

CUDA LIBRARIES: CUSOLVER CONT.

A4)

Eigenvalues are,

-11.07 -6.23 0.86 8.87 16.09

*and eigenvectors those already mentioned in Q1 of the hands-on session 4
(see below version for download)*

→ https://tinyurl.com/cudafordummies/ii/t/chck_cusolver_syevd.cu
→ <https://tinyurl.com/cudafordummies/ii/ho4/notes-ho4.pdf>

SECOND STEPS WITH CUDA CONT.

NVIDIA NSIGHT COMPUTE CLI

Exercise

Q5) *Perform the 3-step profiling process using NVIDIA Nsight Compute CLI and identify reasons for improvement when considering the series mmm_example_[1-3].cu*

- <https://docs.nvidia.com/nsight-compute/NsightComputeCli/index.html>
- https://tinyurl.com/cudafordummies/ii/t/mmm_example_1.cu
- https://tinyurl.com/cudafordummies/ii/t/mmm_example_2.cu
- https://tinyurl.com/cudafordummies/ii/t/mmm_example_3.cu

15 min

SECOND STEPS WITH CUDA CONT.

NVIDIA NSIGHT COMPUTE CLI CONT.

A5)

- i) *Adding a baseline (using the most simplistic approach) and doing relative comparisons to this baseline is a very straightforward way of identifying improvements in different implementations; for example, mmm_example_2.cu versus mmm_example_1.cu (baseline) immediately points out improvements in terms of memory access, but degradation in SM performance mainly because of fewer instructions per cycle, which is a consequence of a reduced number of warps (eligible as well as active) per cycle; also simply looking at Duration [msecond] reveals a factor of 5x for the exe-time while memory throughput has dropped from 5 GB/s to 1 GB/s;*
- ii) *mmm_example_3.cu versus mmm_example_1.cu (baseline) demonstrates a significant reduction in exe-time (Duration) at better SM utilization and improved memory access (70% and 90% of theoretical max, aka SOL); best roofline position of all; here we get an increased number of instructions/cycle and better memory throughput of 7 GB/s and a significant involvement of shared memory;*

SECOND STEPS WITH CUDA CONT.

CUDA STREAMS

Exercise

- Q6)** *Check whether we could make use of CUDA managed unified memory, i.e. `cudaMallocManaged()`, within applications using CUDA streams, for example `stream_test.cu`*
Evaluate concurrency with the help of nsys profile `a.out`

→ https://tinyurl.com/cudafordummies/ii/t/stream_test.cu

15 min

SECOND STEPS WITH CUDA CONT.

CUDA STREAMS CONT.

A6)

Yes, in principle (see below version for download).

However, care must be taken to synchronize individual streams before managed unified memory can be accessed on the host in the usual manner.

→ https://tinyurl.com/cudafordummies/ii/t/stream_test_v2.cu